

Stata Learning Modules

Ethan Deng

Version 1.0

April 20, 2015

1 Fundamentals of Using Stata (part I)

1.1 A Sample Stata Session

Please refer to [Manuals13](#).

1.2 Descriptive information and statistics

This module shows common commands for showing descriptive information and descriptive statistics about data files.

1.2.1 Getting an overview of your file

The `sysuse` command loads a specified Stata-format dataset that was shipped with Stata. Here we will use the `auto` data file.

```
1 sysuse auto
```

The `describe` command shows you basic information about a Stata data file. As you can see, it tells us the number of observations in the file, the number of variables, the names of the variables, and more.

```
1 describe
2 Contains data from auto.dta
3 obs:          74
4 vars:          12                      17 Feb 1999 10:49
5 size:          3,108 (99.6% of memory free)
6 -----
7 1. make        str17  %17s
8 2. price       int    %9.0g
9 3. mpg         byte   %9.0g
10 4. rep78       byte   %9.0g
11 5. hdroom      float  %9.0g
12 6. trunk       byte   %9.0g
13 7. weight      int    %9.0g
14 8. length      int    %9.0g
15 9. turn        byte   %9.0g
16 10. displ      int    %9.0g
17 11. gratio     float  %9.0g
18 12. foreign    byte   %9.0g
19 -----
20 Sorted by:
```

The `codebook` command is a great tool for getting a quick overview of the variables in the data file. It produces a kind of electronic codebook from the data file. Have a look at what it produces below.

```

1 codebook
2 make ----- (unlabeled)
3           type:  string (str17)
4
5           unique values:  74                coded missing:  0 / 74
6
7           examples:  "Cad. Deville"
8                     "Dodge Magnum"
9                     "Merc. XR-7"
10                    "Pont. Catalina"
11
12           warning:  variable has embedded blanks
13
14 price ----- (unlabeled)
15           type:  numeric (int)
16
17           range:  [3291,15906]                units:  1
18           unique values:  74                coded missing:  0 / 74
19
20           mean:  6165.26
21           std. dev:  2949.5
22
23           percentiles:      10%      25%      50%      75%      90%
24                           3895      4195      5006.5      6342      11385
25 //(omitted)

```

Another useful command for getting a quick overview of a data file is the `inspect` command. Here is what the `inspect` command produces for the auto data file.

```

1 inspect
2 price:
3 -----
4
5 | #           Negative      Total   Integers   Non-
6 | #           Zero         -       -         Integers
7 | #           Positive     74      74        -
8 | #           -----
9 | #           Total       74      74        -
10 | # # . . .           Missing   -
11 +-----
12 3291           15906           74
13 (74 unique values)
14
15 mpg:
16 -----
17
18 | #           Negative      Total   Integers   Non-
19 | #           Zero         -       -         Integers
20 | #           Positive     74      74        -
21 | # #           -----
22 | # # #           Total     74      74        -
23 | # # # # .       Missing   -
24 +-----
25 12           41           74

```

```

26 (21 unique values)
27 //(omitted)

```

The `list` command is useful for viewing all or a range of observations. Here we look at *make*, *price*, *mpg*, *rep78* and *foreign* for the first 10 observations.

```

1 list make price mpg rep78 foreign in 1/10
2
3      make      price      mpg      rep78      foreign
4  1.   Dodge Magnum   5886      16         2         0
5  2.    Datsun 510    5079      24         4         1
6  3.   Ford Mustang   4187      21         3         0
7  4.  Linc. Versailles 13466      14         3         0
8  5.   Plym. Sapporo   6486      26         .         0
9  6.   Plym. Arrow    4647      28         3         0
10 7.   Cad. Eldorado  14500      14         2         0
11 8.    AMC Spirit    3799      22         .         0
12 9.   Pont. Catalina   5798      18         4         0
13 10.  Chev. Nova     3955      19         3         0

```

1.2.2 Creating tables

The `tabulate` command is useful for obtaining frequency tables. Below, we make a table for *rep78* and a table for *foreign*. The command can also be shortened to `tab`.

```

1 tabulate rep78
2      rep78 |      Freq.      Percent      Cum.
3 -----+-----
4          1 |          2         2.90         2.90
5          2 |          8        11.59        14.49
6          3 |         30        43.48        57.97
7          4 |         18        26.09        84.06
8          5 |         11        15.94       100.00
9 -----+-----
10         Total |         69       100.00
11 tabulate foreign
12      foreign |      Freq.      Percent      Cum.
13 -----+-----
14          0 |         52        70.27        70.27
15          1 |         22        29.73       100.00
16 -----+-----
17         Total |         74       100.00

```

The `tab1` command can be used as a shortcut to request tables for a series of variables (instead of typing the `tabulate` command over and over again for each variable of interest).

```

1 tab1 rep78 foreign
2 -> tabulation of rep78
3
4      rep78 |      Freq.      Percent      Cum.
5 -----+-----
6          1 |          2         2.90         2.90
7          2 |          8        11.59        14.49
8          3 |         30        43.48        57.97
9          4 |         18        26.09        84.06
10          5 |         11        15.94       100.00
11 -----+-----

```

```

12      Total |          69      100.00
13
14 -> tabulation of foreign
15
16      foreign |          Freq.      Percent      Cum.
17 -----+-----
18          0 |          52      70.27      70.27
19          1 |          22      29.73      100.00
20 -----+-----
21      Total |          74      100.00

```

We can use the `plot` option to make a plot to visually show the tabulated values.

```

1 tabulate rep78, plot
2      rep78 |          Freq.
3 -----+-----
4          1 |          2 |**
5          2 |          8 |*****
6          3 |         30 |*****
7          4 |         18 |*****
8          5 |         11 |*****
9 -----+-----
10      Total |          69

```

We can also make crosstabs using `tabulate`. Let's look at the repair history broken down by *foreign* and *domestic* cars.

```

1 tabulate rep78 foreign
2      |          foreign
3      rep78 |          0          1 |          Total
4 -----+-----
5          1 |          2          0 |          2
6          2 |          8          0 |          8
7          3 |         27          3 |         30
8          4 |          9          9 |         18
9          5 |          2          9 |         11
10 -----+-----
11      Total |         48         21 |         69

```

With the `column` option, we can request column percentages. Notice that about 86% of the foreign cars received a rating of 4 or 5. Only about 23% of domestic cars were rated that highly.

```

1 tabulate rep78 foreign, column
2      |          foreign
3      rep78 |          0          1 |          Total
4 -----+-----
5          1 |          2          0 |          2
6          |         4.17         0.00 |         2.90
7 -----+-----
8          2 |          8          0 |          8
9          |        16.67         0.00 |        11.59
10 -----+-----
11          3 |         27          3 |         30
12          |        56.25        14.29 |         43.48
13 -----+-----
14          4 |          9          9 |         18
15          |        18.75        42.86 |        26.09
16 -----+-----

```

17	5		2	9		11
18			4.17	42.86		15.94
19	-----+-----+-----					
20	Total		48	21		69
21			100.00	100.00		100.00

We can use the `nofreq` option to suppress the frequencies, and just focus on the percentages.

```
1 tabulate rep78 foreign, column nofreq
```

2			foreign		
3	rep78		0	1	Total
4	-----+-----+-----				
5	1		4.17	0.00	2.90
6	2		16.67	0.00	11.59
7	3		56.25	14.29	43.48
8	4		18.75	42.86	26.09
9	5		4.17	42.86	15.94
10	-----+-----+-----				
11	Total		100.00	100.00	100.00

Note that the order of the options does not matter. Just remember that the options must come after the comma.

```
1 tabulate rep78 foreign, nofreq column
```

2			foreign		
3	rep78		0	1	Total
4	-----+-----+-----				
5	1		4.17	0.00	2.90
6	2		16.67	0.00	11.59
7	3		56.25	14.29	43.48
8	4		18.75	42.86	26.09
9	5		4.17	42.86	15.94
10	-----+-----+-----				
11	Total		100.00	100.00	100.00

1.2.3 Generating summary statistics with summarize

For summary statistics, we can use the `summarize` command. Let's generate some summary statistics on *mpg*.

```
1 summarize mpg
```

2	Variable		Obs	Mean	Std. Dev.	Min	Max
3	-----+-----						
4	mpg		74	21.2973	5.785503	12	41

We can use the `detail` option of the `summarize` command to get more detailed summary statistics.

1	summarize mpg, detail					
2	mpg					
3	-----					
4	Percentiles		Smallest			
5	1%	12	12			
6	5%	14	12			
7	10%	14	14	Obs	74	
8	25%	18	14	Sum of Wgt.	74	
9						
10	50%	20			Mean	21.2973
11			Largest		Std. Dev.	5.785503
12	75%	25	34			
13	90%	29	35	Variance	33.47205	

14	95%	34	35	Skewness	.9487176
15	99%	41	41	Kurtosis	3.975005

To get these values separately for *foreign* and *domestic*, we could use the `by foreign:` prefix as shown below. Note that we first had to `sort` the data before using `by foreign:`.

```

1 sort foreign
2 by foreign: summarize mpg
3 -> foreign= 0
4 Variable |      Obs      Mean   Std. Dev.      Min      Max
5 -----+-----
6      mpg |      52   19.82692   4.743297      12      34
7
8 -> foreign= 1
9 Variable |      Obs      Mean   Std. Dev.      Min      Max
10 -----+-----
11      mpg |      22   24.77273   6.611187      14      41

```

This is not the most efficient way to do this. Another way, which does not require the data to be sorted, is by using the `summarize()` option as part of the `tabulate` command.

```

1 tabulate foreign, summarize(mpg)
2      |      Summary of mpg
3 foreign |      Mean   Std. Dev.      Freq.
4 -----+-----
5      0 |   19.826923   4.7432972      52
6      1 |   24.772727   6.6111869      22
7 -----+-----
8      Total |   21.297297   5.7855032      74

```

Here is another example, showing the average price of cars for each level of repair history.

```

1 tabulate rep78, summarize(price)
2      |      Summary of price
3 rep78 |      Mean   Std. Dev.      Freq.
4 -----+-----
5      1 |   4564.5    522.55191      2
6      2 |   5967.625   3579.3568      8
7      3 |   6429.2333   3525.1398     30
8      4 |   6071.5    1709.6083     18
9      5 |    5913     2615.7628     11
10 -----+-----
11      Total |   6146.0435   2912.4403     69

```

1.2.4 Summary

- `describe`: provide information about the current data file, including the number of variables and observations and a listing of the variables in a data file.
- `codebook`: produce codebook like information for the current data file.
- `inspect`: provide a quick overview of data file.
- `list` make mpg: list out the variables make and mpg.
- `tabulate` mpg: make a table of mpg.
- `tabulate` rep78 foreign: make a two way table of rep78 by foreign.
- `summarize` mpg price: produce summary statistics of mpg and price.
- To produce summary statistics for mpg separately for foreign and domestic cars, use

```

1 sort foreign
2 by foreign: summarize(mpg)

```

- `tabulate foreign, summarize(mpg)`: produce summary statistics for mpg by foreign (prior sorting not required).

1.3 Getting help using Stata

This module shows resources you can use to help you learn and use Stata.

1.3.1 Stata online help

When you know the name of the command you want to use (e.g., `summarize`), you can use the Stata help to get a quick summary of the command and its syntax. You can do this in two ways:

1. type `help summarize` in the command window, or
2. click **Help, Stata Command**, then type `summarize`.

Here is what help `summarize` looks like.

```

1 help summarize
2 help summarize                                dialog: summarize
3 -----
4
5 Title
6
7 [R] summarize -- Summary statistics
8
9
10 Syntax
11
12 summarize [varlist] [if] [in] [weight] [, options]
13
14 options      description
15 -----
16 Main
17 detail       display additional statistics
18 meanonly     suppress the display; only calculate the
19              mean; programmer's option
20 format       use variable's display format
21 separator(#) draw separator line after every # variables;
22              default is separator(5)
23 -----
24 varlist may contain time-series operators; see tsvarlist.
25 by may be used with summarize; see by.
26 aweights, fweights, and iweights are allowed. However,
27 iweights may not be used with the detail option; see weight.
28 //(omitted)

```

If you use the pull-down menu to get help for a command, it shows the same basic information but related commands and topics are hotlinks you can click.

When you want to search for a keyword, e.g. `memory`, you can use Stata to search for help topics that contain that keyword. You can do this in two ways:

1. Type search `memory` in the command window, or
2. Click **Help, Search**, then `memory`.

Here is what search memory looks like.

```
1 search memory
2
3 GS      . . . . . Getting Started manual
4
5 [U]     Chapter 7 . . . . . Setting the size of memory
6         (help memory)
7
8 [R]     compress . . . . . Compress data in memory
9         (help compress)
10 //(omitted)
```

As you can see, there are lots of help topics that refer to memory. Some of the topics give you a command, and then you can get help for that command. Notice that those topics start with **GS**[U] or **[R]**. Those are indicating which Stata manual you could find the command (GS=Getting Started, U=Users Guide, R=Reference Guide).

The next set of topics all start with **FAQ** because these are Frequently Asked Questions from the Stata web site. You can see the title of the FAQ and the address of the FAQ. Lastly, there is a topic that starts with **STB** which stands for Stata Technical Bulletin. These refer to add-on programs that you can install into Stata. There are dozens, if not hundreds of specialized and useful programs that you can get from the Stata Technical Bulletin.

You can access this same kind of help from the pull-down menus by clicking **Help** then **Search** then type `memory`. Note how the related commands, the FAQs, and the STB all have hotlinks you can click. For example, you can click on a FAQ and it will bring up that FAQ in your web browser. Or, you could click on an STB and it would walk you through the steps of installing that STB into your copy of Stata. As you can see, there are real advantages to using the pull-down menus for getting help because it is so easy to click on the related topics.

1.3.2 Stata sample data files

Stata has some very useful data files available to you for learning and practicing Stata. For example, you can type

```
1 sysuse auto
```

to use the auto data file that comes with Stata. You can type

```
1 sysuse dir
```

to see the entire list of data files that ship with Stata. You can type

```
1 help dta_contents
```

to see all of the sample data files that you can easily access from within Stata.

1.3.3 Stata web pages

The Stata web page is a wonderful resource. You can visit the main page at <http://www.stata.com>.

The User Support page (click **User Support** from main page) has a great set of resources, including

- FAQs
- NetCourses
- StataList: How to subscribe
- StataList: Archives
- Statalist ado-file Archives
- Stata Bookstore

In the bookstore, you can find books on Stata. A good intro book on Stata is **Statistics with Stata**.

2 Fundamentals of Using Stata (part II)

2.1 Using IF with Stata commands

This module shows the use of **if** with common Stata commands.

Let's use the auto data file.

```
1 sysuse auto
```

For this module, we will focus on the variables *make*, *rep78*, *foreign*, *mpg*, and *price*. We can use the **keep** command to keep just these five variables.

```
1 keep make rep78 foreign mpg price
```

Let's make a table of *rep78* by *foreign* to look at the repair histories of the foreign and domestic cars.

```
1 tabulate rep78 foreign
2      |      foreign
3      rep78 |      0      1 |      Total
4 -----+-----+-----
5      1 |      2      0 |      2
6      2 |      8      0 |      8
7      3 |     27      3 |     30
8      4 |      9      9 |     18
9      5 |      2      9 |     11
10 -----+-----+-----
11      Total |     48     21 |     69
```

Suppose we wanted to focus on just the cars with repair histories of four or better. We can use **if** suffix to do this.

```
1 tabulate rep78 foreign if rep78 >=4
2      |      foreign
3      rep78 |      0      1 |      Total
4 -----+-----+-----
5      4 |      9      9 |     18
6      5 |      2      9 |     11
7 -----+-----+-----
8      Total |     11     18 |     29
```

Let's make the above table using the **column** and **nofreq** options. The command **column** requests column percentages while the command **nofreq** suppresses cell frequencies. Note that **column** and **nofreq** come after the comma. These are options on the **tabulate** command and options need to be placed after a comma.

```
1 tabulate rep78 foreign if rep78 >= 4, column nofreq
2      |      foreign
3      rep78 |      0      1 |      Total
4 -----+-----+-----
5      4 |    81.82    50.00 |    62.07
6      5 |    18.18    50.00 |    37.93
7 -----+-----+-----
8      Total |   100.00   100.00 |   100.00
```

The use of **if** is not limited to the **tabulate** command. Here, we use it with the **list** command.

```
1 list if rep78 >= 4
2      make      price      mpg      rep78      foreign
3  3.      AMC Spirit      3799      22      .      0
```

```

4   5.      Buick Electra      7827      15      4      0
5   7.      Buick Opel       4453      26      .      0
6  15.     Chev. Impala      5705      16      4      0
7  20.     Dodge Colt       3984      30      5      0
8  24.     Ford Fiesta       4389      28      4      0
9  29.     Merc. Bobcat      3829      22      4      0
10 30.     Merc. Cougar      5379      14      4      0
11 //(omitted)

```

Did you see that some of the observations had a value of '.' for rep78? These are missing values. For example, the value of *rep78* for the AMC Spirit is missing. **Stata treats a missing value as positive infinity**, the highest number possible. So, when we said `list if rep78 >= 4`, Stata included the observations where *rep78* was '.' as well.

If we wanted to include just the valid (non-missing) observations that are greater than or equal to 4, we can do the following to tell Stata we want only observations where *rep78* >= 4 and *rep78* is not missing.

```

1 list if rep78 >= 4 & !missing(rep78)
2           make      price      mpg      rep78      foreign
3   5.      Buick Electra      7827      15      4      0
4  15.     Chev. Impala      5705      16      4      0
5  20.     Dodge Colt       3984      30      5      0
6  24.     Ford Fiesta       4389      28      4      0
7  29.     Merc. Bobcat      3829      22      4      0
8  30.     Merc. Cougar      5379      14      4      0
9  33.     Merc. XR-7        6303      14      4      0
10 35.     Olds 98           8814      21      4      0
11 //(omitted)

```

This code will also yield the same output as above.

```

1 list if rep78 >= 4 & rep78 != .

```

We can use `if` with most Stata commands. Here, we get summary statistics for *price* for cars with repair histories of 1 or 2. Note the double equal (==) represents **IS EQUAL TO** and the pipe (|) represents **OR**.

```

1 summarize price if rep78 == 1 | rep78 == 2
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      10      5687   3216.375     3667    14500

```

A simpler way to say this would be ...

```

1 summarize price if rep78 <= 2
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      10      5687   3216.375     3667    14500

```

Likewise, we can do this for cars with repair history of 3, 4 or 5.

```

1 summarize price if rep78 == 3 | rep78 == 4 | rep78 == 5
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      59     6223.847   2880.454     3291    15906

```

Additionally, we can use this code to designate a range of values. Here is a summary of *price* for the values 3 through 5 in *rep78*.

```

1 summarize price if inrange(rep78,3,5)
2 Variable |      Obs      Mean   Std. Dev.      Min      Max

```

```

3 -----+-----
4 price |          59    6223.847    2880.454    3291    15906

```

Let's simplify this by saying `rep78 >= 3`.

```

1 summarize price if rep78 >= 3
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4 price |      64    6239.984    2925.843     3291    15906

```

Did you see the mistake we made? We accidentally included the missing values because we forgot to exclude them. We really needed to say.

```

1 summarize price if rep78 >= 3 & !missing(rep78)
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4 price |      59    6223.847    2880.454     3291    15906

```

2.1.1 Taking a random sample

It is also possible to take a simple random sample of your data using the `sample` command. This information can be found on our STATA FAQ page: [How can I draw a random sample of my data?](#)

2.1.2 Summary

Most Stata commands can be followed by `if`, for example

```

1 summarize if rep78 == 2
2 summarize if rep78 >= 2
3 summarize if rep78 > 2
4 summarize if rep78 <= 2
5 summarize if rep78 < 2
6 summarize if rep78 != 2

```

`if` expressions can be connected with `|` for **OR**, `&` for **AND**.

2.1.3 Missing Values

Missing values are represented as `.` and are the highest value possible. Therefore, when values are missing, be careful with commands like

```

1 summarize if rep78 > 3
2 summarize if rep78 >= 3
3 summarize if rep78 != 3

```

to omit missing values, use

```

1 summarize if rep78 > 3 & !missing(rep78)
2 summarize if rep78 >= 3 & !missing(rep78)
3 summarize if rep78 != 3 & !missing(rep78)

```

2.2 A statistical sampler in Stata

Version info: Code for this page was tested in Stata 12.

This module will give a brief overview of some common statistical tests in Stata. Let's use the `auto` data file that we will use for our examples.

`auto`

```
1 sysuse auto
```

2.2.1 t-tests

Let's do a t-test comparing the miles per gallon (*mpg*) of foreign and domestic cars.

```
1 ttest mpg , by(foreign)
2 Two-sample t test with equal variances
3
4 -----
5      Group |      Obs      Mean      Std. Err.      Std. Dev.      [95% Conf. Interval]
6 -----+-----
7           0 |      52    19.82692     .657777     4.743297     18.50638     21.14747
8           1 |      22    24.77273     1.40951     6.611187     21.84149     27.70396
9 -----+-----
10    combined |      74    21.2973     .672551     5.785503     19.9569     22.63769
11 -----+-----
12      diff |           -4.945804     1.362162           -7.661225     -2.230384
13 -----
14 Degrees of freedom: 72
15
16                      Ho: mean(0) - mean(1) = diff = 0
17
18      Ha: diff < 0 Ha: diff ~= "0" Ha: diff > 0
19      t =   -3.6308           t =   -3.6308           t =   -3.6308
20      P < t =   0.0003           P > |t| =   0.0005           P > t =   0.9997
```

As you see in the output above, the domestic cars had significantly lower *mpg* (19.8) than the foreign cars (24.7).

2.2.2 Chi-square

Let's compare the repair rating (*rep78*) of the foreign and domestic cars. We can make a crosstab of *rep78* by *foreign*. We may want to ask whether these variables are independent. We can use the *chi2* option to request a chi-square test of independence as well as the crosstab.

```
1 tabulate rep78 foreign, chi2
2
3      rep78 |      foreign
4      rep78 |      0      1 |      Total
5 -----+-----
6           1 |      2      0 |      2
7           2 |      8      0 |      8
8           3 |     27      3 |     30
9           4 |      9      9 |     18
10          5 |      2      9 |     11
11 -----+-----
12      Total |     48     21 |     69
13
14      Pearson chi2(4) = 27.2640 Pr = 0.000
```

The chi-square is not really valid when you have empty cells. In such cases when you have empty cells, or cells with small frequencies, you can request Fisher's exact test with the *exact* option.

```
1 tabulate rep78 foreign, chi2 exact
2
3      rep78 |      foreign
4      rep78 |      0      1 |      Total
5 -----+-----
```

5	1		2	0		2
6	2		8	0		8
7	3		27	3		30
8	4		9	9		18
9	5		2	9		11
10	-----+-----+-----					
11	Total		48	21		69
12						
13	Pearson chi2 (4) = 27.2640 Pr = 0.000					
14	Fisher's exact = 0.000					

2.2.3 Correlation

We can use the `correlate` command to get the correlations among variables. Let's look at the correlations among *price* *mpg* *weight* and *rep78*. (We use *rep78* in the correlation even though it is not continuous to illustrate what happens when you use correlate with variables with missing data.)

```
1 correlate price mpg weight rep78
2 (obs=69)
```

3			price	mpg	weight	rep78
4						
5						
6	price		1.0000			
7	mpg		-0.4559	1.0000		
8	weight		0.5478	-0.8055	1.0000	
9	rep78		0.0066	0.4023	-0.4003	1.0000

Note that the output above said (obs=69). The `correlate` command drops data on a listwise basis, meaning that if any of the variables are missing, then the entire observation is omitted from the correlation analysis.

We can use `pwcorr` (pairwise correlations) if we want to obtain correlations that deletes missing data on a pairwise basis instead of a listwise basis. We will use the `obs` option to show the number of observations used for calculating each correlation.

```
1 pwcorr price mpg weight rep78, obs
```

2			price	mpg	weight	rep78
3						
4	price		1.0000			
5			74			
6						
7	mpg		-0.4686	1.0000		
8			74	74		
9						
10	weight		0.5386	-0.8072	1.0000	
11			74	74	74	
12						
13	rep78		0.0066	0.4023	-0.4003	1.0000
14			69	69	69	69
15						

Note how the correlations that involve *rep78* have an N of 69 compared to the other correlations that have an N of 74. This is because *rep78* has five missing values, so it only had 69 valid observations, but the other variables had no missing data so they had 74 valid observations.

2.2.4 Regression

Let's look at doing regression analysis in Stata. For this example, let's drop the cases where *rep78* is 1 or 2 or missing.

```
1 drop if (rep78 <= 2) | (rep78 ==.)
2 (15 observations deleted)
```

Now, let's predict *mpg* from *price* and *weight*. As you see below, *weight* is a significant predictor of *mpg*, but *price* is not.

```
1 regress mpg price weight
2
3 Source |      SS      df      MS                Number of obs =      59
4 -----+-----
5 Model | 1375.62097      2  687.810483            F( 2, 56) =    47.87
6 Residual | 804.616322     56  14.3681486            Prob > F      =    0.0000
7 -----+-----
8 Total | 2180.23729     58  37.5902981            R-squared     =    0.6310
9                                     Adj R-squared =    0.6178
10                                    Root MSE     =    3.7905
11
12 -----+-----
13 mpg |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
14 -----+-----
15 price | -.0000139   .0002108     -0.066   0.948    - .0004362   .0004084
16 weight | -.005828   .0007301     -7.982   0.000    - .0072906  -.0043654
17 _cons | 39.08279   1.855011     21.069   0.000     35.36676   42.79882
```

What if we wanted to predict *mpg* from *rep78* as well. *rep78* is really more of a categorical variable than it is a continuous variable. To include it in the regression, we should convert *rep78* into dummy variables. Fortunately, Stata makes dummy variables easily using *tabulate*. The *gen(rep)* option tells Stata that we want to generate dummy variables from *rep78* and we want the stem of the dummy variables to be *rep*.

```
1 tabulate rep78, gen(rep)
2
3 rep78 |      Freq.      Percent      Cum.
4 -----+-----
5 3 |      30      50.85      50.85
6 4 |      18      30.51      81.36
7 5 |      11      18.64     100.00
8 -----+-----
9 Total |      59     100.00
```

Stata has created *rep1* (1 if *rep78* is 3), *rep2* (1 if *rep78* is 4) and *rep3* (1 if *rep78* is 5). We can use the *tabulate* command to verify that the dummy variables were created properly.

```
1 tabulate rep78 rep1
2
3 rep78 |      rep78==      3.0000
4 -----+-----
5 3 |      0      30 |      30
6 4 |      18      0 |      18
7 5 |      11      0 |      11
8 -----+-----
9 Total |      29      30 |      59
10 tabulate rep78 rep2
11
12 rep78 |      rep78==      4.0000
13 -----+-----
14 3 |      0      30 |      30
15 4 |      18      0 |      18
16 5 |      11      0 |      11
17 -----+-----
18 Total |      29      30 |      59
```

```

14      3 |      30      0 |      30
15      4 |      0      18 |      18
16      5 |      11      0 |      11
17 -----+-----+-----
18      Total |      41      18 |      59
19 tabulate rep78 rep3
20      | rep78==      5.0000
21      rep78 |      0      1 |      Total
22 -----+-----+-----
23      3 |      30      0 |      30
24      4 |      18      0 |      18
25      5 |      0      11 |      11
26 -----+-----+-----
27      Total |      48      11 |      59

```

Now we can include *rep1* and *rep2* as dummy variables in the regression model.

```

1 regress mpg price weight rep1 rep2
2
3      Source |      SS      df      MS      Number of obs =      59
4 -----+-----+-----+-----+-----
5      Model | 1435.91975      4 358.979938      F( 4, 54) = 26.04
6      Residual | 744.317536     54 13.7836581      Prob > F      = 0.0000
7 -----+-----+-----+-----
8      Total | 2180.23729     58 37.5902981      R-squared      = 0.6586
9                                     Adj R-squared   = 0.6333
10                                    Root MSE       = 3.7126
11
12      mpg |      Coef.   Std. Err.    t    P>|t|    [95% Conf. Interval]
13 -----+-----+-----+-----+-----
14      price | -.0001126   .0002133    -0.53  0.600   - .0005403   .0003151
15      weight | -.005107    .0008236    -6.20  0.000   - .0067584   -.0034557
16      rep1 | -2.886288   1.504639    -1.92  0.060   -5.902908    .1303314
17      rep2 | -2.88417    1.484817    -1.94  0.057   -5.861048    .0927086
18      _cons | 39.89189    1.892188    21.08  0.000    36.09828    43.68555

```

2.2.5 Analysis of variance

If you wanted to do an analysis of variance looking at the differences in *mpg* among the three repair groups, you can use the *oneway* command to do this.

```

1 oneway mpg rep78
2
3      Source      Analysis of Variance
4      Source      SS      df      MS      F      Prob > F
5 -----+-----+-----+-----+-----
6 Between groups    506.325167      2 253.162583      8.47      0.0006
7 Within groups    1673.91212     56 29.8912879
8 -----+-----+-----+-----
9 Total            2180.23729     58 37.5902981
10
11 Bartlett's test for equal variances:  chi2(2) = 9.9384  Prob>chi2 = 0.007

```

If you include the *tabulate* option, you get mean *mpg* for the three groups, which shows that the group with the best repair rating (*rep78* of 5) also has the highest *mpg* (27.3).

```

1 oneway mpg rep78, tabulate
2

```

3			Summary of mpg		
4	rep78		Mean	Std. Dev.	Freq.
5	-----+-----				
6	3		19.433333	4.1413252	30
7	4		21.666667	4.9348699	18
8	5		27.363636	8.7323849	11
9	-----+-----				
10	Total		21.59322	6.1310927	59
11					
12	Analysis of Variance				
13	Source		SS	df	MS
14	-----+-----				
15	Between groups		506.325167	2	253.162583
16	Within groups		1673.91212	56	29.8912879
17	-----+-----				
18	Total		2180.23729	58	37.5902981
19					
20	Bartlett's test for equal variances: $\chi^2(2) = 9.9384$ Prob> $\chi^2 = 0.007$				

If you want to include covariates, you need to use the `anova` command. The continuous(price weight) option tells Stata that those variables are covariates.

1 anova mpg rep78 c.price c.weight

2

3

4

5

6

7

8

9

10

11

12

13

14

15

Number of obs = 59

R-squared = 0.6586

Root MSE = 3.71263

Adj R-squared = 0.6333

Source | Partial SS df MS F Prob > F

-----+-----

Model | 1435.91975 4 358.979938 26.04 0.0000

|

rep78 | 60.2987853 2 30.1493926 2.19 0.1221

price | 3.8421233 1 3.8421233 0.28 0.5997

weight | 529.932889 1 529.932889 38.45 0.0000

|

Residual | 744.317536 54 13.7836581

-----+-----

Total | 2180.23729 58 37.5902981

2.3 An overview of Stata syntax

This module shows the general structure of Stata commands. We will demonstrate this using `summarize` as an example, although this general structure applies to most Stata commands.

Note: This code was tested in Stata 12.

Let's first use the auto data file.

```
1 use auto
```

As you have seen, we can type `summarize` and it will give us summary statistics for all of the variables in the data file.

1	<code>summarize</code>					
2	Variable		Obs	Mean	Std. Dev.	Min
3	-----+-----					
4	make		0			
5	price		74	6165.257	2949.496	3291
6	mpg		74	21.2973	5.785503	12

7	rep78		69	3.405797	.9899323	1	5
8	hdroom		74	2.993243	.8459948	1.5	5
9	trunk		74	13.75676	4.277404	5	23
10	weight		74	3019.459	777.1936	1760	4840
11	length		74	187.9324	22.26634	142	233
12	turn		74	39.64865	4.399354	31	51
13	displ		74	197.2973	91.83722	79	425
14	gratio		74	3.014865	.4562871	2.19	3.89
15	foreign		74	.2972973	.4601885	0	1

It is also possible to obtain means for specific variables. For example, below we get summary statistics just for *mpg* and *price*.

```
1 summarize mpg price
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	74	21.2973	5.785503	12	41
price	74	6165.257	2949.496	3291	15906

We could further tell Stata to limit the summary statistics to just foreign cars by adding an if qualifier.

```
1 summarize mpg price if (foreign == 1)
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	22	24.77273	6.611187	14	41
price	22	6384.682	2621.915	3748	12990

The *if* qualifier can contain more than one condition. Here, we ask for summary statistics for the foreign cars which get less than 30 miles per gallon.

```
1 summarize mpg price if foreign == 1 & mpg < 30
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	17	21.94118	3.896643	14	28
price	17	6996.235	2674.552	3895	12990

We can use the **detail** option to ask Stata to give us more detail in the summary statistics. Notice that the **detail** option goes after the comma. If the comma were omitted, Stata would give an error.

```
1 summarize mpg price if foreign == 1 & mpg < 30 , detail
```

mpg					

Percentiles					
Smallest					
1%	14	14			
5%	14	17			
10%	17	17	Obs		17
25%	18	18	Sum of Wgt.		17
50%					
	23		Mean		21.94118
Largest					
			Std. Dev.		3.896643
75%	25	25			
90%	26	25	Variance		15.18382
95%	28	26	Skewness		-.4901235
99%	28	28	Kurtosis		2.201759
price					

Percentiles					
Smallest					

20	1%	3895	3895		
21	5%	3895	4296		
22	10%	4296	4499	Obs	17
23	25%	5079	4697	Sum of Wgt.	17
24					
25	50%	6229		Mean	6996.235
26			Largest	Std. Dev.	2674.552
27	75%	8129	9690		
28	90%	11995	9735	Variance	7153229
29	95%	12990	11995	Skewness	.9818272
30	99%	12990	12990	Kurtosis	2.930843

Note that even though we built these parts up one at a time, they don't have to go together. Let's look at some other forms of the summarize command.

You can tell Stata which observation numbers you want using the in qualifier. Here we ask for summaries of observations 1 to 10. This is useful if you have a big data file and want to try out a command on a subset of observations.

```
1 summarize in 1/10
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
make	0				
price	10	5517.4	2063.518	3799	10372
mpg	10	19.5	3.27448	15	26
rep78	8	3.125	.3535534	3	4
hdroom	10	3.3	.7527727	2	4.5
trunk	10	14.7	3.88873	10	21
weight	10	3271	558.3796	2230	4080
length	10	194	19.32759	168	222
turn	10	40.2	3.259175	34	43
displ	10	223.9	71.77503	121	350
gratio	10	2.907	.3225264	2.41	3.58
foreign	10	0	0	0	0

Also, recall that you can ask Stata to perform summaries for foreign and domestic cars separately using by, as shown below.

```
1 sort foreign
2 by foreign: summarize
3 -> foreign= 0
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
make	0				
price	52	6072.423	3097.104	3291	15906
mpg	52	19.82692	4.743297	12	34
rep78	48	3.020833	.837666	1	5
hdroom	52	3.153846	.9157578	1.5	5
trunk	52	14.75	4.306288	7	23
weight	52	3317.115	695.3637	1800	4840
length	52	196.1346	20.04605	147	233
turn	52	41.44231	3.967582	31	51
displ	52	233.7115	85.26299	86	425
gratio	52	2.806538	.3359556	2.19	3.58
foreign	52	0	0	0	0

```
18
19 -> foreign= 1
```

Variable	Obs	Mean	Std. Dev.	Min	Max
make	0				
price	22	6384.682	2621.915	3748	12990
mpg	22	24.77273	6.611187	14	41
rep78	21	4.285714	.7171372	3	5
hdroom	22	2.613636	.4862837	1.5	3.5
trunk	22	11.40909	3.216906	5	16
weight	22	2315.909	433.0035	1760	3420
length	22	168.5455	13.68255	142	193
turn	22	35.40909	1.501082	32	38
displ	22	111.2273	24.88054	79	163
gratio	22	3.507273	.2969076	2.98	3.89
foreign	22	1	0	1	1

Let's review all those pieces.

A command can be preceded with a **by** prefix, as shown below.

```
1 by foreign: summarize
```

There are many parts that can come after a command. They are each presented separately below.

For example, **summarize** followed by the names of variables.

```
1 summarize mpg price
```

summarize with **in** specifying a range of records to be summarized.

```
1 summarize in 1/10
```

summarize with simple **if** specifying records to summarize.

```
1 summarize if foreign == 1
```

summarize with complex **if** specifying records to summarize.

```
1 summarize if foreign == 1 & mpg > 30
```

summarize followed by option(s).

```
1 summarize , detail
```

So, putting it all together, the general syntax of the **summarize** command can be described as:

```
1 [by varlist:] summarize [varlist] [in range] [if exp] , [options]
```

Understanding the overall syntax of Stata commands helps you remember them and use them more effectively, and it also aids you understand the help files in Stata. All the extra stuff about **by**, **if** and **in** could be confusing. Let's have a look at the help file for **summarize**. It makes more sense knowing what the **by**, **if** and **in** parts mean.

```
1 help summarize
2 -----
3 help for summarize (manual: [R] summarize)
4 -----
5
6 Summary statistics
7 -----
8
9 [by varlist:] summarize [varlist] [weight] [if exp] [in range]
10 [, { detail | meanonly } format ]
```

2.4 Missing data

2.4.1 Introduction

This module will explore missing data in STATA, focusing on numeric missing data. It will describe how to indicate missing data in your raw data files, as well as how missing data are handled in STATA logical commands and assignment statements.

We will illustrate some of the missing data properties in STATA using data from a reaction time study with eight subjects indicated by the variable *id*, and the subjects reaction times were measured at three time points (*trial1* *trial2* *trial3*). The input data file is shown below.

```
1 input id trial1 trial2 trial3
2 1 1.5 1.4 1.6
3 2 1.5 . 1.9
4 3 . 2.0 1.6
5 4 . . 2.2
6 5 1.9 2.1 2
7 6 1.8 2.0 1.9
8 7 . . .
9 end
10 list
```

You might notice that some of the reaction times are coded using a single '.' as is the case for subject 2. The person measuring time for that trial did not measure the response time properly, therefore the data for the second trial is missing.

```
1      +-----+
2      | id   trial1   trial2   trial3 |
3      |-----|
4  1. | 1     1.5     1.4     1.6 |
5  2. | 2     1.5     .     1.9 |
6  3. | 3     .     2     1.6 |
7  4. | 4     .     .     2.2 |
8  5. | 5     1.9     2.1     2 |
9      |-----|
10 6. | 6     1.8     2     1.9 |
11 7. | 7     .     .     . |
12      +-----+
```

2.4.2 How STATA handles missing data in STATA procedures

As a general rule, STATA commands that perform computations of any type handle missing data by omitting the missing values. However, the way that missing values are omitted is not always consistent across commands, so let's take a look at some examples.

First, let's `summarize` our reaction time variables and see how STATA handles the missing values.

```
1 summarize trial1 trial2 trial3
```

As you see in the output below, `summarize` computed means using 4 observations for *trial1* and *trial2* and 6 observations for *trial3*. In short, the `summarize` command performed the computations on all the available data.

Variable	Obs	Mean	Std. Dev.	Min	Max
trial1	4	1.675	.2061553	1.5	1.9
trial2	4	1.875	.3201562	1.4	2.1
trial3	6	1.866667	.233809	1.6	2.2

A second example, shows how the `tabulation` or `tab1` command handles missing data. Like `summarize`, `tab1` uses just available data. Note that the percentages are computed based on the total number of non-missing cases.

```

1 tab1 trial1 trial2 trial3
2 -> tabulation of trial1
3
4      trial1 |      Freq.      Percent      Cum.
5 -----+-----
6          1.5 |          2       50.00       50.00
7          1.8 |          1       25.00       75.00
8          1.9 |          1       25.00      100.00
9 -----+-----
10         Total |          4      100.00
11
12 -> tabulation of trial2
13
14      trial2 |      Freq.      Percent      Cum.
15 -----+-----
16          1.4 |          1       25.00       25.00
17          2   |          2       50.00       75.00
18          2.1 |          1       25.00      100.00
19 -----+-----
20         Total |          4      100.00
21
22 -> tabulation of trial3
23
24      trial3 |      Freq.      Percent      Cum.
25 -----+-----
26          1.6 |          2       33.33       33.33
27          1.9 |          2       33.33       66.67
28          2   |          1       16.67       83.33
29          2.2 |          1       16.67      100.00
30 -----+-----
31         Total |          6      100.00

```

It is possible that you might want the percentages to be computed out of the total number of observations, and the percentage missing for each variable shown in the table. This can be achieved by including the `missing` option after the `tabulation` command,

```

1 tab1 trial1 trial2 trial3, m
2 -> tabulation of trial1
3
4      trial1 |      Freq.      Percent      Cum.
5 -----+-----
6          1.5 |          2       28.57       28.57
7          1.8 |          1       14.29       42.86
8          1.9 |          1       14.29       57.14
9          .   |          3       42.86      100.00
10 -----+-----
11         Total |          7      100.00
12
13 -> tabulation of trial2
14
15      trial2 |      Freq.      Percent      Cum.
16 -----+-----
17          1.4 |          1       14.29       14.29
18          2   |          2       28.57       42.86

```

```

19      2.1 |      1      14.29      57.14
20      .  |      3      42.86     100.00
21 -----+-----
22      Total |      7     100.00
23
24 -> tabulation of trial3
25
26      trial3 |      Freq.      Percent      Cum.
27 -----+-----
28      1.6  |      2      28.57      28.57
29      1.9  |      2      28.57      57.14
30      2    |      1      14.29      71.43
31      2.2  |      1      14.29      85.71
32      .    |      1      14.29     100.00
33 -----+-----
34      Total |      7     100.00

```

Let's look at how the `correlate` command handles missing data. We would expect that it would perform the computations based on the available data, and omit the missing values. Here is an example command.

```
1 corr trial1 trial2 trial3
```

The output is show below. Note how the missing values were excluded. For each pair variables, the `corr` command used the number of pairs that had valid data. For the pair formed by *trial1* and *trial2*, there were 3 pairs with valid data. For the pairing of *trial1* and *trial3* there were 4 valid pairs, and likewise there were 4 valid pairs for *trial3* and *trial2*. Using all of the valid pairs of data is called pairwise deletion of missing data.

```

1      | trial1 trial2 trial3
2 -----+-----
3      trial1 | 1.0000
4          |      4
5          |
6      trial2 | 0.9939 1.0000
7          |      3      4
8          |
9      trial3 | 0.7001 0.6439 1.0000
10         |      4      4      6

```

It is possible to ask STATA to only perform the correlations on the observations that had complete data for all of the variables on the var statement. For example, you might want the correlations of the reaction times just for the observations that had non-missing data on all of the trials. This is called `listwise` deletion of missing data meaning that when any of the variables are missing, the entire observation is omitted from the analysis. You can request `listwise` deletion within `pwcorr` as illustrated below.

```

1 pwcorr trial1 trial2 trial3, listwise obs
2      | trial1 trial2 trial3
3 -----+-----
4      trial1 | 1.0000
5          |      3
6          |
7      trial2 | 0.9939 1.0000
8          |      3      3
9          |
10     trial3 | 1.0000 0.9939 1.0000
11         |      3      3      3

```

2.4.3 Summary of how missing values are handled in STATA procedures

- **summarize**: For each variable, the number of non-missing values are used.
- **tabulation**: By default, missing values are excluded and percentages are based on the number of non-missing values. If you use the missing option on the **tab** command, the percentages are based on the total number of observations (non-missing and missing) and the percentage of missing values are reported in the table.
- **corr**: By default, correlations are computed based on the number of pairs with non-missing data (**pairwise** deletion of missing data). The **pwcorr** command can be used to request that correlations be computed only for observations that have non-missing data for all variables listed after the **pwcorr** command (**listwise** deletion of missing data).
- **reg**: If any of the variables listed after the **reg** command are missing, the observations missing that value(s) are excluded from the analysis (i.e., **listwise** deletion of missing data).
- For other procedures, see the STATA manual for information on how missing data are handled.

2.4.4 Missing values in assignment statements

It is important to understand how missing values are handled in assignment statements. Consider the example shown below.

```
1 gen sum1 = trial1 + trial2 + trial3
```

The **list** command below illustrates how missing values are handled in assignment statements. The variable *sum1* is based on the variables *trial1*, *trial2* and *trial3*. If any of those variables were missing, the value for *sum1* was set to missing. Therefore *sum1* is missing for observations 2, 3 and 4, as is the case for observation 7.

```
1 list
2      +-----+
3      | id   trial1   trial2   trial3   sum1 |
4      |-----|
5  1. |  1     1.5     1.4     1.6     4.5 |
6  2. |  2     1.5          .     1.9     . |
7  3. |  3          .     2     1.6     . |
8  4. |  4          .          .     2.2     . |
9  5. |  5     1.9     2.1          2     6 |
10     |-----|
11  6. |  6     1.8          2     1.9     5.7 |
12  7. |  7          .          .          . |
13     +-----+
```

As a general rule, computations involving missing values yield missing values. For example,

```
1 2 + 2 yields 4
2 2 + . yields .
3 2 / 2 yields 1
4 . / 2 yields .
5 2 * 3 yields 6
6 2 * . yields .
```

whenever you add, subtract, multiply, divide, etc., values that involve missing data, the result is missing.

In our reaction time experiment, the total reaction time *sum1* is missing for four out of seven cases. We could try totaling the data for the non-missing trials by using the **rowtotal** function as shown in the example below.

```
1 egen sum2 = rowtotal(trial1 trial2 trial3)
2 list
```

The results below show that sum2 now contains the sum of the non-missing trials.

1		+-----+							
2			id	trial1	trial2	trial3	sum1	sum2	
3		-----							
4	1.		1	1.5	1.4	1.6	4.5	4.5	
5	2.		2	1.5	.	1.9	.	3.4	
6	3.		3	.	2	1.6	.	3.6	
7	4.		4	.	.	2.2	.	2.2	
8	5.		5	1.9	2.1	2	6	6	
9		-----							
10	6.		6	1.8	2	1.9	5.7	5.7	
11	7.		7	0	
12		+-----+							

Note that the `rowtotal` function treats missing as a zero value. When summing several variables it may not be reasonable to treat missing as zero if an observations is missing on all variables to be summed. The `rowtotal` function with the `missing` option will return a missing value if an observation is missing on all variables.

```
1 egen sum3 = rowtotal(trial1 trial2 trial3) , missing
2
3
4
5
6
7
8
9
10
11
12
13
14
```

	id	trial1	trial2	trial3	sum1	sum2	sum3
1.	1	1.5	1.4	1.6	4.5	4.5	4.5
2.	2	1.5	.	1.9	.	3.4	3.4
3.	3	.	2	1.6	.	3.6	3.6
4.	4	.	.	2.2	.	2.2	2.2
5.	5	1.9	2.1	2	6	6	6
6.	6	1.8	2	1.9	5.7	5.7	5.7
7.	7	0	.

Other statements work similarly. For example, observed what happened when we try to create an average variable without using a function (as in the example below). If any of the variables `trial1`, `trial2` or `trial3` are missing, the value for `avg1` are set to missing.

1	<code>egen avg1 = (trial1 + trial2 + trial3)/3</code>
---	---

Alternatively, the `rowmean` function averages the data for the non-missing trials in the same way as the `rowtotal` function.

1	<code>egen avg2 = rowmean(trial1 trial2 trial3)</code>
---	--

Note: Had there been large number of trials, say 50 trials, then it would be annoying to have to type `avg=rowmean(trial1 trial2 trial3 trial4 ...)`. Here is a shortcut you could use in this kind of situation:

```
1 egen avg3 = rowmean(trial1 - trial3)
2 list
3 +-----+
4 | id   trial1  trial2  trial3  avg1  avg2  avg3 |
5 |-----|
6 1. | 1      1.5      1.4      1.6      1.5   1.5   1.5 |
7 2. | 2      1.5      .        1.9      .     1.7   1.7 |
8 3. | 3      .        2        1.6      .     1.8   1.8 |
9 4. | 4      .        .        2.2      .     2.2   2.2 |
```



```

10 5. | 5      1.9      2.1      2      2      2      2 |
11   |-----|
12 6. | 6      1.8      2      1.9      1.9      1.9      1.9 |
13 7. | 7      .      .      .      .      .      . |
14   |-----+

```

Finally, you can use the `rowmiss` and `rownonmiss` functions to determine the number of missing and the number of non-missing values, respectively, in a list of variables. This is illustrated below.

```

1 egen miss = rowmiss(trial1 - trial3)
2 egen nomiss = rownonmiss(trial1 - trial3)
3 list

```

For variable *nomiss*, observations 1, 5 and 6 had three valid values, observations 2 and 3 had two valid values, observation 4 had only one valid value and observation 7 had no valid values. The variable *miss* shows the opposite, it provides a count of the number of missing values.

```

1      +-----+
2      | id   trial1   trial2   trial3   miss   nomiss |
3      |-----|
4 1. | 1     1.5     1.4     1.6     0     3 |
5 2. | 2     1.5     .     1.9     1     2 |
6 3. | 3     .     2     1.6     1     2 |
7 4. | 4     .     .     2.2     2     1 |
8 5. | 5     1.9     2.1     2     0     3 |
9      |-----|
10 6. | 6     1.8     2     1.9     0     3 |
11 7. | 7     .     .     .     3     0 |
12      +-----+

```

2.4.5 Missing values in logical statements

It is important to understand how missing values are handled in logical statements. For example, say that you want to create a 0/1 variable for *trial1* that is 1 if it is 1.5 or less, and 0 if it is over 1.5. We show this below (incorrectly, as you will see).

```

1 gen newvar1 =(trial2 <1.5)
2 list trial2 newvar1

```

It appears that something went wrong with our newly created variable *newvar1*! The observations with missing values for *trial2* were assigned a zero for *newvar1*.

```

1      +-----+
2      | trial2   newvar1 |
3      |-----|
4 1. | 1.4     1 |
5 2. | .     0 |
6 3. | 2     0 |
7 4. | .     0 |
8 5. | 2.1     0 |
9      |-----|
10 6. | 2     0 |
11 7. | .     0 |
12      +-----+

```

Let's explore why this happened by looking at the frequency table of *trial2*.

As you can see in the output, missing values are at the listed after the highest value 2.1 This is because STATA treats a missing value as the largest possible value (e.g., positive infinity) and that value is greater than 2.1, so then the values for *newvar1* become 0.

```

1 tab trial2, missing
2      trial2 |      Freq.      Percent      Cum.
3 -----+-----
4      1.4 |          1      14.29      14.29
5      2 |          2      28.57      42.86
6      2.1 |          1      14.29      57.14
7      . |          3      42.86     100.00
8 -----+-----
9      Total |          7     100.00

```

Now that we understand how STATA treats missing values, we will explicitly exclude missing values to make sure they are treated properly, as shown below.

```

1 gen newvar2 =(trial2 <1.5) if trial2 !=.
2 list trial2 newvar1 newvar2

```

As you can see in the STATA output below, the new variable *newvar2* has missing values for observations that are also missing for *trial2*.

```

1      +-----+
2      | trial2  newvar1  newvar2 |
3      |-----|
4  1. |    1.4        1        1 |
5  2. |     .         0         . |
6  3. |     2         0         0 |
7  4. |     .         0         . |
8  5. |    2.1        0         0 |
9      |-----|
10  6. |     2         0         0 |
11  7. |     .         0         . |
12      +-----+

```

2.4.6 Missing values in logical statements

When creating or recoding variables that involve missing values, always pay attention to whether the variable includes missing values.

2.4.7 For more information

- See the STATA FAQ: [How can I recode missing values into different categories?](#)
- See the STATA FAQ: [Can I quickly see how many missing values a variable has?](#) for more information on examining the number of missing and non-missing values for a particular variable or set of variables.