

Stata Learning Modules

Ethan Deng

Version 1.0

April 23, 2015

Notice: All the notes are from the website [UCLA Stata Learning Modules!](#) The latest version of this notes (Highlighted PDF format) can be found on [EthanDeng's Page](#).

1 Fundamentals of Using Stata (part I)

1.1 A Sample Stata Session: [Manuals13](#)

1.2 Descriptive information and statistics

This module shows common commands for showing descriptive information and descriptive statistics about data files.

1.2.1 Getting an overview of your file

The `sysuse` command loads a specified Stata-format dataset that was shipped with Stata. Here we will use the auto data file.

```
1 sysuse auto
```

The `describe` command shows you basic information about a Stata data file. As you can see, it tells us the number of observations in the file, the number of variables, the names of the variables, and more.

```
1 describe
2 Contains data from auto.dta
3 obs:          74
4 vars:          12          17 Feb 1999 10:49
5 size:          3,108 (99.6% of memory free)
6 -----
7 1. make      str17  %17s
8 2. price     int    %9.0g
9 3. mpg       byte   %9.0g
10 4. rep78     byte   %9.0g
11 5. hdroom    float  %9.0g
12 6. trunk     byte   %9.0g
13 7. weight    int    %9.0g
14 8. length    int    %9.0g
15 9. turn      byte   %9.0g
16 10. displ    int    %9.0g
17 11. gratio   float  %9.0g
18 12. foreign  byte   %9.0g
19 -----
20 Sorted by:
```

The `codebook` command is a great tool for getting a quick overview of the variables in the data file. It produces a kind of electronic codebook from the data file. Have a look at what it produces below.

```

1 codebook
2 make ----- (unlabeled)
3         type:  string (str17)
4
5         unique values:  74                coded missing:  0 / 74
6
7         examples:  "Cad. Deville"
8                   "Dodge Magnum"
9                   "Merc. XR-7"
10                  "Pont. Catalina"
11
12         warning:  variable has embedded blanks
13
14 price ----- (unlabeled)
15         type:  numeric (int)
16
17         range:  [3291,15906]                units:  1
18         unique values:  74                coded missing:  0 / 74
19
20         mean:  6165.26
21         std. dev:  2949.5
22
23         percentiles:      10%      25%      50%      75%      90%
24                        3895      4195      5006.5      6342      11385
25 //(omitted)

```

Another useful command for getting a quick overview of a data file is the `inspect` command. Here is what the `inspect` command produces for the auto data file.

```

1 inspect
2 price:
3 -----
4
5 | #          Negative      Total   Integers   Non-
6 | #          Zero         -       -         Integers
7 | #          Positive     74      74         -
8 | #          -----
9 | #          Total       74      74         -
10 | # # . . .      Missing   -
11 +-----
12 3291          15906      74
13 (74 unique values)
14
15 mpg:
16 -----
17
18 | #          Negative      Total   Integers   Non-
19 | #          Zero         -       -         Integers
20 | #          Positive     74      74         -
21 | # #          -----
22 | # # #          Total     74      74         -
23 | # # # # .      Missing   -
24 +-----
25 12          41          74

```

```

26 (21 unique values)
27 //(omitted)

```

The `list` command is useful for viewing all or a range of observations. Here we look at *make*, *price*, *mpg*, *rep78* and *foreign* for the first 10 observations.

```

1 list make price mpg rep78 foreign in 1/10
2
3      make      price      mpg      rep78      foreign
4  1.   Dodge Magnum   5886      16         2         0
5  2.    Datsun 510    5079      24         4         1
6  3.   Ford Mustang   4187      21         3         0
7  4.  Linc. Versailles 13466      14         3         0
8  5.   Plym. Sapporo   6486      26         .         0
9  6.   Plym. Arrow    4647      28         3         0
10 7.   Cad. Eldorado  14500      14         2         0
11 8.    AMC Spirit    3799      22         .         0
12 9.   Pont. Catalina   5798      18         4         0
13 10.  Chev. Nova     3955      19         3         0

```

1.2.2 Creating tables

The `tabulate` command is useful for obtaining frequency tables. Below, we make a table for *rep78* and a table for *foreign*. The command can also be shortened to `tab`.

```

1 tabulate rep78
2      rep78 |      Freq.      Percent      Cum.
3 -----+-----
4          1 |          2         2.90         2.90
5          2 |          8        11.59        14.49
6          3 |         30        43.48        57.97
7          4 |         18        26.09        84.06
8          5 |         11        15.94       100.00
9 -----+-----
10      Total |         69       100.00
11 tabulate foreign
12      foreign |      Freq.      Percent      Cum.
13 -----+-----
14          0 |         52        70.27        70.27
15          1 |         22        29.73       100.00
16 -----+-----
17      Total |         74       100.00

```

The `tab1` command can be used as a shortcut to request tables for a series of variables (instead of typing the `tabulate` command over and over again for each variable of interest).

```

1 tab1 rep78 foreign
2 -> tabulation of rep78
3
4      rep78 |      Freq.      Percent      Cum.
5 -----+-----
6          1 |          2         2.90         2.90
7          2 |          8        11.59        14.49
8          3 |         30        43.48        57.97
9          4 |         18        26.09        84.06
10         5 |         11        15.94       100.00
11 -----+-----

```

```

12      Total |          69      100.00
13
14 -> tabulation of foreign
15
16      foreign |          Freq.      Percent      Cum.
17 -----+-----
18          0 |          52      70.27      70.27
19          1 |          22      29.73      100.00
20 -----+-----
21      Total |          74      100.00

```

We can use the `plot` option to make a plot to visually show the tabulated values.

```

1 tabulate rep78, plot
2      rep78 |          Freq.
3 -----+-----
4          1 |          2 |**
5          2 |          8 |*****
6          3 |         30 |*****
7          4 |         18 |*****
8          5 |         11 |*****
9 -----+-----
10      Total |          69

```

We can also make crosstabs using `tabulate`. Let's look at the repair history broken down by *foreign* and *domestic* cars.

```

1 tabulate rep78 foreign
2      |          foreign
3      rep78 |          0          1 |          Total
4 -----+-----
5          1 |          2          0 |          2
6          2 |          8          0 |          8
7          3 |         27          3 |         30
8          4 |          9          9 |         18
9          5 |          2          9 |         11
10 -----+-----
11      Total |         48         21 |         69

```

With the `column` option, we can request column percentages. Notice that about 86% of the foreign cars received a rating of 4 or 5. Only about 23% of domestic cars were rated that highly.

```

1 tabulate rep78 foreign, column
2      |          foreign
3      rep78 |          0          1 |          Total
4 -----+-----
5          1 |          2          0 |          2
6          |         4.17         0.00 |         2.90
7 -----+-----
8          2 |          8          0 |          8
9          |        16.67         0.00 |        11.59
10 -----+-----
11          3 |         27          3 |         30
12          |        56.25        14.29 |        43.48
13 -----+-----
14          4 |          9          9 |         18
15          |        18.75        42.86 |        26.09
16 -----+-----

```

17	5		2	9		11
18			4.17	42.86		15.94
19	-----+-----+-----					
20	Total		48	21		69
21			100.00	100.00		100.00

We can use the `nofreq` option to suppress the frequencies, and just focus on the percentages.

```
1 tabulate rep78 foreign, column nofreq
2
3      |      foreign
4 rep78 |      0      1 |      Total
5 -----+-----+-----
6      1 |      4.17      0.00 |      2.90
7      2 |     16.67      0.00 |     11.59
8      3 |     56.25     14.29 |     43.48
9      4 |     18.75     42.86 |     26.09
10     5 |      4.17     42.86 |     15.94
11 -----+-----+-----
12 Total |    100.00    100.00 |    100.00
```

Note that the order of the options does not matter. Just remember that the options must come after the comma.

```
1 tabulate rep78 foreign, nofreq column
2
3      |      foreign
4 rep78 |      0      1 |      Total
5 -----+-----+-----
6      1 |      4.17      0.00 |      2.90
7      2 |     16.67      0.00 |     11.59
8      3 |     56.25     14.29 |     43.48
9      4 |     18.75     42.86 |     26.09
10     5 |      4.17     42.86 |     15.94
11 -----+-----+-----
12 Total |    100.00    100.00 |    100.00
```

1.2.3 Generating summary statistics with summarize

For summary statistics, we can use the `summarize` command. Let's generate some summary statistics on *mpg*.

```
1 summarize mpg
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----+-----+-----+-----+-----
4      mpg |       74    21.2973    5.785503       12       41
```

We can use the `detail` option of the `summarize` command to get more detailed summary statistics.

```
1 summarize mpg, detail
2
3      mpg
4 -----+-----+-----+-----+-----+-----
5 Percentiles      Smallest
6 1%             12             12
7 5%             14             12
8 10%            14             14      Obs              74
9 25%            18             14      Sum of Wgt.      74
10 50%            20
11                      Largest      Mean              21.2973
12 75%            25             34      Std. Dev.        5.785503
13 90%            29             35      Variance          33.47205
```

14	95%	34	35	Skewness	.9487176
15	99%	41	41	Kurtosis	3.975005

To get these values separately for *foreign* and *domestic*, we could use the `by foreign:` prefix as shown below. Note that we first had to `sort` the data before using `by foreign:`.

```

1 sort foreign
2 by foreign: summarize mpg
3 -> foreign= 0
4 Variable |      Obs      Mean   Std. Dev.      Min      Max
5 -----+-----
6      mpg |      52   19.82692   4.743297      12      34
7
8 -> foreign= 1
9 Variable |      Obs      Mean   Std. Dev.      Min      Max
10 -----+-----
11      mpg |      22   24.77273   6.611187      14      41

```

This is not the most efficient way to do this. Another way, which does not require the data to be sorted, is by using the `summarize()` option as part of the `tabulate` command.

```

1 tabulate foreign, summarize(mpg)
2      |      Summary of mpg
3 foreign |      Mean   Std. Dev.      Freq.
4 -----+-----
5      0 |   19.826923   4.7432972      52
6      1 |   24.772727   6.6111869      22
7 -----+-----
8      Total |   21.297297   5.7855032      74

```

Here is another example, showing the average price of cars for each level of repair history.

```

1 tabulate rep78, summarize(price)
2      |      Summary of price
3 rep78 |      Mean   Std. Dev.      Freq.
4 -----+-----
5      1 |   4564.5    522.55191      2
6      2 |   5967.625   3579.3568      8
7      3 |   6429.2333   3525.1398     30
8      4 |   6071.5    1709.6083     18
9      5 |    5913     2615.7628     11
10 -----+-----
11      Total |   6146.0435   2912.4403     69

```

1.2.4 Summary

- `describe`: provide information about the current data file, including the number of variables and observations and a listing of the variables in a data file.
- `codebook`: produce codebook like information for the current data file.
- `inspect`: provide a quick overview of data file.
- `list` make mpg: list out the variables make and mpg.
- `tabulate` mpg: make a table of mpg.
- `tabulate` rep78 foreign: make a two way table of rep78 by foreign.
- `summarize` mpg price: produce summary statistics of mpg and price.
- To produce summary statistics for mpg separately for foreign and domestic cars, use

```

1 sort foreign
2 by foreign: summarize(mpg)

```

- `tabulate foreign, summarize(mpg)`: produce summary statistics for mpg by foreign (prior sorting not required).

1.3 Getting help using Stata

This module shows resources you can use to help you learn and use Stata.

1.3.1 Stata online help

When you know the name of the command you want to use (e.g., `summarize`), you can use the Stata help to get a quick summary of the command and its syntax. You can do this in two ways:

1. type `help summarize` in the command window, or
2. click **Help, Stata Command**, then type `summarize`.

Here is what help `summarize` looks like.

```

1 help summarize
2 help summarize                                dialog: summarize
3 -----
4
5 Title
6
7 [R] summarize -- Summary statistics
8
9
10 Syntax
11
12 summarize [varlist] [if] [in] [weight] [, options]
13
14 options          description
15 -----
16 Main
17 detail           display additional statistics
18 meanonly         suppress the display; only calculate the
19                  mean; programmer's option
20 format           use variable's display format
21 separator(#)     draw separator line after every # variables;
22                  default is separator(5)
23 -----
24 varlist may contain time-series operators; see tsvarlist.
25 by may be used with summarize; see by.
26 aweights, fweights, and iweights are allowed. However,
27 iweights may not be used with the detail option; see weight.
28 //(omitted)

```

If you use the pull-down menu to get help for a command, it shows the same basic information but related commands and topics are hotlinks you can click.

When you want to search for a keyword, e.g. `memory`, you can use Stata to search for help topics that contain that keyword. You can do this in two ways:

1. Type search `memory` in the command window, or
2. Click **Help, Search**, then `memory`.

Here is what search memory looks like.

```
1 search memory
2
3 GS      . . . . . Getting Started manual
4
5 [U]     Chapter 7 . . . . . Setting the size of memory
6         (help memory)
7
8 [R]     compress . . . . . Compress data in memory
9         (help compress)
10 //(omitted)
```

As you can see, there are lots of help topics that refer to memory. Some of the topics give you a command, and then you can get help for that command. Notice that those topics start with **GS**[U] or **[R]**. Those are indicating which Stata manual you could find the command (GS=Getting Started, U=Users Guide, R=Reference Guide).

The next set of topics all start with **FAQ** because these are Frequently Asked Questions from the Stata web site. You can see the title of the FAQ and the address of the FAQ. Lastly, there is a topic that starts with **STB** which stands for Stata Technical Bulletin. These refer to add-on programs that you can install into Stata. There are dozens, if not hundreds of specialized and useful programs that you can get from the Stata Technical Bulletin.

You can access this same kind of help from the pull-down menus by clicking **Help** then **Search** then type `memory`. Note how the related commands, the FAQs, and the STB all have hotlinks you can click. For example, you can click on a FAQ and it will bring up that FAQ in your web browser. Or, you could click on an STB and it would walk you through the steps of installing that STB into your copy of Stata. As you can see, there are real advantages to using the pull-down menus for getting help because it is so easy to click on the related topics.

1.3.2 Stata sample data files

Stata has some very useful data files available to you for learning and practicing Stata. For example, you can type

```
1 sysuse auto
```

to use the auto data file that comes with Stata. You can type

```
1 sysuse dir
```

to see the entire list of data files that ship with Stata. You can type

```
1 help dta_contents
```

to see all of the sample data files that you can easily access from within Stata.

1.3.3 Stata web pages

The Stata web page is a wonderful resource. You can visit the main page at <http://www.stata.com>.

The User Support page (click **User Support** from main page) has a great set of resources, including

- FAQs
- NetCourses
- StataList: How to subscribe
- StataList: Archives
- Statalist ado-file Archives
- Stata Bookstore

In the bookstore, you can find books on Stata. A good intro book on Stata is **Statistics with Stata**.

2 Fundamentals of Using Stata (part II)

2.1 Using IF with Stata commands

This module shows the use of **if** with common Stata commands.

Let's use the auto data file.

```
1 sysuse auto
```

For this module, we will focus on the variables *make*, *rep78*, *foreign*, *mpg*, and *price*. We can use the **keep** command to keep just these five variables.

```
1 keep make rep78 foreign mpg price
```

Let's make a table of *rep78* by *foreign* to look at the repair histories of the foreign and domestic cars.

```
1 tabulate rep78 foreign
2
3      rep78 |      foreign
4      -----+-----
5      1 |      2      0 |      2
6      2 |      8      0 |      8
7      3 |     27      3 |     30
8      4 |      9      9 |     18
9      5 |      2      9 |     11
10     -----+-----
11     Total |     48     21 |     69
```

Suppose we wanted to focus on just the cars with repair histories of four or better. We can use **if** suffix to do this.

```
1 tabulate rep78 foreign if rep78 >= 4
2
3      rep78 |      foreign
4      -----+-----
5      4 |      9      9 |     18
6      5 |      2      9 |     11
7      -----+-----
8     Total |     11     18 |     29
```

Let's make the above table using the **column** and **nofreq** options. The command **column** requests column percentages while the command **nofreq** suppresses cell frequencies. Note that **column** and **nofreq** come after the comma. These are options on the **tabulate** command and options need to be placed after a comma.

```
1 tabulate rep78 foreign if rep78 >= 4, column nofreq
2
3      rep78 |      foreign
4      -----+-----
5      4 |    81.82    50.00 |    62.07
6      5 |    18.18    50.00 |    37.93
7      -----+-----
8     Total |   100.00   100.00 |   100.00
```

The use of **if** is not limited to the **tabulate** command. Here, we use it with the **list** command.

```
1 list if rep78 >= 4
2
3      make      price      mpg      rep78      foreign
4  3.      AMC Spirit      3799      22      .      0
```

```

4   5.      Buick Electra      7827      15      4      0
5   7.      Buick Opel       4453      26      .      0
6  15.     Chev. Impala      5705      16      4      0
7  20.     Dodge Colt       3984      30      5      0
8  24.     Ford Fiesta       4389      28      4      0
9  29.     Merc. Bobcat      3829      22      4      0
10 30.     Merc. Cougar      5379      14      4      0
11 //(omitted)

```

Did you see that some of the observations had a value of '.' for rep78? These are missing values. For example, the value of *rep78* for the AMC Spirit is missing. **Stata treats a missing value as positive infinity**, the highest number possible. So, when we said `list if rep78 >= 4`, Stata included the observations where *rep78* was '.' as well.

If we wanted to include just the valid (non-missing) observations that are greater than or equal to 4, we can do the following to tell Stata we want only observations where *rep78* >= 4 and *rep78* is not missing.

```

1 list if rep78 >= 4 & !missing(rep78)
2           make      price      mpg      rep78      foreign
3   5.      Buick Electra      7827      15      4      0
4  15.     Chev. Impala      5705      16      4      0
5  20.     Dodge Colt       3984      30      5      0
6  24.     Ford Fiesta       4389      28      4      0
7  29.     Merc. Bobcat      3829      22      4      0
8  30.     Merc. Cougar      5379      14      4      0
9  33.     Merc. XR-7        6303      14      4      0
10 35.     Olds 98           8814      21      4      0
11 //(omitted)

```

This code will also yield the same output as above.

```

1 list if rep78 >= 4 & rep78 != .

```

We can use `if` with most Stata commands. Here, we get summary statistics for *price* for cars with repair histories of 1 or 2. Note the double equal (==) represents **IS EQUAL TO** and the pipe (|) represents **OR**.

```

1 summarize price if rep78 == 1 | rep78 == 2
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      10      5687   3216.375     3667    14500

```

A simpler way to say this would be ...

```

1 summarize price if rep78 <= 2
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      10      5687   3216.375     3667    14500

```

Likewise, we can do this for cars with repair history of 3, 4 or 5.

```

1 summarize price if rep78 == 3 | rep78 == 4 | rep78 == 5
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4   price |      59     6223.847   2880.454     3291    15906

```

Additionally, we can use this code to designate a range of values. Here is a summary of *price* for the values 3 through 5 in *rep78*.

```

1 summarize price if inrange(rep78,3,5)
2 Variable |      Obs      Mean   Std. Dev.      Min      Max

```

```

3 -----+-----
4 price |          59    6223.847    2880.454    3291    15906

```

Let's simplify this by saying `rep78 >= 3`.

```

1 summarize price if rep78 >= 3
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4 price |      64    6239.984    2925.843     3291    15906

```

Did you see the mistake we made? We accidentally included the missing values because we forgot to exclude them. We really needed to say.

```

1 summarize price if rep78 >= 3 & !missing(rep78)
2 Variable |      Obs      Mean   Std. Dev.      Min      Max
3 -----+-----
4 price |      59    6223.847    2880.454     3291    15906

```

2.1.1 Taking a random sample

It is also possible to take a simple random sample of your data using the `sample` command. This information can be found on our STATA FAQ page: [How can I draw a random sample of my data?](#)

2.1.2 Summary

Most Stata commands can be followed by `if`, for example

```

1 summarize if rep78 == 2
2 summarize if rep78 >= 2
3 summarize if rep78 > 2
4 summarize if rep78 <= 2
5 summarize if rep78 < 2
6 summarize if rep78 != 2

```

`if` expressions can be connected with `|` for **OR**, `&` for **AND**.

2.1.3 Missing Values

Missing values are represented as `.` and are the highest value possible. Therefore, when values are missing, be careful with commands like

```

1 summarize if rep78 > 3
2 summarize if rep78 >= 3
3 summarize if rep78 != 3

```

to omit missing values, use

```

1 summarize if rep78 > 3 & !missing(rep78)
2 summarize if rep78 >= 3 & !missing(rep78)
3 summarize if rep78 != 3 & !missing(rep78)

```

2.2 A statistical sampler in Stata

Version info: Code for this page was tested in Stata 12.

This module will give a brief overview of some common statistical tests in Stata. Let's use the `auto` data file that we will use for our examples.

`auto`

```
1 sysuse auto
```

2.2.1 t-tests

Let's do a t-test comparing the miles per gallon (*mpg*) of foreign and domestic cars.

```
1 ttest mpg , by(foreign)
2 Two-sample t test with equal variances
3
4 -----
5      Group |      Obs      Mean      Std. Err.      Std. Dev.      [95% Conf. Interval]
6 -----+-----
7           0 |       52   19.82692   .6577777   4.743297   18.50638   21.14747
8           1 |       22   24.77273   1.40951   6.611187   21.84149   27.70396
9 -----+-----
10 combined |       74   21.2973   .6725511   5.785503   19.9569   22.63769
11 -----+-----
12      diff |           -4.945804   1.362162           -7.661225   -2.230384
13 -----
14 Degrees of freedom: 72
15
16                      Ho: mean(0) - mean(1) = diff = 0
17
18      Ha: diff < 0 Ha: diff ~= "0" Ha: diff > 0
19      t =   -3.6308           t =   -3.6308           t =   -3.6308
20      P < t =   0.0003           P > |t| =   0.0005           P > t =   0.9997
```

As you see in the output above, the domestic cars had significantly lower *mpg* (19.8) than the foreign cars (24.7).

2.2.2 Chi-square

Let's compare the repair rating (*rep78*) of the foreign and domestic cars. We can make a crosstab of *rep78* by *foreign*. We may want to ask whether these variables are independent. We can use the *chi2* option to request a chi-square test of independence as well as the crosstab.

```
1 tabulate rep78 foreign, chi2
2
3      rep78 |      foreign
4      rep78 |      0      1 |      Total
5 -----+-----
6           1 |       2       0 |       2
7           2 |       8       0 |       8
8           3 |      27       3 |      30
9           4 |       9       9 |      18
10          5 |       2       9 |      11
11 -----+-----
12      Total |      48      21 |      69
13
14      Pearson chi2(4) = 27.2640 Pr = 0.000
```

The chi-square is not really valid when you have empty cells. In such cases when you have empty cells, or cells with small frequencies, you can request Fisher's exact test with the *exact* option.

```
1 tabulate rep78 foreign, chi2 exact
2
3      rep78 |      foreign
4      rep78 |      0      1 |      Total
5 -----+-----
```

5	1		2	0		2
6	2		8	0		8
7	3		27	3		30
8	4		9	9		18
9	5		2	9		11
10	-----+-----+-----					
11	Total		48	21		69
12						
13	Pearson chi2 (4) = 27.2640 Pr = 0.000					
14	Fisher's exact = 0.000					

2.2.3 Correlation

We can use the `correlate` command to get the correlations among variables. Let's look at the correlations among *price* *mpg* *weight* and *rep78*. (We use *rep78* in the correlation even though it is not continuous to illustrate what happens when you use correlate with variables with missing data.)

```
1 correlate price mpg weight rep78
2 (obs=69)
```

3			price	mpg	weight	rep78
4						
5						
6	price		1.0000			
7	mpg		-0.4559	1.0000		
8	weight		0.5478	-0.8055	1.0000	
9	rep78		0.0066	0.4023	-0.4003	1.0000

Note that the output above said (obs=69). The `correlate` command drops data on a listwise basis, meaning that if any of the variables are missing, then the entire observation is omitted from the correlation analysis.

We can use `pwcorr` (pairwise correlations) if we want to obtain correlations that deletes missing data on a pairwise basis instead of a listwise basis. We will use the `obs` option to show the number of observations used for calculating each correlation.

```
1 pwcorr price mpg weight rep78, obs
2
```

3			price	mpg	weight	rep78
4						
5						
6						
7	price		1.0000			
8			74			
9						
10	mpg		-0.4686	1.0000		
11			74	74		
12						
13	weight		0.5386	-0.8072	1.0000	
14			74	74	74	
15						
16	rep78		0.0066	0.4023	-0.4003	1.0000
17			69	69	69	69
18						

Note how the correlations that involve *rep78* have an N of 69 compared to the other correlations that have an N of 74. This is because *rep78* has five missing values, so it only had 69 valid observations, but the other variables had no missing data so they had 74 valid observations.

2.2.4 Regression

Let's look at doing regression analysis in Stata. For this example, let's drop the cases where *rep78* is 1 or 2 or missing.

```
1 drop if (rep78 <= 2) | (rep78 ==.)
2 (15 observations deleted)
```

Now, let's predict *mpg* from *price* and *weight*. As you see below, *weight* is a significant predictor of *mpg*, but *price* is not.

```
1 regress mpg price weight
2
3 Source |      SS      df      MS                Number of obs =      59
4 -----+-----
5 Model | 1375.62097      2  687.810483            F( 2, 56) = 47.87
6 Residual | 804.616322     56  14.3681486            Prob > F      = 0.0000
7 -----+-----
8 Total | 2180.23729     58  37.5902981            R-squared     = 0.6310
9                                     Adj R-squared = 0.6178
10                                    Root MSE     = 3.7905
11
12 -----+-----
13 mpg |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
14 -----+-----
15 price | -.0000139   .0002108     -0.066   0.948    - .0004362   .0004084
16 weight | -.005828   .0007301     -7.982   0.000    - .0072906  - .0043654
17 _cons | 39.08279   1.855011     21.069   0.000     35.36676   42.79882
```

What if we wanted to predict *mpg* from *rep78* as well. *rep78* is really more of a categorical variable than it is a continuous variable. To include it in the regression, we should convert *rep78* into dummy variables. Fortunately, Stata makes dummy variables easily using *tabulate*. The *gen(rep)* option tells Stata that we want to generate dummy variables from *rep78* and we want the stem of the dummy variables to be *rep*.

```
1 tabulate rep78, gen(rep)
2
3 rep78 |      Freq.      Percent      Cum.
4 -----+-----
5 3 |      30      50.85      50.85
6 4 |      18      30.51      81.36
7 5 |      11      18.64     100.00
8 -----+-----
9 Total |      59     100.00
```

Stata has created *rep1* (1 if *rep78* is 3), *rep2* (1 if *rep78* is 4) and *rep3* (1 if *rep78* is 5). We can use the *tabulate* command to verify that the dummy variables were created properly.

```
1 tabulate rep78 rep1
2
3 rep78 |      0      1 |      Total
4 -----+-----
5 3 |      0      30 |      30
6 4 |     18      0 |      18
7 5 |     11      0 |      11
8 -----+-----
9 Total |     29     30 |      59
10 tabulate rep78 rep2
11
12 rep78 |      0      1 |      Total
13 -----+-----
```

14	3		30	0		30
15	4		0	18		18
16	5		11	0		11
17	-----+-----+-----					
18	Total		41	18		59
19	tabulate rep78 rep3					
20			rep78==	5.0000		
21	rep78		0	1		Total
22	-----+-----+-----					
23	3		30	0		30
24	4		18	0		18
25	5		0	11		11
26	-----+-----+-----					
27	Total		48	11		59

Now we can include *rep1* and *rep2* as dummy variables in the regression model.

1	regress mpg price weight rep1 rep2						
2							
3	Source		SS	df	MS	Number of obs	= 59
4	-----+-----						
5	Model		1435.91975	4	358.979938	F(4, 54)	= 26.04
6	Residual		744.317536	54	13.7836581	Prob > F	= 0.0000
7	-----+-----						
8	Total		2180.23729	58	37.5902981	R-squared	= 0.6586
9							
10	-----+-----						
11	mpg		Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
12	-----+-----						
13	price		-.0001126	.0002133	-0.53	0.600	-.0005403 .0003151
14	weight		-.005107	.0008236	-6.20	0.000	-.0067584 -.0034557
15	rep1		-2.886288	1.504639	-1.92	0.060	-5.902908 .1303314
16	rep2		-2.88417	1.484817	-1.94	0.057	-5.861048 .0927086
17	_cons		39.89189	1.892188	21.08	0.000	36.09828 43.6855
18	-----+-----						

2.2.5 Analysis of variance

If you wanted to do an analysis of variance looking at the differences in *mpg* among the three repair groups, you can use the *oneway* command to do this.

1	oneway mpg rep78					
2	Analysis of Variance					
3	Source	SS	df	MS	F	Prob > F
4	-----+-----					
5	Between groups	506.325167	2	253.162583	8.47	0.0006
6	Within groups	1673.91212	56	29.8912879		
7	-----+-----					
8	Total	2180.23729	58	37.5902981		
9						
10	Bartlett's test for equal variances: chi2(2) = 9.9384 Prob>chi2 = 0.007					

If you include the *tabulate* option, you get mean *mpg* for the three groups, which shows that the group with the best repair rating (*rep78* of 5) also has the highest *mpg* (27.3).

```
1 oneway mpg rep78, tabulate
2
```

```

3      |           Summary of mpg
4      rep78 |           Mean   Std. Dev.   Freq.
5      -----+-----
6      3 | 19.433333  4.1413252   30
7      4 | 21.666667  4.9348699   18
8      5 | 27.363636  8.7323849   11
9      -----+-----
10     Total | 21.59322  6.1310927   59
11
12           Analysis of Variance
13     Source           SS           df           MS           F           Prob > F
14     -----+-----
15     Between groups   506.325167         2    253.162583     8.47     0.0006
16     Within groups   1673.91212         56    29.8912879
17     -----+-----
18     Total           2180.23729         58    37.5902981
19
20     Bartlett's test for equal variances:  chi2(2) = 9.9384  Prob>chi2 = 0.007

```

If you want to include covariates, you need to use the `anova` command. The continuous(price weight) option tells Stata that those variables are covariates.

```

1  anova mpg rep78 c.price c.weight
2              Number of obs =      59      R-squared      = 0.6586
3              Root MSE      = 3.71263      Adj R-squared = 0.6333
4
5              Source | Partial SS   df           MS           F           Prob > F
6              -----+-----
7              Model | 1435.91975   4    358.979938     26.04     0.0000
8              |
9              rep78 | 60.2987853   2    30.1493926     2.19     0.1221
10             price | 3.8421233    1    3.8421233      0.28     0.5997
11             weight | 529.932889   1    529.932889    38.45     0.0000
12             |
13             Residual | 744.317536   54    13.7836581
14             -----+-----
15             Total | 2180.23729   58    37.5902981

```

2.3 An overview of Stata syntax

This module shows the general structure of Stata commands. We will demonstrate this using `summarize` as an example, although this general structure applies to most Stata commands.

Note: This code was tested in Stata 12.

Let's first use the auto data file.

```
1 use auto
```

As you have seen, we can type `summarize` and it will give us summary statistics for all of the variables in the data file.

```

1  summarize
2  Variable |      Obs       Mean   Std. Dev.       Min       Max
3  -----+-----
4     make |         0
5     price |        74    6165.257   2949.496     3291    15906
6     mpg   |        74     21.2973    5.785503        12        41

```


7	rep78		69	3.405797	.9899323	1	5
8	hdroom		74	2.993243	.8459948	1.5	5
9	trunk		74	13.75676	4.277404	5	23
10	weight		74	3019.459	777.1936	1760	4840
11	length		74	187.9324	22.26634	142	233
12	turn		74	39.64865	4.399354	31	51
13	displ		74	197.2973	91.83722	79	425
14	gratio		74	3.014865	.4562871	2.19	3.89
15	foreign		74	.2972973	.4601885	0	1

It is also possible to obtain means for specific variables. For example, below we get summary statistics just for *mpg* and *price*.

```
1 summarize mpg price
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	74	21.2973	5.785503	12	41
price	74	6165.257	2949.496	3291	15906

We could further tell Stata to limit the summary statistics to just foreign cars by adding an if qualifier.

```
1 summarize mpg price if (foreign == 1)
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	22	24.77273	6.611187	14	41
price	22	6384.682	2621.915	3748	12990

The *if* qualifier can contain more than one condition. Here, we ask for summary statistics for the foreign cars which get less than 30 miles per gallon.

```
1 summarize mpg price if foreign == 1 & mpg < 30
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
mpg	17	21.94118	3.896643	14	28
price	17	6996.235	2674.552	3895	12990

We can use the **detail** option to ask Stata to give us more detail in the summary statistics. Notice that the **detail** option goes after the comma. If the comma were omitted, Stata would give an error.

```
1 summarize mpg price if foreign == 1 & mpg < 30 , detail
```

mpg					

Percentiles		Smallest			
1%	14	14			
5%	14	17			
10%	17	17	Obs		17
25%	18	18	Sum of Wgt.		17
50%	23		Mean		21.94118
		Largest	Std. Dev.		3.896643
75%	25	25			
90%	26	25	Variance		15.18382
95%	28	26	Skewness		-.4901235
99%	28	28	Kurtosis		2.201759
price					

Percentiles		Smallest			

20	1%	3895	3895		
21	5%	3895	4296		
22	10%	4296	4499	Obs	17
23	25%	5079	4697	Sum of Wgt.	17
24					
25	50%	6229		Mean	6996.235
26			Largest	Std. Dev.	2674.552
27	75%	8129	9690		
28	90%	11995	9735	Variance	7153229
29	95%	12990	11995	Skewness	.9818272
30	99%	12990	12990	Kurtosis	2.930843

Note that even though we built these parts up one at a time, they don't have to go together. Let's look at some other forms of the summarize command.

You can tell Stata which observation numbers you want using the in qualifier. Here we ask for summaries of observations 1 to 10. This is useful if you have a big data file and want to try out a command on a subset of observations.

```
1 summarize in 1/10
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
make	0				
price	10	5517.4	2063.518	3799	10372
mpg	10	19.5	3.27448	15	26
rep78	8	3.125	.3535534	3	4
hdroom	10	3.3	.7527727	2	4.5
trunk	10	14.7	3.88873	10	21
weight	10	3271	558.3796	2230	4080
length	10	194	19.32759	168	222
turn	10	40.2	3.259175	34	43
displ	10	223.9	71.77503	121	350
gratio	10	2.907	.3225264	2.41	3.58
foreign	10	0	0	0	0

Also, recall that you can ask Stata to perform summaries for foreign and domestic cars separately using by, as shown below.

```
1 sort foreign
2 by foreign: summarize
3 -> foreign= 0
```

Variable	Obs	Mean	Std. Dev.	Min	Max
-----+-----					
make	0				
price	52	6072.423	3097.104	3291	15906
mpg	52	19.82692	4.743297	12	34
rep78	48	3.020833	.837666	1	5
hdroom	52	3.153846	.9157578	1.5	5
trunk	52	14.75	4.306288	7	23
weight	52	3317.115	695.3637	1800	4840
length	52	196.1346	20.04605	147	233
turn	52	41.44231	3.967582	31	51
displ	52	233.7115	85.26299	86	425
gratio	52	2.806538	.3359556	2.19	3.58
foreign	52	0	0	0	0

```
18
19 -> foreign= 1
```

Variable	Obs	Mean	Std. Dev.	Min	Max
make	0				
price	22	6384.682	2621.915	3748	12990
mpg	22	24.77273	6.611187	14	41
rep78	21	4.285714	.7171372	3	5
hdroom	22	2.613636	.4862837	1.5	3.5
trunk	22	11.40909	3.216906	5	16
weight	22	2315.909	433.0035	1760	3420
length	22	168.5455	13.68255	142	193
turn	22	35.40909	1.501082	32	38
displ	22	111.2273	24.88054	79	163
gratio	22	3.507273	.2969076	2.98	3.89
foreign	22	1	0	1	1

Let's review all those pieces.

A command can be preceded with a **by** prefix, as shown below.

```
1 by foreign: summarize
```

There are many parts that can come after a command. They are each presented separately below.

For example, **summarize** followed by the names of variables.

```
1 summarize mpg price
```

summarize with **in** specifying a range of records to be summarized.

```
1 summarize in 1/10
```

summarize with simple **if** specifying records to summarize.

```
1 summarize if foreign == 1
```

summarize with complex **if** specifying records to summarize.

```
1 summarize if foreign == 1 & mpg > 30
```

summarize followed by option(s).

```
1 summarize , detail
```

So, putting it all together, the general syntax of the **summarize** command can be described as:

```
1 [by varlist:] summarize [varlist] [in range] [if exp] , [options]
```

Understanding the overall syntax of Stata commands helps you remember them and use them more effectively, and it also aids you understand the help files in Stata. All the extra stuff about **by**, **if** and **in** could be confusing. Let's have a look at the help file for **summarize**. It makes more sense knowing what the **by**, **if** and **in** parts mean.

```
1 help summarize
2 -----
3 help for summarize (manual: [R] summarize)
4 -----
5
6 Summary statistics
7 -----
8
9 [by varlist:] summarize [varlist] [weight] [if exp] [in range]
10 [, { detail | meanonly } format ]
```

2.4 Missing data

2.4.1 Introduction

This module will explore missing data in STATA, focusing on numeric missing data. It will describe how to indicate missing data in your raw data files, as well as how missing data are handled in STATA logical commands and assignment statements.

We will illustrate some of the missing data properties in STATA using data from a reaction time study with eight subjects indicated by the variable *id*, and the subjects reaction times were measured at three time points (*trial1* *trial2* *trial3*). The input data file is shown below.

```
1 input id trial1 trial2 trial3
2 1 1.5 1.4 1.6
3 2 1.5 . 1.9
4 3 . 2.0 1.6
5 4 . . 2.2
6 5 1.9 2.1 2
7 6 1.8 2.0 1.9
8 7 . . .
9 end
10 list
```

You might notice that some of the reaction times are coded using a single '.' as is the case for subject 2. The person measuring time for that trial did not measure the response time properly, therefore the data for the second trial is missing.

```
1      +-----+
2      | id   trial1   trial2   trial3 |
3      |-----|
4  1. | 1     1.5     1.4     1.6 |
5  2. | 2     1.5     .     1.9 |
6  3. | 3     .     2     1.6 |
7  4. | 4     .     .     2.2 |
8  5. | 5     1.9     2.1     2 |
9      |-----|
10 6. | 6     1.8     2     1.9 |
11 7. | 7     .     .     . |
12      +-----+
```

2.4.2 How STATA handles missing data in STATA procedures

As a general rule, STATA commands that perform computations of any type handle missing data by omitting the missing values. However, the way that missing values are omitted is not always consistent across commands, so let's take a look at some examples.

First, let's `summarize` our reaction time variables and see how STATA handles the missing values.

```
1 summarize trial1 trial2 trial3
```

As you see in the output below, `summarize` computed means using 4 observations for *trial1* and *trial2* and 6 observations for *trial3*. In short, the `summarize` command performed the computations on all the available data.

Variable	Obs	Mean	Std. Dev.	Min	Max
trial1	4	1.675	.2061553	1.5	1.9
trial2	4	1.875	.3201562	1.4	2.1
trial3	6	1.866667	.233809	1.6	2.2

A second example, shows how the `tabulation` or `tab1` command handles missing data. Like `summarize`, `tab1` uses just available data. Note that the percentages are computed based on the total number of non-missing cases.

```

1 tab1 trial1 trial2 trial3
2 -> tabulation of trial1
3
4      trial1 |      Freq.      Percent      Cum.
5 -----+-----
6          1.5 |          2       50.00       50.00
7          1.8 |          1       25.00       75.00
8          1.9 |          1       25.00      100.00
9 -----+-----
10         Total |          4      100.00
11
12 -> tabulation of trial2
13
14      trial2 |      Freq.      Percent      Cum.
15 -----+-----
16          1.4 |          1       25.00       25.00
17          2   |          2       50.00       75.00
18          2.1 |          1       25.00      100.00
19 -----+-----
20         Total |          4      100.00
21
22 -> tabulation of trial3
23
24      trial3 |      Freq.      Percent      Cum.
25 -----+-----
26          1.6 |          2       33.33       33.33
27          1.9 |          2       33.33       66.67
28          2   |          1       16.67       83.33
29          2.2 |          1       16.67      100.00
30 -----+-----
31         Total |          6      100.00

```

It is possible that you might want the percentages to be computed out of the total number of observations, and the percentage missing for each variable shown in the table. This can be achieved by including the `missing` option after the `tabulation` command,

```

1 tab1 trial1 trial2 trial3, m
2 -> tabulation of trial1
3
4      trial1 |      Freq.      Percent      Cum.
5 -----+-----
6          1.5 |          2       28.57       28.57
7          1.8 |          1       14.29       42.86
8          1.9 |          1       14.29       57.14
9          .   |          3       42.86      100.00
10 -----+-----
11         Total |          7      100.00
12
13 -> tabulation of trial2
14
15      trial2 |      Freq.      Percent      Cum.
16 -----+-----
17          1.4 |          1       14.29       14.29
18          2   |          2       28.57       42.86

```

```

19      2.1 |      1      14.29      57.14
20      .  |      3      42.86     100.00
21 -----+-----
22      Total |      7     100.00
23
24 -> tabulation of trial3
25
26      trial3 |      Freq.      Percent      Cum.
27 -----+-----
28      1.6  |      2      28.57     28.57
29      1.9  |      2      28.57     57.14
30      2    |      1      14.29     71.43
31      2.2  |      1      14.29     85.71
32      .    |      1      14.29     100.00
33 -----+-----
34      Total |      7     100.00

```

Let's look at how the `correlate` command handles missing data. We would expect that it would perform the computations based on the available data, and omit the missing values. Here is an example command.

```
1 corr trial1 trial2 trial3
```

The output is show below. Note how the missing values were excluded. For each pair variables, the `corr` command used the number of pairs that had valid data. For the pair formed by *trial1* and *trial2*, there were 3 pairs with valid data. For the pairing of *trial1* and *trial3* there were 4 valid pairs, and likewise there were 4 valid pairs for *trial3* and *trial2*. Using all of the valid pairs of data is called pairwise deletion of missing data.

```

1      | trial1 trial2 trial3
2 -----+-----
3      trial1 | 1.0000
4          |      4
5          |
6      trial2 | 0.9939 1.0000
7          |      3      4
8          |
9      trial3 | 0.7001 0.6439 1.0000
10         |      4      4      6

```

It is possible to ask STATA to only perform the correlations on the observations that had complete data for all of the variables on the var statement. For example, you might want the correlations of the reaction times just for the observations that had non-missing data on all of the trials. This is called `listwise` deletion of missing data meaning that when any of the variables are missing, the entire observation is omitted from the analysis. You can request `listwise` deletion within `pwcorr` as illustrated below.

```

1 pwcorr trial1 trial2 trial3, listwise obs
2      | trial1 trial2 trial3
3 -----+-----
4      trial1 | 1.0000
5          |      3
6          |
7      trial2 | 0.9939 1.0000
8          |      3      3
9          |
10     trial3 | 1.0000 0.9939 1.0000
11         |      3      3      3

```

2.4.3 Summary of how missing values are handled in STATA procedures

- **summarize**: For each variable, the number of non-missing values are used.
- **tabulation**: By default, missing values are excluded and percentages are based on the number of non-missing values. If you use the missing option on the **tab** command, the percentages are based on the total number of observations (non-missing and missing) and the percentage of missing values are reported in the table.
- **corr**: By default, correlations are computed based on the number of pairs with non-missing data (**pairwise** deletion of missing data). The **pwcorr** command can be used to request that correlations be computed only for observations that have non-missing data for all variables listed after the **pwcorr** command (**listwise** deletion of missing data).
- **reg**: If any of the variables listed after the **reg** command are missing, the observations missing that value(s) are excluded from the analysis (i.e., **listwise** deletion of missing data).
- For other procedures, see the STATA manual for information on how missing data are handled.

2.4.4 Missing values in assignment statements

It is important to understand how missing values are handled in assignment statements. Consider the example shown below.

```
1 gen sum1 = trial1 + trial2 + trial3
```

The **list** command below illustrates how missing values are handled in assignment statements. The variable *sum1* is based on the variables *trial1*, *trial2* and *trial3*. If any of those variables were missing, the value for *sum1* was set to missing. Therefore *sum1* is missing for observations 2, 3 and 4, as is the case for observation 7.

```
1 list
2      +-----+
3      | id   trial1   trial2   trial3   sum1 |
4      |-----|
5  1. |  1     1.5     1.4     1.6     4.5 |
6  2. |  2     1.5           1.9     . |
7  3. |  3           .     2     1.6     . |
8  4. |  4           .           2.2     . |
9  5. |  5     1.9     2.1           2     6 |
10     |-----|
11  6. |  6     1.8           2     1.9     5.7 |
12  7. |  7           .           .           . |
13     +-----+
```

As a general rule, computations involving missing values yield missing values. For example,

```
1 2 + 2 yields 4
2 2 + . yields .
3 2 / 2 yields 1
4 . / 2 yields .
5 2 * 3 yields 6
6 2 * . yields .
```

whenever you add, subtract, multiply, divide, etc., values that involve missing data, the result is missing.

In our reaction time experiment, the total reaction time *sum1* is missing for four out of seven cases. We could try totaling the data for the non-missing trials by using the **rowtotal** function as shown in the example below.

```
1 egen sum2 = rowtotal(trial1 trial2 trial3)
2 list
```

The results below show that sum2 now contains the sum of the non-missing trials.

1		+-----+							
2			id	trial1	trial2	trial3	sum1	sum2	
3		-----							
4	1.		1	1.5	1.4	1.6	4.5	4.5	
5	2.		2	1.5	.	1.9	.	3.4	
6	3.		3	.	2	1.6	.	3.6	
7	4.		4	.	.	2.2	.	2.2	
8	5.		5	1.9	2.1	2	6	6	
9		-----							
10	6.		6	1.8	2	1.9	5.7	5.7	
11	7.		7	0	
12		+-----+							

Note that the `rowtotal` function treats missing as a zero value. When summing several variables it may not be reasonable to treat missing as zero if an observations is missing on all variables to be summed. The `rowtotal` function with the `missing` option will return a missing value if an observation is missing on all variables.

```
1 egen sum3 = rowtotal(trial1 trial2 trial3) , missing
2
3
4
5
6
7
8
9
10
11
12
13
14
```

	id	trial1	trial2	trial3	sum1	sum2	sum3
1.	1	1.5	1.4	1.6	4.5	4.5	4.5
2.	2	1.5	.	1.9	.	3.4	3.4
3.	3	.	2	1.6	.	3.6	3.6
4.	4	.	.	2.2	.	2.2	2.2
5.	5	1.9	2.1	2	6	6	6
6.	6	1.8	2	1.9	5.7	5.7	5.7
7.	7	0	.

Other statements work similarly. For example, observed what happened when we try to create an average variable without using a function (as in the example below). If any of the variables `trial1`, `trial2` or `trial3` are missing, the value for `avg1` are set to missing.

1	<code>egen avg1 = (trial1 + trial2 + trial3)/3</code>
---	---

Alternatively, the `rowmean` function averages the data for the non-missing trials in the same way as the `rowtotal` function.

1	<code>egen avg2 = rowmean(trial1 trial2 trial3)</code>
---	--

Note: Had there been large number of trials, say 50 trials, then it would be annoying to have to type `avg=rowmean(trial1 trial2 trial3 trial4 ...)`. Here is a shortcut you could use in this kind of situation:

```
1 egen avg3 = rowmean(trial1 - trial3)
2 list
3
4 +-----+
5 | id   trial1  trial2  trial3  avg1  avg2  avg3 |
6 +-----+
7 1. | 1      1.5      1.4      1.6      1.5    1.5    1.5 |
8 2. | 2      1.5      .        1.9      .      1.7    1.7 |
9 3. | 3      .        2        1.6      .      1.8    1.8 |
10 4. | 4      .        .        2.2      .      2.2    2.2 |
```



```

10  5. | 5      1.9      2.1      2      2      2      2 |
11  |-----|
12  6. | 6      1.8      2      1.9      1.9      1.9      1.9 |
13  7. | 7      .      .      .      .      .      . |
14  +-----+

```

Finally, you can use the `rowmiss` and `rownonmiss` functions to determine the number of missing and the number of non-missing values, respectively, in a list of variables. This is illustrated below.

```

1 egen miss = rowmiss(trial1 - trial3)
2 egen nomiss = rownonmiss(trial1 - trial3)
3 list

```

For variable *nomiss*, observations 1, 5 and 6 had three valid values, observations 2 and 3 had two valid values, observation 4 had only one valid value and observation 7 had no valid values. The variable *miss* shows the opposite, it provides a count of the number of missing values.

```

1      +-----+
2      | id   trial1   trial2   trial3   miss   nomiss |
3      |-----|
4  1. | 1     1.5     1.4     1.6     0     3 |
5  2. | 2     1.5     .     1.9     1     2 |
6  3. | 3     .     2     1.6     1     2 |
7  4. | 4     .     .     2.2     2     1 |
8  5. | 5     1.9     2.1     2     0     3 |
9      |-----|
10  6. | 6     1.8     2     1.9     0     3 |
11  7. | 7     .     .     .     3     0 |
12  +-----+

```

2.4.5 Missing values in logical statements

It is important to understand how missing values are handled in logical statements. For example, say that you want to create a 0/1 variable for *trial1* that is 1 if it is 1.5 or less, and 0 if it is over 1.5. We show this below (incorrectly, as you will see).

```

1 gen newvar1 =(trial2 <1.5)
2 list trial2 newvar1

```

It appears that something went wrong with our newly created variable *newvar1*! The observations with missing values for *trial2* were assigned a zero for *newvar1*.

```

1      +-----+
2      | trial2   newvar1 |
3      |-----|
4  1. | 1.4     1 |
5  2. | .     0 |
6  3. | 2     0 |
7  4. | .     0 |
8  5. | 2.1    0 |
9      |-----|
10  6. | 2     0 |
11  7. | .     0 |
12  +-----+

```

Let's explore why this happened by looking at the frequency table of *trial2*.

As you can see in the output, missing values are at the listed after the highest value 2.1 This is because STATA treats a missing value as the largest possible value (e.g., positive infinity) and that value is greater than 2.1, so then the values for *newvar1* become 0.

```

1 tab trial2, missing
2      trial2 |      Freq.      Percent      Cum.
3 -----+-----
4      1.4 |          1      14.29      14.29
5      2 |          2      28.57      42.86
6      2.1 |          1      14.29      57.14
7      . |          3      42.86     100.00
8 -----+-----
9      Total |          7     100.00

```

Now that we understand how STATA treats missing values, we will explicitly exclude missing values to make sure they are treated properly, as shown below.

```

1 gen newvar2 =(trial2 <1.5) if trial2 !=.
2 list trial2 newvar1 newvar2

```

As you can see in the STATA output below, the new variable *newvar2* has missing values for observations that are also missing for *trial2*.

```

1      +-----+
2      | trial2  newvar1  newvar2 |
3      |-----|
4  1. |    1.4        1        1 |
5  2. |     .         0         . |
6  3. |     2         0         0 |
7  4. |     .         0         . |
8  5. |    2.1        0         0 |
9      |-----|
10  6. |     2         0         0 |
11  7. |     .         0         . |
12      +-----+

```

2.4.6 Missing values in logical statements

When creating or recoding variables that involve missing values, always pay attention to whether the variable includes missing values.

2.4.7 For more information

- See the STATA FAQ: [How can I recode missing values into different categories?](#)
- See the STATA FAQ: [Can I quickly see how many missing values a variable has?](#) for more information on examining the number of missing and non-missing values for a particular variable or set of variables.

3 Graphics

3.1 Introduction to Graphs in Stata

This module will introduce some basic graphs in Stata 12, including histograms, boxplots, scatterplots, and scatterplot matrices.

Let's use the auto data file for making some graphs.

```
1 sysuse auto.dta
```

The `histogram` command can be used to make a simple histogram of `mpg`

```
1 histogram mpg
```




Figure 1: `histogram`, the right graph with option `discrete`

If you are creating a histogram for a categorical variable such as `rep78`, you can add the option `discrete`. As you can see below, when you specify this option, the midpoint of each bin labels the respective bar.

```
1 hist rep78, percent discrete
```

The `graph box` command can be used to produce a boxplot which can help you examine the distribution of `mpg`. If `mpg` were normally distributed, the line (the median) would be in the middle of the box (the 25th and 75th percentiles, Q1 and Q3) and the ends of the whiskers (the upper and lower adjacent values, which are the most extreme values within $Q3 + 1.5(Q3 - Q1)$ and $Q1 - 1.5(Q3 - Q1)$, respectively) would be equidistant from the box. The boxplot for `mpg` shows positive skew. The median is pulled to the low end of the box.

```
1 graph box mpg
```




Figure 2: `graph box`, the right graph with option `by`

The boxplot can be done separately for foreign and domestic cars using the `by()` or `over()` option.

```
1 graph box mpg, by(foreign)
```

```
1 graph box mpg, over(foreign)
```

As you can see in the graph above, there are a pair of outliers in the box plots produced. These can be removed from the box plot using the `noout` command in Stata.

```
1 graph box mpg, over(foreign) noout
```

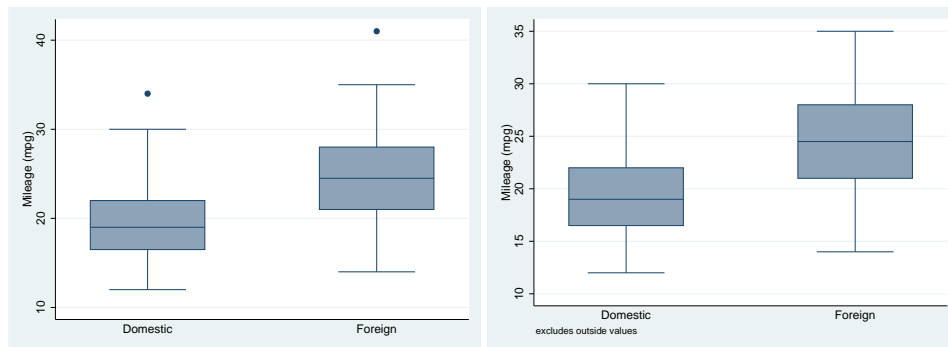



Figure 3: `graph box` with `over`, the right graph with option `noout`

The graph no longer includes the outlying values. Stata also includes a message at the bottom of the graph noting that outside values were excluded.

Stata can also produce pie charts.

```
1 graph pie, over(rep78) plabel(_all name) title("Repair Record 1978")
```



The `graph pie` command with the `over` option creates a pie chart representing the frequency of each group or value of `rep78`. The `plabel` option places the value labels for `rep78` inside each slice of the pie chart.

A two way scatter plot can be used to show the relationship between `mpg` and `weight`. As we would expect, there is a negative relationship between `mpg` and `weight`.

```
1 graph twoway scatter mpg weight
```

Note that you can save typing like this

```
1 twoway scatter mpg weight
```

We can show the regression line predicting `mpg` from `weight` like this.

```
1 twoway lfit mpg weight
```





Figure 4: `twoway` graph, the left graph is `scatter` and the right is `lfit`

We can combine these graphs like shown below.

```
1 twoway (scatter mpg weight) (lfit mpg weight)
```



We can add labels to the points labeling them by make as shown below. Note that `mlabel` is an option on the `scatter` command.

```
1 twoway (scatter mpg weight, mlabel(make) ) (lfit mpg weight)
```

The marker label position can be changed using the `mlabangle()` option.

```
1 twoway (scatter mpg weight, mlabel(make) mlabangle(45)) (lfit mpg weight)
```




Figure 5: `twoway` graph `mlabel`, the right graph with option `mlabangle`

We can combine separate graphs for foreign and domestic cars as shown below, and we have requested confidence bands around the predicted values by using `lfitci` in place of `lfit`. Note that the `by` option is at the end of the command.

```
1 twoway (scatter mpg weight) (lfitci mpg weight), by(foreign)
```

You can request a scatter plot matrix with the `graph matrix` command. Here we examine the relationships among *mpg*, *weight* and *price*.

```
1 graph matrix mpg weight price
```

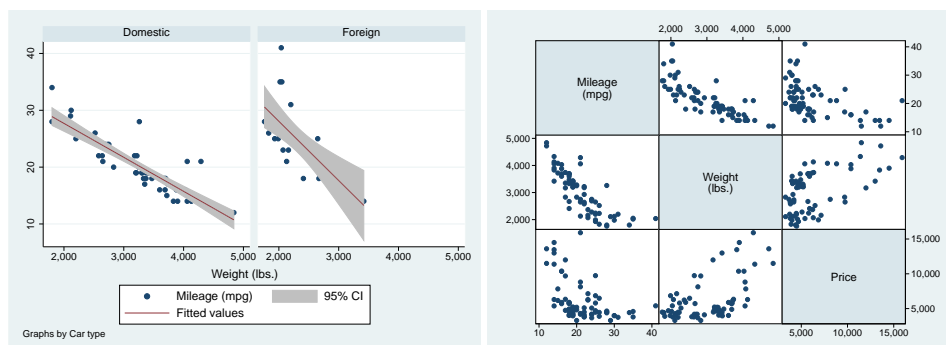



Figure 6: `twoway` graph with `lfitci` and `graph matrix`

3.2 Graphics: Overview of Twoway Plots

This module shows examples of the different kinds of graphs that can be created with the `graph twoway` command. This is illustrated by showing the command and the resulting graph. For more information, see the [Stata Graphics Manual](#) available over the web and from within Stata by typing `help graph`, and in particular the section on [Two Way Scatterplots](#).

Basic twoway scatterplot V.S. Line Plot

```
1 sysuse sp500
2 graph twoway scatter close date // the left graph
3 graph twoway line close date // the right graph
```



Connected Line Plot


```
1 graph twoway connected close date
```

Immediate scatterplot

```
1 graph twoway scatteri ///
2 965.8 15239 (3) "Low 965.8" ///
3 1373.73 15005 (3) "High 1373.73" , msymbol(i)
```

Scatterplot and Immediate Scatterplot

```
1 graph twoway ///
2   (scatter close date) ///
3   (scatteri 965.8 15239 (3) "Low, 9/21, 965.8" ///
4     1373.7 15005 (3) "High, 1/30, 1373.7", msymbol(i) )
```



Left graph is twoway scatter plot; the center is the immediate scatter plot; the right is the combination of the two.

Area Graph

```
1 drop if _n > 57
2 graph twoway area close date, sort
```

Bar plot


```
1 graph twoway bar close date
```

Spike plot


```
1 graph twoway spike close date
```

Dropline plot

```
1 graph twoway dropline close date
```



Notes: The left plot is **area** plot; the right one is **bar** plot.



Notes: The left plot is **spike** plot; the right one is **dropline** plot.

Dot plot

```
1 graph twoway dot change date
```

Range plot with area shading

```
1 graph twoway rarea high low date
```

Range plot with bars

```
1 graph twoway rbar high low date
```



Left graph is `dot` plot; the center is `rarea` plot; the right is `rbar` plot.

Range plot with spikes


```
1 graph twoway rspike high low date
```

Range plot with capped spikes

```
1 graph twoway rcap high low date
```

Range plot with spikes capped with symbols

```
1 graph twoway rcapsym high low date
```



Left graph is `rspike` plot; the center is `rcap` plot; the right is `rcapsym` plot.

Range plot with markers


```
1 graph twoway rscatter high low date
```

Range plot with lines

```
1 graph twoway rline high low date
```

Range plot with lines and markers

```
1 graph twoway rconnected high low date
```

Left graph is `rscatter` plot; the center is `rline` plot; the right is `rconnected` plot.

Median band line plot


```
1 use "http://www.ats.ucla.edu/stat/stata/notes/hsb2.dta", clear
2 graph twoway mband read write
```

Spline line plot

```
1 graph twoway mspline read write
```

LOWESS line plot

```
1 graph twoway lowess read write
```



Left graph is `mband` plot; the center is `mspline` plot; the right is `lowess` plot.

Linear prediction plot


```
1 graph twoway lfit read write
```

Quadratic prediction plot

```
1 graph twoway qfit read write
```

Fractional polynomial plot

```
1 graph twoway fpfit read write
```



Left graph is `lfit` plot; the center is `qfit` plot; the right is `fpfit` plot.

Linear prediction plot with confidence intervals


```
1 graph twoway lfitci read write
```

Quadratic plot with confidence intervals

```
1 graph twoway qfitci read write
```

Fractional polynomial plot with CIs

```
1 graph twoway fptestci read write
```



Left graph is `lfitci` plot; the center is `qfitci` plot; the right is `fptestci` plot.

Histogram


```
1 graph twoway histogram read
```

Kernel density plot

```
1 graph twoway kdensity read
```

Function plot

```
1 graph twoway function y=normalden(x), range(-4 4)
```



Left graph is `histogram` plot; the center is `kdensity` plot; the right is `function` plot.


3.3 Graphics: Twoway Scatterplots

This module shows some of the options when using the `twoway` command to produce scatterplots. This is illustrated by showing the command and the resulting graph. This includes hotlinks to the [Stata Graphics Manual](#) available over the web and from within Stata by typing `help graph`.

3.3.1 Two Way Scatterplots


Basic twoway scatterplot

```
1 twoway (scatter read write)
```





3.3.2 Schemes

```
1 twoway (scatter read write) , scheme(economist) // use the economist scheme (left)
2 twoway (scatter read write) , scheme(s1mono)    // use the s1mono scheme (right)
```



3.3.3 Marker Placement Options (i.e. Jitter)

```
1 twoway (scatter write read, jitter(3)) // with jitter option (left)
2 twoway (scatter write read)           // without jitter (right)
```




3.3.4 Marker Label Options

```
1 // Using small black square symbols.
2 twoway (scatter write read, msymbol(square) msize(small) mcolor(black))
```

```

3 // With markers red on the inside, black medium thick outline
4 twoway (scatter write read, mfcolor(red) mlcolor(black) mlwidth(medthick))


```



```

1 // Identifying Observations with Marker Labels
2 twoway (scatter read write, mlabel(id))
3 // Using large red marker labels at 12 O'clock
4 twoway (scatter read write if id <=10, mlabel(id) mlabposition(12) mlabsize(large)
         mlabcolor(red))


```



```

1 // Markers at 90 degree angle at 12 O'clock with a gap of 5
2 twoway (scatter read write if id <=10, mlabel(ses) mlabangle(90) mlabposition(12) mlabgap
         (5))
3 // If mlabgap option is omitted
4 twoway (scatter read write if id <=10, mlabel(ses) mlabangle(90) mlabposition(12))

```



```


1 // Modifying marker position separately for variables (1)
2 generate pos = 3

```

```

3 replace pos = 1 if (id == 5)
4 replace pos = 5 if (id == 6)
5 replace pos = 9 if (id == 3)
6 twoway (scatter read write if id <= 10, mlabel(ses) mlabv(pos))
7 // If option mlabv is not used
8 twoway (scatter read write if id <= 10, mlabel(ses))

```




3.3.5 Connect Options

```

1 // Connecting with straight line
2 egen mread = mean(read), by(write)
3 twoway (scatter mread write, connect(1) sort)
4 // If the sort option is omitted
5 twoway (scatter mread write, connect(1))
6 // Thick black dotted connecting line
7 twoway (scatter mread write, connect(1) clwidth(thick) clcolor(black) clpattern(dot) sort)


```



```

1 // Show gaps in line when there are missing values
2 egen sdread = sd(read), by(write)
3 twoway (scatter sdread write, connect(1) sort cmissing(n))
4 // Omitting cmissing option
5 twoway (scatter sdread write, connect(1) sort)

```



3.4 Graphics: Combining Twoway Scatterplots

This module shows examples of combining twoway scatterplots. This is illustrated by showing the command and the resulting graph. This includes hotlinks to the [Stata Graphics Manual](#) available over the web and from within Stata by typing `help graph`.

The data set used in these examples can be obtained using the following command:




```
1 use "http://www.ats.ucla.edu/stat/stata/notes/hsb2.dta", clear
```

This illustrates combining graphs in the following situations.

- Plots for separate groups (using `by`)
- Combining separate plots together into a single plot
- Combining separate graphs together into a single graph


3.4.1 Plots for separate groups

```
1 // Separate graphs by gender (male and female)
2 twoway (scatter read write), by(female)
3 // Separate graphs by ses and gender
4 twoway (scatter read write), by(female ses)
5 // Swapping position of ses and gender
6 twoway (scatter read write), by(ses female, cols(2))
```



3.4.2 Combining scatterplots and linear fit in one graph

```
1 // Scatterplot with linear fit
2 twoway (scatter read write) ///
3       (lfit read write) ,
4       ytitle(Reading Score)
5 // Graphs separated by SES and female with linear fit lines and points identified by id
6 twoway (scatter read write, mlabel(id)) ///
7       (lfit read write, range(30 70)) ,
8       ytitle(Reading Score) by(ses female)
9 // Graph for high ses females with linear fit with and without obs 51
10 twoway (scatter read write, mlabel(id)) ///
11       (lfit read write, range(30 70)) ///
12       (lfit read write if id != 51, range(30 70)) if female==1 & ses==3,
13       ytitle(Reading Score) legend(lab(3 "Fitted values without Obs 51"))
```

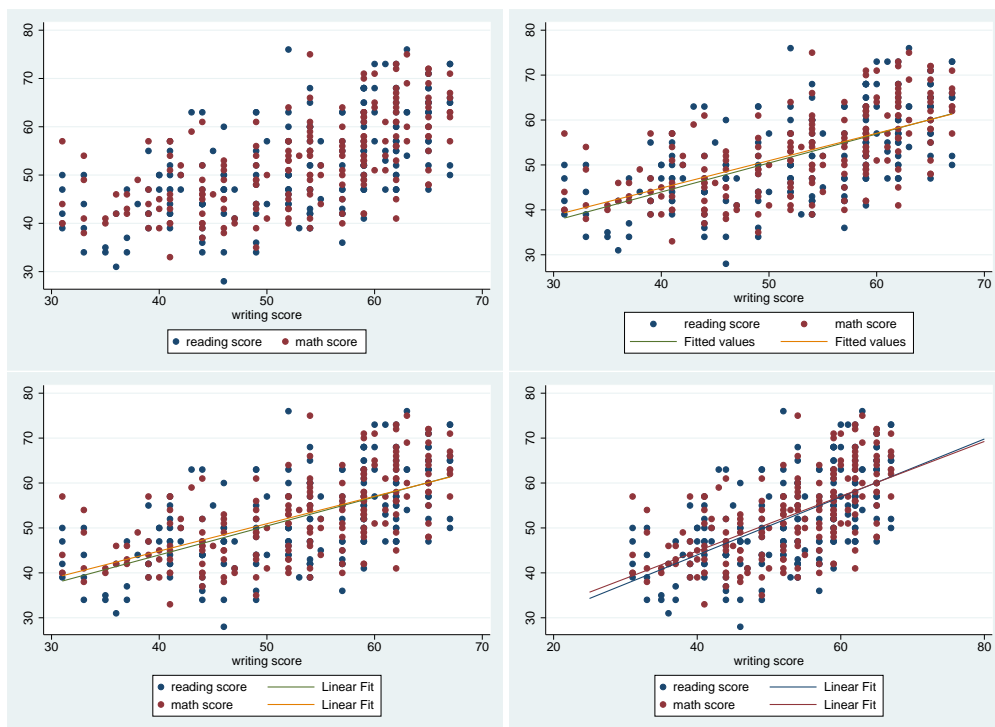


3.4.3 Combining scatterplots with multiple variables and linear fits

```

1 // Reading and math score by writing score
2 twoway (scatter read write) /// plot(1,1)
3       (scatter math write)
4 // Reading and math score by writing score with fit lines
5 twoway (scatter read write) (scatter math write) /// plot(1,2)
6       (lfit read write) (lfit math write) ///
7 // Adding legend to above graph
8 twoway (scatter read write) /// plot(2,1)
9       (scatter math write) ///
10      (lfit read write) ///
11      (lfit math write), ///
12      legend(label(3 "Linear Fit") label(4 "Linear Fit")) ///
13      legend(order(1 3 2 4))
14 // Final version of graph making line style same as dot style, and ranges the same
15 twoway (scatter read write) /// plot(2,2)
16       (scatter math write) ///
17       (lfit read write, pstyle(p1) range(25 80) ) ///
18       (lfit math write, pstyle(p2) range(25 80) ), ///
19       legend(label(3 "Linear Fit") label(4 "Linear Fit")) ///
20       legend(order(1 3 2 4))

```




3.4.4 Combining scatterplots and linear fit for separate groups

```

1 // Overlay graph of males and females in one graph
2 separate write, by(female)
3 twoway (scatter write0 read) (scatter write1 read), ///
4         ytitle(Writing Score) legend(order(1 "Males" 2 "Females"))
5 // Overlay graph of males and females in one graph with linear fit lines
6 twoway (scatter write0 read) (scatter write1 read) ///
7         (lfit write0 read) (lfit write1 read), ///
8         ytitle(Writing Score) ///
9         legend(order(1 "Males" 2 "Females" 3 "Lfit Males" 4 "Lfit Females"))

```



3.4.5 Combining separate graphs into one graph

First, we make 3 graphs (not shown)

```

1 // Making the Graphs
2 twoway (scatter read write) (lfit read write), name(scatterx)
3 regress read write
4 rvfplot, name(rvf)
5 lvr2plot, name(lvr)

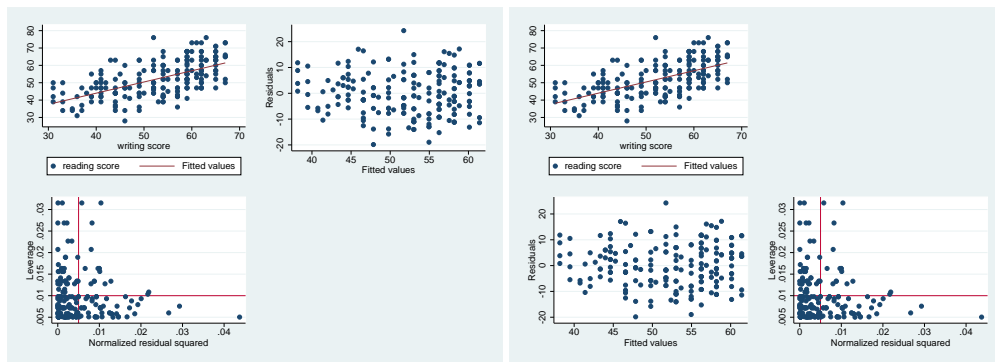
```

Now we can use `graph combine` to combine these into one graph, we can also move the place where the empty graph is located, as shown below.

```

1 // combine the three graphs into one graph
2 graph combine scatterx rvf lvr
3 // Combining the graphs differently
4 graph combine scatterx rvf lvr, hole(2)


```



3.5 Graphics: Common Graph Options

This module shows examples of the different kinds of graphs that can be created with the `graph twoway` command. This is illustrated by showing the command and the resulting graph. For more information, see the [Stata Graphics Manual](#) available over the web and from within Stata by typing `help graph`, and in particular the section on *Two Way Scatterplots*.

```
1 // Adding a title
2 graph twoway scatter read write, ///
3   title("Scatterplot of Reading and Writing")
4 // Black title, positioned at 11 O'Clock
5 graph twoway scatter read write, ///
6   title("Scatterplot of Reading and Writing", ///
7     color(black) position(11))
8 // Title at 5 O'Clock, medium size text, positioned within the graph
9 graph twoway scatter read write, ///
10  title("Scatterplot of Reading and Writing", ///
11    size(medium) position(5) ring(0))
```



```
1 // Title in a box with cyan background, magenta border and a medium margin around the
   title
2 graph twoway scatter read write, ///
3   title("Scatterplot of Reading and Writing", ///
4     box bcolor(cyan) blcolor(magenta) bmargin(medium))
5 // Two line title with a gap of 3 between the titles
6 graph twoway scatter read write, ///
7   title("Scatterplot of Reading and Writing" ///
8     "Sample of 200 Students", linegap(3) )
```





```
1 // Graph with title, subtitle, caption, and a note
2 graph twoway scatter read write, ///
3   title("Scatterplot of Reading and Writing") ///
4   subtitle("Sample of 200 Students") ///
```

```

5  note(High School and Beyond Data) ///
6  caption(From www.ats.ucla.edu)
7  // Moving and sizing note and caption
8  graph twoway scatter read write, ///
9  title("Scatterplot of Reading and Writing") ///
10 subtitle("Sample of 200 Students") ///
11 note(High School and Beyond Data, size(medium) position(5)) ///
12 caption(From www.ats.ucla.edu, size(vsmall) position(5))



```



```

1 // Modifying title on x and y axis
2 twoway scatter read write, ///
3 ytitle(Score on Writing Test) ///
4 xtitle(Score on Reading Test)
5
6 // Complete example
7 twoway scatter read write, ///
8 title("Scatterplot of Reading and Writing") ///
9 subtitle("Sample of 200 Students") ///
10 note(High School and Beyond Data, size(medium) position(5)) ///
11 caption(From www.ats.ucla.edu, size(vsmall) position(5)) ///
12 ytitle(Score on Writing Test) ///
13 xtitle(Score on Reading Test)



```



```

1 // Sizing graph to have 4 by 2 aspect ratio
2 twoway scatter read write, ysize(2) xsize(4)
3
4 // Making text scaled 1.5 times normal size
5 graph twoway scatter read write, scale(1.5)




```



```

1 // Graph with sand color outside graph, gray inside graph
2 graph twoway scatter read write, ///
3   title("Scatterplot of Reading and Writing") ///
4   graphregion( color(sand) ) plotregion( fcolor(gray) )
5
6 // Graph with sand color outside graph, gray inside graph, red outer border, blue inner
  border
7 graph twoway scatter read write, ///
8   title("Scatterplot of Reading and Writing") ///
9   graphregion( fcolor(red) ifcolor(sand) ) ///
10  plotregion( fcolor(blue) ifcolor(gray))
11
12 // Graph with colors for many border elements
13 graph twoway scatter read write, ///
14   title("Scatterplot of Reading and Writing") ///
15   graphregion( fcolor(red)   lcolor(yellow)   lwidth(thick)   ///
16               ifcolor(sand) ilcolor(orange) ilwidth(thick)) ///
17   plotregion( fcolor(blue)  lcolor(green)   lwidth(thick)   ///
18               ifcolor(gray) ilcolor(purple) ilwidth(thick))

```



4 Reading Data in Stata

4.1 Using and saving files in Stata

4.1.1 Using and saving Stata data files

The `use` command gets a Stata data file from disk and places it in memory so you can analyze and/or modify it. A data file must be read into memory before you can analyze it. It is kind of like when you open a *Word* document; you need to read a *Word* document into *Word* before you can work with it. The `use` command below gets the Stata data file called `auto.dta` from disk and places it in memory so we can analyze and/or modify it. Since Stata data files end with `.dta` you need only say `use auto` and Stata knows to read in the file called `auto.dta`.

```
1 use auto
```

The `describe` command tells you information about the data that is currently sitting in memory.

```
1 describe
2
3 Contains data from auto.dta
4   obs:           74
5   vars:           12                      17 Feb 1999 10:49
6   size:           3,108 (99.6% of memory free)
7 -----
8   1. make         str17   %17s
9   2. price        int     %9.0g
10  3. mpg          byte    %9.0g
11  4. rep78        byte    %9.0g
12  5. hdroom       float   %9.0g
13  6. trunk        byte    %9.0g
14  7. weight       int     %9.0g
15  8. length       int     %9.0g
16  9. turn         byte    %9.0g
17 10. displ       int     %9.0g
18 11. gratio      float   %9.0g
19 12. foreign      byte    %9.0g
20 -----
21 Sorted by:
```

Now that the data is in memory, we can analyze it. For example, the `summarize` command gives summary statistics for the data currently in memory.

```
1 summarize
2
3 Variable |      Obs      Mean   Std. Dev.      Min      Max
4 -----+-----
5   make |         0
6   price |       74   6165.257   2949.496     3291   15906
7   mpg  |       74    21.2973    5.785503      12      41
8   rep78 |       69    3.405797    .9899323       1       5
9   hdroom |       74    2.993243    .8459948     1.5       5
10  trunk |       74    13.75676    4.277404       5      23
11  weight |       74   3019.459    777.1936    1760    4840
12  length |       74   187.9324   22.26634     142     233
13   turn |       74    39.64865    4.399354      31     51
14  displ |       74   197.2973   91.83722      79     425
15  gratio |       74    3.014865    .4562871     2.19     3.89
16  foreign |       74    .2972973    .4601885       0       1
```

Let's make a change to the data in memory. We will compute a variable called *price2* which will be double the value of *price*.

```
1 generate price2 = 2*price
```

If we use the `describe` command again, we see the variable we just created is part of the data in memory. We also see a note from Stata saying dataset has changed since last saved. Stata knows that the data in memory has changed, and would need to be saved to avoid losing the changes. It is like when you are editing a *Word* document; if you don't save the data, any changes you make will be lost. If we shut the computer off before saving the changes, the changes we made would be lost.

```
1 describe
2 Contains data from auto.dta
3   obs:           74
4   vars:           13                               17 Feb 1999 10:49
5   size:           3,404 (99.6% of memory free)
6 -----
7   1. make         str17   %17s
8   2. price        int     %9.0g
9   3. mpg          byte    %9.0g
10  4. rep78        byte    %9.0g
11  5. hdroom       float   %9.0g
12  6. trunk        byte    %9.0g
13  7. weight       int     %9.0g
14  8. length       int     %9.0g
15  9. turn         byte    %9.0g
16 10. displ        int     %9.0g
17 11. gratio       float   %9.0g
18 12. foreign      byte    %9.0g
19 13. price2       float   %9.0g
20 -----
21 Sorted by:
22   Note: dataset has changed since last saved
```

The `save` command is used to save the data in memory permanently on disk. Let's save this data and call it `auto2` (Stata will save it as `auto2.dta`).

```
1 save auto2
2
3 file auto2.dta saved
```

Let's make another change to the dataset. We will compute a variable called *price3* which will be three times the value of *price*.

```
1 generate price3 = 3*price
```

Let's try to save this data again to `auto2`

```
1 save auto2
2 file auto2.dta already exists
3 r(602);
```

Did you see how Stata said file `auto2.dta` already exists? Stata is worried that you will accidentally overwrite your data file. You need to use the `replace` option to tell Stata that you know that the file exists and you want to replace it.

```

1 save auto2, replace
2
3 file auto2.dta saved

```

Let's make another change to the data in memory by creating a variable called *price4* that is four times the *price*.

```

1 generate price4 = price*4

```

Suppose we want to use the original auto file and we don't care if we lose the changes we just made in memory (i.e., losing the variable *price4*). We can try to use the auto file.

```

1 sysuse auto
2
3 no; data in memory would be lost
4 r(4);

```

See how Stata refused to *use* the file, saying no; *data in memory* would be lost? Stata did not want you to lose the changes that you made to the data sitting in memory. If you really want to discard the changes in memory, then you need to use the *clear* option on the *use* command, as shown below.

```

1 sysuse auto, clear

```

Stata tries to protect you from losing your data by doing the following:

1. If you want to save a file over an existing file, you need to use the *replace* option, e.g., *save auto, replace*.
2. If you try to *use* a file and the file in memory has unsaved changes, you need to use the *clear* option to tell Stata that you want to discard the changes, e.g., *use auto, clear*.

Before we move on to the next topic, let's clear out the data in memory.

```

1 clear

```

4.1.2 Using files larger than 1 megabyte

When you use a data file, Stata reads the entire file into memory. By default, Stata limits the size of data in memory to 1 megabyte (PC version 6.0 Intercooled). You can view the amount of memory that Stata has reserved for data with the *memory* command.

```

1 memory
2
3 Total memory                1,048,576 bytes    100.00%
4
5 overhead (pointers)         0                0.00%
6 data                        0                0.00%
7 -----
8 data + overhead              0                0.00%
9
10 programs, saved results, etc. 1,152            0.11%
11 -----
12 Total                       1,152            0.11%
13
14 Free                       1,047,424          99.89%

```

If you try to *use* a file which exceeds the amount of memory Stata has allocated for data, it will give you an error message like this.

```

1 no room to add more observations
2 r(901);

```

You can increase the amount of memory that Stata has allocated to data using the `set memory` command. For example, if you had a data file which was 1.5 megabytes, you can set the memory to, say, 2 megabytes shown below.

```
1 set memory 2m
2
3 (2048k)
```

Once you have increased the memory, you should be able to `use` the data file if you have allocated enough memory for it.

4.1.3 Summary

- `sysuse auto`: To `use` the auto file from disk and read it into memory
- `save auto`: To save the file auto from memory to disk
- `save auto, replace`: To save a file if the file auto already exists
- `sysuse auto, clear`: to use a file auto and `clear` out the current data in memory
- `clear`: If you want to `clear` out the data in memory, you want to lose the changes
- `set memory 2m`: To allocate 2 megabytes of memory for a data file.
- `memory`: To view the allocation of memory to data and how much is used.

4.2 Inputting your data into Stata

This module will show how to input your data into Stata. This covers inputting data with comma delimited, tab delimited, space delimited, and fixed column data.

Note: all of the sample input files for this page were created by us and are not included with Stata. You can create them yourself to try out this code by copying and pasting the data into a text file.

4.2.1 Typing data into the Stata editor

One of the easiest methods for getting data into Stata is using the Stata data editor, which resembles an Excel spreadsheet. It is useful when your data is on paper and needs to be typed in, or if your data is already typed into an Excel spreadsheet. To learn more about the Stata data editor, see the `edit` module.

4.2.2 Comma/tab separated file with variable names on line 1

Two common file formats for raw data are **comma separated** files and **tab separated** files. Such files are commonly made from spreadsheet programs like *Excel*. Consider the **comma delimited** file shown below.

```
1 type auto2.raw
2 make, mpg, weight, price
3 AMC Concord, 22, 2930, 4099
4 AMC Pacer, 17, 3350, 4749
5 AMC Spirit, 22, 2640, 3799
6 Buick Century, 20, 3250, 4816
7 Buick Electra, 15,4080, 7827
```

This file has two characteristics:

- The first line has the names of the variables separated by commas,
- The following lines have the values for the variables, also separated by commas.

This kind of file can be read using the `insheet` command, as shown below.

```
1 insheet using auto2.raw
2
```

```
3 (4 vars, 5 obs)
```

We can check to see if the data came in right using the `list` command.

```
1 list
2           make      mpg   weight   price
3  1.   AMC Concord    22    2930    4099
4  2.   AMC Pacer     17    3350    4749
5  3.   AMC Spirit    22    2640    3799
6  4. Buick Century    20    3250    4816
7  5. Buick Electra   15    4080    7827
```

Since you will likely have more observations, you can use `in` to list just a subset of observations. Below, we `list` observations 1 through 3.

```
1 list in 1/3
2           make      mpg   weight   price
3  1.   AMC Concord    22    2930    4099
4  2.   AMC Pacer     17    3350    4749
5  3.   AMC Spirit    22    2640    3799
```

Now that the file has been read into Stata, you can save it with the `save` command (we will skip doing that step).

The exact same `insheet` command could be used to read a `tab` delimited file. The `insheet` command is clever because it can figure out whether you have a comma delimited or `tab` delimited file, and then read it. (However, `insheet` could not handle a file that uses a mixture of commas and tabs as delimiters.)

Before starting the next section, let's clear out the existing data in memory.

```
1 clear
```

4.2.3 Comma/tab separated file (no variable names in file)

Consider a file that is identical to the one we examined in the previous section, but it does not have the variable names on line 1

```
1 type auto3.raw
2 AMC Concord, 22, 2930, 4099
3 AMC Pacer, 17, 3350, 4749
4 AMC Spirit, 22, 2640, 3799
5 Buick Century, 20, 3250, 4816
6 Buick Electra, 15,4080, 7827
```

This file can be read using the `insheet` command as shown below.

```
1 insheet using auto3.raw
2 (4 vars, 5 obs)
```

But where did Stata get the variable names? If Stata does not have names for the variables, it names them `v1`, `v2`, `v3` etc., as you can see below.

```
1 list
2
3           v1      v2      v3      v4
4  1.   AMC Concord    22    2930    4099
5  2.   AMC Pacer     17    3350    4749
6  3.   AMC Spirit    22    2640    3799
7  4. Buick Century    20    3250    4816
8  5. Buick Electra   15    4080    7827
```


Let's clear out the data in memory, and then try reading the data again.

```
1 clear
```

Now, let's try reading the data and tell Stata the names of the variables on the `insheet` command.

```
1 insheet make mpg weight price using auto3.raw
2 (4 vars, 5 obs)
```

As the list command shows, Stata used the variable names supplied on the `insheet` command.

```
1 list
2
3           make      mpg    weight    price
4  1.   AMC Concord      22     2930     4099
5  2.     AMC Pacer      17     3350     4749
6  3.     AMC Spirit      22     2640     3799
7  4.   Buick Century      20     3250     4816
8  5.   Buick Electra      15     4080     7827
```

The `insheet` command works equally well on files which use tabs as separators. Stata examines the file and determines whether commas or tabs are being used as separators and reads the file appropriately.

Now that the file has been read into Stata, you can save it with the `save` command (we will skip doing that step).

Let's clear out the data in memory before going to the next section.

```
1 clear
```

4.2.4 Space separated file

Consider a file where the variables are separated by spaces like the one shown below.

```
1 type auto4.raw
2 "AMC Concord" 22 2930 4099
3 "AMC Pacer" 17 3350 4749
4 "AMC Spirit" 22 2640 3799
5 "Buick Century" 20 3250 4816
6 "Buick Electra" 15 4080 7827
```

Note that the make of car is contained within quotation marks. This is necessary because the names contain spaces within them. Without the quotes, Stata would think AMC is the *make* and Concord is the *mpg*. If the *make* did not have spaces embedded within them, the quotation marks would not be needed.

This file can be read with the `infile` command as shown below.

```
1 infile str13 make mpg weight price using auto4.raw
2 (5 observations read)
```

You may be asking yourself, where did the `str13` come from? Since *make* is a character variable, we need to tell Stata that it is a character variable, and how long it can be. The `str13` tells Stata it is a **string** variable and that it could be up to 13 characters wide.

The `list` command confirms that the data was read correctly.

```
1 list
2           make      mpg    weight    price
3  1.   AMC Concord      22     2930     4099
4  2.     AMC Pacer      17     3350     4749
5  3.     AMC Spirit      22     2640     3799
6  4.   Buick Century      20     3250     4816
7  5.   Buick Electra      15     4080     7827
```

Now that the file has been read into Stata, you can save it with the `save` command (we will skip doing that step). Let's clear out the data in memory before moving on to the next section.

```
1 clear
```

4.2.5 Fixed format file

Consider a file using fixed column data like the one shown below.

```
1 type auto5.raw
2 AMC Concord    22 2930 4099
3 AMC Pacer      17 3350 4749
4 AMC Spirit     22 2640 3799
5 Buick Century  20 3250 4816
6 Buick Electra  15 4080 7827
```

Note that the variables are clearly defined by which column(s) they are located. Also, note that the *make* of car is not contained within quotation marks. The quotations are not needed because the columns define where the *make* begins and ends, and the embedded spaces no longer create confusion.

This file can be read with the `infix` command as shown below.

```
1 infix str make 1-13 mpg 15-16 weight 18-21 price 23-26 using auto5.raw
2
3 (5 observations read)
```

Here again we need to tell Stata that *make* is a **string** variable by preceding *make* with `str`. We did not need to indicate the length since Stata can infer that *make* can be up to 13 characters wide based on the column locations.

The `list` command confirms that the data was read correctly.

```
1 list
2
3      make      mpg      weight      price
4 1.  AMC Concord      22      2930      4099
5 2.   AMC Pacer      17      3350      4749
6 3.   AMC Spirit      22      2640      3799
7 4. Buick Century      20      3250      4816
8 5. Buick Electra      15      4080      7827
```

Now that the file has been read into Stata, you can save it with the `save` command (we will skip doing that step). Let's clear out the data in memory before moving on to the next section.

```
1 clear
```

4.2.6 Other methods of getting data into Stata

This does not cover all possible methods of getting raw data into Stata, but does cover many common situations. See the Stata Users Guide for more comprehensive information on reading raw data into Stata.

Another method that should be mentioned is the use of data conversion programs. These programs can convert data from one file format into another file format. For example, they could directly create a Stata file from an Excel Spreadsheet, a Lotus Spreadsheet, an Access database, a Dbase database, a SAS data file, an SPSS system file, etc. Two such examples are Stat Transfer and DBMS Copy. Both of these products are available on SSC PCs and DBMS Copy is available on Nicco and Aristotle.

Finally, if you are using Nicco, Aristotle or the RS/6000 Cluster, there is a command specifically for converting SAS data into Stata called `sas2stata`. If you have SAS data you want to convert to Stata, this may be a useful way to get your SAS data into Stata.

4.2.7 Summary

- **edit**: Bring up the Stata data editor for typing data in.
- **insheet using** auto2.raw, **clear**: Read in the comma or tab delimited file called auto2.raw taking the variable names from the first line of data.
- **insheet** make mpg weight price **using** auto3.raw, **clear**: Read in the comma or tab delimited file called auto3.raw naming the variables mpg weight and price.
- **infile str13** make mpg weight price **using** auto4.raw, **clear**: Read in the space separated file named auto4.raw. The variable make is surrounded by quotes because it has embedded blanks.
- **infix str** make 1-13 mpg 15-16 weight 18-21 **using** auto5.raw, **clear**: Read in the fixed format file named auto5.raw.
- **DBMS/Copy, Stat Transfer, sas2stata, and Stata Users Guide.**: Other methods

4.3 Using dates in Stata

This module will show how to use date variables, date functions, and date display formats in Stata.

4.3.1 Converting dates from raw data using the `date()` function

The trick to inputting dates in Stata is to forget they are dates, and treat them as character strings, and then later convert them into a Stata date variable. You might have the following date data in your raw data file.

```
1 type dates1.raw
2 John 1 Jan 1960
3 Mary 11 Jul 1955
4 Kate 12 Nov 1962
5 Mark 8 Jun 1959
```

You can read these data by typing:

```
1 infix str name 1-4 str bday 6-17 using dates1.raw
2 (4 observations read)
```

Using the `list` command, you can see that the date information has been read correctly into `bday`.

```
1 list
2           name           bday
3 1.      John      1 Jan 1960
4 2.      Mary     11 Jul 1955
5 3.      Kate     12 Nov 1962
6 4.      Mark      8 Jun 1959
```

Since `bday` is a string variable, you cannot do any kind of date computations with it until you make a date variable from it. You can generate a date version of `bday` using the `date()` function. The example below creates a date variable called `birthday` from the character variable `bday`. The syntax is slightly different depending on which version of Stata you are using. The difference is in how the pattern is specified. In Stata 9 it should be lower case (e.g., “dmy”) and in Stata 10, it should be upper case for day, month, and year (e.g., “DMY”) but lower case if you want to specify hours, minutes or seconds (e.g., “DMYhms”). Our data are in the order day, month, year, so we “useDMY” (“ordmy” if you are using Stata 9) within the `date()` command. (Unless otherwise noted, all other Stata commands on this page are the same for versions 9 and 10.)

In Stata **version 9**:

```
1 generate birthday=date(bday, "dmy")
```

In Stata **version 10**:

```
1 generate birthday=date(bday, "DMY")
```

Let's have a look at both bday and birthday.

```
1 list
2
3      name      bday      birthday
4  1.   John   1 Jan 1960           0
5  2.   Mary  11 Jul 1955        -1635
6  3.   Kate  12 Nov 1962        1046
7  4.   Mark   8 Jun 1959        -207
```

The values for *birthday* may seem confusing. The value of *birthday* for John is 0 and the value of *birthday* for Mark is -207. Dates are actually stored as **the number of days from Jan 1, 1960** which is convenient for the computer storing and performing date computations, but is difficult for you and I to read.

We can tell Stata that *birthday* should be displayed using the %d format to make it easier for humans to read.

```
1 format birthday %d
2 list
3      name      bday      birthday
4  1.   John   1 Jan 1960   01jan1960
5  2.   Mary  11 Jul 1955  11jul1955
6  3.   Kate  12 Nov 1962  12nov1962
7  4.   Mark   8 Jun 1959   08jun1959
```

The `date()` function is very flexible and can handle dates written in almost any manner. For example, consider the file `dates2.raw`.

```
1 type dates2.raw
2 John Jan 1 1960
3 Mary 07/11/1955
4 Kate 11.12.1962
5 Mark Jun/8 1959
```

These dates are messy, but they are consistent. Even though the formats look different, it is always a month day year separated by a delimiter (e.g., space slash dot or dash). We can try using the syntax from above to read in our new dates. Note that, as discussed above, for Stata version 10 the order of the date is declared in upper case letters (i.e., "MDY") while for version 9 it is declared in all lower case (i.e., "mdy").

```
1 clear
2 infix str name 1-4 str bday 6-17 using dates2.raw
3
4 (4 observations read)
5
6 generate birthday=date(bday, "MDY")
7
8 format birthday %d
9 list
10
11      name      bday      birthday
12  1.   John   Jan 1 1960   01jan1960
13  2.   Mary  07/11/1955  11jul1955
14  3.   Kate  11.12.1962  12nov1962
15  4.   Mark  Jun/8 1959   08jun1959
```

Stata was able to read those dates without a problem. Let's try an even tougher set of dates. For example, consider the dates in `dates3.raw`.

```

1 type dates3.raw
2 4-12-1990
3 4.12.1990
4 Apr 12, 1990
5 Apr12,1990
6 April 12, 1990
7 4/12.1990
8 Apr121990

```

Let's try reading these dates and see how Stata handles them. Again, remember that for Stata version 10 dates are declared "MDY" while for version 9 they are declared "mdy".

```

1 clear
2 infix str bday 1-20 using dates3.raw
3 (7 observations read)
4 generate birthday=date(bday,"MDY")
5 (1 missing value generated)
6 format birthday %d
7 list

```

	bday	birthday
1.	4-12-1990	12apr1990
2.	4.12.1990	12apr1990
3.	Apr 12, 1990	12apr1990
4.	Apr12,1990	12apr1990
5.	April 12, 1990	12apr1990
6.	4/12.1990	12apr1990
7.	Apr121990	.

As you can see, Stata was able to handle almost all of those crazy date formats. It was able to handle Apr12,1990 even though there was not a delimiter between the month and day (Stata was able to figure it out since the month was character and the day was a number). The only date that did not work was Apr121990 and that is because there was no delimiter between the day and year. As you can see, the `date()` function can handle just about any date as long as there are delimiters separating the month day and year. In certain cases Stata can read all numeric dates entered without delimiters, see [help dates](#) for more information.

4.3.2 Converting dates from raw data using the `mdy()` function

In some cases, you may have the month, day, and year stored as numeric variables in a dataset. For example, you may have the following data for birth dates from `dates4.raw`.

```

1 type dates4.raw
2 7 11 1948
3 1 1 1960
4 10 15 1970
5 12 10 1971

```

You can read in this data using the following syntax to create a separate variable for month, day and year.

```

1 clear
2 infix month 1-2 day 4-5 year 7-10 using dates4.raw
3 (4 observations read)
4 list

```

	month	day	year
1.	7	11	1948
2.	1	1	1960

8	3.	10	15	1970
9	4.	12	10	1971

A Stata date variable can be created using the `mdy()` function as shown below.

```
1 generate birthday=mdy(month,day,year)
```

Let's format *birthday* using the `%d` format so it displays better.

```
1 format birthday %d
2 list
```

3		month	day	year	birthday
4	1.	7	11	1948	11jul1948
5	2.	1	1	1960	01jan1960
6	3.	10	15	1970	15oct1970
7	4.	12	10	1971	10dec1971

Consider the data in `dates5.raw`, which is the same as `dates4.raw` except that only two digits are used to signify the year.

```
1 type dates5.raw
2 7 11 48
3 1 1 60
4 10 15 70
5 12 10 71
```

Let's try reading these dates just like we read `dates4.raw`.

```
1 clear
2 infix month 1-2 day 4-5 year 7-10 using dates5.raw
3 (4 observations read)
4 generate birthday=mdy(month,day,year)
5 (4 missing values generated)
6 format birthday %d
7 list
```

8		month	day	year	birthday
9	1.	7	11	48	.
10	2.	1	1	60	.
11	3.	10	15	70	.
12	4.	12	10	71	.

As you can see, the values for *birthday* are all missing. This is because Stata assumes that the years were literally 48, 60, 70 and 71 (it does not assume they are 1948, 1960, 1970 and 1971). You can force Stata to assume the century portion is 1900 by adding 1900 to the year as shown below (note that we use `replace` instead of `generate` since the variable *birthday* already exists).

```
1 replace birthday=mdy(month,day,year+1900)
2 (4 real changes made)
3 format birthday %d
4 list
```

5		month	day	year	birthday
6	1.	7	11	48	11jul1948
7	2.	1	1	60	01jan1960
8	3.	10	15	70	15oct1970
9	4.	12	10	71	10dec1971

4.3.3 Computations with elapsed dates

Date variables make computations involving dates very convenient. For example, to calculate everyone's age on January 1, 2000 simply use the following conversion.

```
1 generate age2000=( mdy(1,1,2000) - birthday ) / 365.25
2 list
3      month      day      year  birthday  age2000
4  1.         7      11      48  11jul1948  51.47433
5  2.         1       1      60  01jan1960   40
6  3.        10      15      70  15oct1970  29.21287
7  4.        12      10      71  10dec1971  28.06023
```

Please note that this formula for age does not work well over very short time spans. For example, the age for a child on their his birthday will be less than one due to using 365.25. There are formulas that are more exact but also much more complex. Here is an example courtesy of Dan Blanchette.

```
1 generate altage = floor([ym(2000, 1) - ym(year(birthday), month(birthday))] - [1 < day(
    birthday)]) / 12)
```

4.3.4 Other date functions

Given a date variable, one can have the month, day and year returned separately if desired, using the `month()`, `day()` and `year()` functions, respectively.

```
1 generate m=month(birthday)
2 generate d=day(birthday)
3 generate y=year(birthday)
4 list m d y birthday
5      m      d      y  birthday
6  1.     7     11    1948  11jul1948
7  2.     1      1    1960  01jan1960
8  3.    10     15    1970  15oct1970
9  4.    12     10    1971  10dec1971
```

If you'd like to return the **day of the week** for a date variable, use the `dow()` function (where 0=Sunday, 1=Monday etc.).

```
1 gen week_d=dow(birthday)
2 list birthday week_d
3      birthday  week_d
4  1. 11jul1948      0
5  2. 01jan1960      5
6  3. 15oct1970      4
7  4. 10dec1971      5
```

4.3.5 Summary

The `date()` function converts strings containing dates to date variables. The syntax varies slightly by version.

- In Stata version 9:

```
1 gen date2 = date(date, "dmy")
```

- In Stata version 10:

```
1 gen date2 = date(date, "DMY")
```

- The `mdy()` function takes three numeric arguments (month, day, year) and converts them to a date variable.

```
1 generate birthday=mdy(month,day,year)
```

- You can display elapsed times as actual dates with display formats such as the `%d` format.

```
1 format birthday %d
```

Other date functions include the `month()`, `day()`, `year()`, and `dow()` functions. For online help with dates, type `help dates` at the command line. For more detailed explanations about how Stata handles dates and date functions, please refer to the Stata Users Guide.

5 Basic Data Management in Stata

5.1 Labeling data

This module will show how to create labels for your data. Stata allows you to label your data file (**data label**), to label the variables within your data file (variable labels), and to label the values for your variables (**value labels**). Let's use a file called `autolab.dta` that does not have any labels.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/autolab.dta", clear
```

Let's use the `describe` command to verify that indeed this file does not have any labels.

```
1 describe
2 Contains data from autolab.dta
3 obs:          74                      1978 Automobile Data
4 vars:          12                     23 Oct 2008 13:36
5 size:         3,478 (99.9% of memory free) (_dta has notes)
6 -----
7           storage  display      value
8 variable name  type   format    label   variable label
9 -----
10 make          str18  %-18s
11 price          int    %8.0gc
12 mpg            int    %8.0g
13 rep78          int    %8.0g
14 headroom       float  %6.1f
15 trunk          int    %8.0g
16 weight         int    %8.0gc
17 length         int    %8.0g
18 turn           int    %8.0g
19 displacement   int    %8.0g
20 gear_ratio     float  %6.2f
21 foreign        byte   %8.0g
22 -----
23 Sorted by:
```

Let's use the `label data` command to add a label describing the data file. This label can be up to 80 characters long.

```
1 label data "This file contains auto data for the year 1978"
```

The `describe` command shows that this label has been applied to the version that is currently in memory.

```
1 describe
2 Contains data from autolab.dta
3 obs:          74                      This file contains auto data for the year
   1978
4 vars:          12                     23 Oct 2008 13:36
5 size:         3,478 (99.9% of memory free) (_dta has notes)
6 -----
7           storage  display      value
8 variable name  type   format    label   variable label
9 -----
10 make          str18  %-18s
11 price          int    %8.0gc
12 mpg            int    %8.0g
13 rep78          int    %8.0g
14 headroom       float  %6.1f
```

```

15 trunk          int      %8.0g
16 weight         int      %8.0gc
17 length         int      %8.0g
18 turn           int      %8.0g
19 displacement    int      %8.0g
20 gear_ratio      float    %6.2f
21 foreign         byte     %8.0g

```

```

22 -----
23 Sorted by:

```

Let's use the `label variable` command to assign labels to the variables *rep78*, *price*, *mpg* and *foreign*.

```

1 label variable rep78 "the repair record from 1978"
2 label variable price "the price of the car in 1978"
3 label variable mpg "the miles per gallon for the car"
4 label variable foreign "the origin of the car, foreign or domestic"

```

The `describe` command shows these labels have been applied to the variables.

```

1 describe
2 Contains data from autolab.dta
3 obs:          74                      This file contains auto data for the year
   1978
4 vars:          12                      23 Oct 2008 13:36
5 size:          3,478 (99.9% of memory free)  (_dta has notes)
6 -----
7               storage  display  value
8 variable name  type    format   label    variable label
9 -----
10 make          str18    %-18s
11 price         int      %8.0gc    the price of the car in 1978
12 mpg           int      %8.0g     the miles per gallon for the car
13 rep78         int      %8.0g     the repair record from 1978
14 headroom      float    %6.1f
15 trunk         int      %8.0g
16 weight        int      %8.0gc
17 length        int      %8.0g
18 turn          int      %8.0g
19 displacement  int      %8.0g
20 gear_ratio     float    %6.2f
21 foreign        byte     %8.0g     the origin of the car, foreign or domestic
22 -----
23 Sorted by:

```

Let's make a value label called *foreignl* to label the values of the variable *foreign*. This is a two step process where you first define the label, and then you assign the label to the variable. The `label define` command below creates the value label called *foreignl* that associates 0 with domestic car and 1 with foreign car.

```

1 label define foreignl 0 "domestic car" 1 "foreign car"

```

The `label values` command below associates the variable *foreign* with the label *foreignl*.

```

1 label values foreign foreignl

```

If we use the `describe` command, we can see that the variable *foreign* has a value label called *foreignl* assigned to it.

```

1 describe
2 Contains data from autolab.dta
3 obs:          74                      This file contains auto data for the year
   1978
4 vars:          12                      23 Oct 2008 13:36
5 size:          3,478 (99.9% of memory free)  (_dta has notes)
6 -----
7           storage  display  value
8 variable name  type  format  label  variable label
9 -----
10 make          str18  %-18s
11 price          int   %8.0gc      the price of the car in 1978
12 mpg           int   %8.0g        the miles per gallon for the car
13 rep78          int   %8.0g        the repair record from 1978
14 headroom       float %6.1f
15 trunk          int   %8.0g
16 weight         int   %8.0gc
17 length         int   %8.0g
18 turn           int   %8.0g
19 displacement   int   %8.0g
20 gear_ratio     float %6.2f
21 foreign        byte  %12.0g      foreign1  the origin of the car, foreign or domestic
22 -----
23 Sorted by:

```

Now when we use the `tabulate foreign` command, it shows the labels domestic car and foreign car instead of just 0 and 1.

```

1 table foreign
2 -----+-----
3 the origin |
4 of the car, |
5 foreign or  |
6 domestic   |      Freq.
7 -----+-----
8 domestic car |      52
9 foreign car  |      22
10 -----+-----

```

Value labels are used in other commands as well. For example, below we issue the `ttest` , `by(foreign)` command, and the output labels the groups as domestic and foreign (instead of 0 and 1).

```

1 ttest mpg , by(foreign)
2 Two-sample t test with equal variances
3
4 -----+-----
5 Group |      Obs      Mean    Std. Err.   Std. Dev.   [95% Conf. Interval]
6 -----+-----
7 domestic |      52    19.82692    .6577777    4.743297    18.50638    21.14747
8 foreign  |      22    24.77273    1.40951    6.611187    21.84149    27.70396
9 -----+-----
10 combined |      74    21.2973    .6725511    5.785503    19.9569    22.63769
11 -----+-----
12 diff |           -4.945804    1.362162           -7.661225    -2.230384

```

```

13 -----
14 Degrees of freedom: 72
15
16             Ho: mean(domestic) - mean(foreign) = diff = 0
17
18     Ha: diff < 0 Ha: diff ~= "0" Ha: diff > 0
19         t =  -3.6308             t =  -3.6308             t =  -3.6308
20     P < t =   0.0003             P > |t| =   0.0005             P > t =   0.9997

```

One very important note: These labels are assigned to the data that is currently in memory. To make these changes permanent, you need to **save** the data. When you **save** the data, all of the labels (data labels, variable labels, value labels) will be saved with the data file.

5.1.1 Summary

- Assign a label to the data file currently in memory.

```
1 label data "1978 auto data"
```

- Assign a label to the variable foreign.

```
1 label variable foreign "the origin of the car, foreign or domestic"
```

- Create the value label foreign1 and assign it to the variable foreign.

```
1 label define foreign1 0 "domestic car" 1 "foreign car"
2 label values foreign foreign1
```

5.2 Creating and recoding variables

This module shows how to create and recode variables. In Stata you can create new variables with **generate** and you can modify the values of an existing variable with **replace** and with **recode**.

5.2.1 Computing new variables using generate and replace

Let's **use** the auto data for our examples. In this section we will see how to compute variables with **generate** and **replace**.

```
1 use auto
```

The variable *length* contains the length of the car in inches. Below we see summary statistics for *length*.

```
1 summarize length
```

Variable	Obs	Mean	Std. Dev.	Min	Max
length	74	187.9324	22.26634	142	233

Let's use the **generate** command to make a new variable that has the length in feet instead of inches, called *len_ft*.

```
1 generate len_ft = length / 12
```

We should emphasize that **generate** is for creating a new variable. For an existing variable, you need to use the **replace** command (not **generate**). As shown below, we use **replace** to repeat the assignment to *len_ft*.

```

1 replace len_ft = length / 12
2
3 (49 real changes made)
4
5
6 summarize length len_ft
7
8 Variable |      Obs      Mean   Std. Dev.      Min      Max
9 -----+-----
10 length  |      74   187.9324   22.26634     142     233
11 len_ft  |      74   15.66104   1.855528   11.83333   19.41667

```

The syntax of `generate` and `replace` are identical, except:

- `generate` works when the variable does not yet exist and will give an error if the variable already exists.
- `replace` works when the variable already exists, and will give an error if the variable does not yet exist.

Suppose we wanted to make a variable called *length2* which has *length* squared.

```

1 generate length2 = length^2
2
3 summarize length2
4
5 Variable |      Obs      Mean   Std. Dev.      Min      Max
6 -----+-----
7 length2  |      74  35807.69  8364.045     20164   54289

```

Or we might want to make *loglen* which is the natural log of *length*.

```

1 generate loglen = log(length)
2
3 summarize loglen
4
5 Variable |      Obs      Mean   Std. Dev.      Min      Max
6 -----+-----
7 loglen   |      74    5.229035   .1201383    4.955827   5.451038

```

Let's get the mean and standard deviation of *length* and we can make Z-scores of *length*.

```

1 summarize length
2
3 Variable |      Obs      Mean   Std. Dev.      Min      Max
4 -----+-----
5 length   |      74   187.9324   22.26634     142     233

```

The mean is 187.93 and the standard deviation is 22.27, so *zlength* can be computed as shown below.

`generate zlength = (length - 187.93) / 22.27`

```

1 summarize zlength
2
3 Variable |      Obs      Mean   Std. Dev.      Min      Max
4 -----+-----
5 zlength  |      74    .0001092   .9998357   -2.062416   2.023799

```

With `generate` and `replace`, you can use

- + - for addition and subtraction
- * / for multiplication and division
- ^ for exponents (e.g., `length^2`)

- () for controlling order of operations.

5.2.2 Recoding new variables using generate and replace

Suppose that we wanted to break *mpg* down into three categories. Let's look at a table of *mpg* to see where we might draw the lines for such categories.

```
1 tabulate mpg
```

mpg	Freq.	Percent	Cum.
12	2	2.70	2.70
14	6	8.11	10.81
15	2	2.70	13.51
16	4	5.41	18.92
17	4	5.41	24.32
18	9	12.16	36.49
19	8	10.81	47.30
20	3	4.05	51.35
21	5	6.76	58.11
22	5	6.76	64.86
23	3	4.05	68.92
24	4	5.41	74.32
25	5	6.76	81.08
26	3	4.05	85.14
28	3	4.05	89.19
29	1	1.35	90.54
30	2	2.70	93.24
31	1	1.35	94.59
34	1	1.35	95.95
35	2	2.70	98.65
41	1	1.35	100.00
Total	74	100.00	

Let's convert *mpg* into three categories to help make this more readable. Here we convert *mpg* into three categories using *generate* and *replace*.

```
1 generate mpg3 = .
2 (74 missing values generated)
3
4 replace mpg3 = 1 if (mpg <= 18)
5 (27 real changes made)
6
7 replace mpg3 = 2 if (mpg >= 19) & (mpg <= 23)
8 (24 real changes made)
9
10 replace mpg3 = 3 if (mpg >= 24) & (mpg < .)
11 (23 real changes made)
```

Let's use *tabulate* to check that this worked correctly. Indeed, you can see that a value of 1 for *mpg3* goes from 12-18, a value of 2 goes from 19-23, and a value of 3 goes from 24-41.

```
1 tabulate mpg mpg3
```

mpg	1	2	3	Total
-----	---	---	---	-------

5	-----+-----+-----+-----+-----
6	12 2 0 0 2
7	14 6 0 0 6
8	15 2 0 0 2
9	16 4 0 0 4
10	17 4 0 0 4
11	18 9 0 0 9
12	19 0 8 0 8
13	20 0 3 0 3
14	21 0 5 0 5
15	22 0 5 0 5
16	23 0 3 0 3
17	24 0 0 4 4
18	25 0 0 5 5
19	26 0 0 3 3
20	28 0 0 3 3
21	29 0 0 1 1
22	30 0 0 2 2
23	31 0 0 1 1
24	34 0 0 1 1
25	35 0 0 2 2
26	41 0 0 1 1
27	-----+-----+-----+-----+-----
28	Total 27 24 23 74

Now, we could use `mpg3` to show a crosstab of `mpg3` by `foreign` to contrast the mileage of the foreign and domestic cars.

1 `tabulate mpg3 foreign, column`

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

mpg3	foreign		
	0	1	Total
1	22	5	27
	42.31	22.73	36.49
2	19	5	24
	36.54	22.73	32.43
3	11	12	23
	21.15	54.55	31.08
Total	52	22	74
	100.00	100.00	100.00

The crosstab above shows that 21% of the domestic cars fall into the **high mileage** category, while 55% of the foreign cars fit into this category.

5.2.3 Recoding variables using `recode`

There is an easier way to recode `mpg` to three categories using `generate` and `recode`. First, we make a copy of `mpg`, calling it `mpg3a`. Then, we use `recode` to convert `mpg3a` into three categories: min-18 into 1, 19-23 into 2, and 24-max into 3.

1	<code>generate mpg3a = mpg</code>
2	

```

3 recode mpg3a (min/18=1) (19/23=2) (24/max=3)
4
5 (74 changes made)

```

Let's double check to see that this worked correctly. We see that it worked perfectly.

```

1 tabulate mpg mpg3a
2
3      |
4      mpg |      1      2      3 |      Total
5 -----+-----
6      12 |      2      0      0 |      2
7      14 |      6      0      0 |      6
8      15 |      2      0      0 |      2
9      16 |      4      0      0 |      4
10     17 |      4      0      0 |      4
11     18 |      9      0      0 |      9
12     19 |      0      8      0 |      8
13     20 |      0      3      0 |      3
14     21 |      0      5      0 |      5
15     22 |      0      5      0 |      5
16     23 |      0      3      0 |      3
17     24 |      0      0      4 |      4
18     25 |      0      0      5 |      5
19     26 |      0      0      3 |      3
20     28 |      0      0      3 |      3
21     29 |      0      0      1 |      1
22     30 |      0      0      2 |      2
23     31 |      0      0      1 |      1
24     34 |      0      0      1 |      1
25     35 |      0      0      2 |      2
26     41 |      0      0      1 |      1
27 -----+-----
28     Total |      27      24      23 |      74

```

5.2.4 Recodes with if

Let's create a variable called *mpgfd* that assesses the mileage of the cars with respect to their origin. Let this be a 0/1 variable called *mpgfd* which is:

- 0 if below the median mpg for its group (foreign/domestic)
- 1 if at/above the median mpg for its group (foreign/domestic).

sort foreign

```

1 by foreign: summarize mpg, detail
2
3 -> foreign=      0
4      mpg
5 -----
6      Percentiles      Smallest
7      1%      12      12
8      5%      14      12
9      10%      14      14      Obs      52
10     25%      16.5      14      Sum of Wgt.      52
11
12     50%      19
13      Largest      Mean      19.82692
14      Std. Dev.      4.743297

```



```

14 75%          22          28
15 90%          26          29      Variance      22.49887
16 95%          29          30      Skewness       .7712432
17 99%          34          34      Kurtosis       3.441459
18
19 -> foreign=      1
20
21          mpg
22 -----
23      Percentiles      Smallest
24 1%          14          14
25 5%          17          17      Obs            22
26 10%         17          17      Sum of Wgt.      22
27 25%          21          18
28 50%         24.5
29
30          Largest      Mean            24.77273
31 75%          28          31      Std. Dev.      6.611187
32 90%          35          35      Variance       43.70779
33 95%          35          35      Skewness       .657329
34 99%          41          41      Kurtosis       3.10734

```

We see that the median is 19 for the domestic (`foreign==0`) cars and 24.5 for the foreign (`foreign==1`) cars. The `generate` and `recode` commands below recode `mpg` into `mpgfd` based on the domestic car median for the domestic cars, and based on the foreign car median for the foreign cars.

```

1 generate mpgfd = mpg
2
3 recode mpgfd (min/18=0) (19/max=1) if foreign==0
4
5 (52 changes made)
6
7 recode mpgfd (min/24=0) (25/max=1) if foreign==1
8
9 (22 changes made)

```

We can check using this below, and the recoded value `mpgfd` looks correct.

```

1 by foreign: tabulate mpg mpgfd
2
3 -> foreign=      0
4
5      mpg |      mpgfd
6 -----+-----
7      12 |      2      0 |      2
8      14 |      5      0 |      5
9      15 |      2      0 |      2
10     16 |      4      0 |      4
11     17 |      2      0 |      2
12     18 |      7      0 |      7
13     19 |      0      8 |      8
14     20 |      0      3 |      3
15     21 |      0      3 |      3
16     22 |      0      5 |      5
17     24 |      0      3 |      3
18     25 |      0      1 |      1
19     26 |      0      2 |      2
20     28 |      0      2 |      2

```

21	29	0	1	1
22	30	0	1	1
23	34	0	1	1
24	-----			
25	Total	22	30	52
26				
27				
28	-> foreign= 1			
29		mpgfd		
30	mpg	0	1	Total
31	-----			
32	14	1	0	1
33	17	2	0	2
34	18	2	0	2
35	21	2	0	2
36	23	3	0	3
37	24	1	0	1
38	25	0	4	4
39	26	0	1	1
40	28	0	1	1
41	30	0	1	1
42	31	0	1	1
43	35	0	2	2
44	41	0	1	1
45	-----			
46	Total	11	11	22

5.2.5 Summary

- Create a new variable `len_ft` which is length divided by 12.

```
1 generate len_ft = length / 12
```

- Change values of an existing variable named `len_ft`.

```
1 replace len_ft = length / 12
```

- Recode `mpg` into `mpg3`, having three categories using `generate` and `replace if`.

```
1 generate mpg3 = .
2 replace mpg3 = 1 if (mpg <=18)
3 replace mpg3 = 2 if (mpg >=19) & (mpg <=23)
4 replace mpg3 = 3 if (mpg >=24) & (mpg <.)
```

- Recode `mpg` into `mpg3a`, having three categories, 1 2 3, using `generate` and `recode`.

```
1 generate mpg3a = mpg
2 recode mpg3a (min/18=1) (19/23=2) (24/max=3)
```

- Recode `mpg` into `mpgfd`, having two categories, but using different cutoffs for foreign and domestic cars.

```
1 generate mpgfd = mpg
2 recode mpgfd (min/18=0) (19/max=1) if foreign==0
3 recode mpgfd (min/24=0) (25/max=1) if foreign==1
```

5.3 Subsetting data

This module shows how you can subset data in Stata. You can subset data by keeping or dropping variables, and you can subset data by keeping or dropping observations. You can also subset data as you [use](#) a data file if you are trying to read a file that is too big to fit into the memory on your computer.

5.3.1 Keeping and dropping variables

Sometimes you do not want all of the variables in a data file. You can use the [keep](#) and [drop](#) commands to subset variables. If we think of your data like a spreadsheet, this section will show how you can remove columns (variables) from your data. Let's illustrate this with the auto data file.

```
1 sysuse auto
```

We can use the [describe](#) command to see its variables.

```
1 describe
2
3 Contains data from C:\Program Files\Stata10\ado\base/a/auto.dta
4 obs:          74          1978 Automobile Data
5 vars:          12          13 Apr 2007 17:45
6 size:         3,478 (99.7% of memory free)  (_dta has notes)
7 -----
8           storage  display  value
9 variable name  type   format  label  variable label
10 -----
11 make          str18  %-18s          Make and Model
12 price         int    %8.0gc          Price
13 mpg           int    %8.0g           Mileage (mpg)
14 rep78         int    %8.0g           Repair Record 1978
15 headroom      float  %6.1f           Headroom (in.)
16 trunk         int    %8.0g           Trunk space (cu. ft.)
17 weight        int    %8.0gc          Weight (lbs.)
18 length        int    %8.0g           Length (in.)
19 turn          int    %8.0g           Turn Circle (ft.)
20 displacement  int    %8.0g           Displacement (cu. in.)
21 gear_ratio    float  %6.2f           Gear Ratio
22 foreign       byte   %8.0g           origin  Car type
23 -----
24 Sorted by:  foreign
```

Suppose we want to just have *make* *mpg* and *price*, we can [keep](#) just those variables, as shown below.

```
1 keep make mpg price
```

If we issue the [describe](#) command again, we see that indeed those are the only variables left.

```
1 describe
2
3 Contains data from C:\Program Files\Stata10\ado\base/a/auto.dta
4 obs:          74          1978 Automobile Data
5 vars:          3          13 Apr 2007 17:45
6 size:         1,924 (99.8% of memory free)  (_dta has notes)
7 -----
8           storage  display  value
9 variable name  type   format  label  variable label
10 -----
11 make          str18  %-18s          Make and Model
```

```

12 price          int    %8.0gc          Price
13 mpg           int    %8.0g          Mileage (mpg)
14 -----
15 Sorted by:
16     Note:  dataset has changed since last saved

```

Remember, this has not changed the file on disk, but only the copy we have in memory. If we saved this file calling it `auto`, it would mean that we would replace the existing file (with all the variables) with this file which just has `make`, `mpg` and `price`. In effect, we would permanently lose all of the other variables in the data file. It is important to be careful when using the `save` command after you have eliminated variables, and it is recommended that you save such files to a file with a new name, e.g., `save auto2`. Let's show how to use the `drop` command to drop variables. First, let's clear out the data in memory and `use` the `auto` data file.

```

1 sysuse auto, clear

```

perhaps we are not interested in the variables `displ` and `gear_ratio`. We can get rid of them using the `drop` command shown below.

```

1 drop displ gear_ratio

```

Again, using `describe` shows that the variables have been eliminated.

```

1 describe
2
3 Contains data from C:\Program Files\Stata10\ado\base/a/auto.dta
4 obs:          74          1978 Automobile Data
5 vars:         10          13 Apr 2007 17:45
6 size:         3,034 (99.7% of memory free)  (_dta has notes)
7 -----
8           storage  display  value
9 variable name  type   format  label      variable label
10 -----
11 make          str18  %-18s          Make and Model
12 price         int    %8.0gc         Price
13 mpg           int    %8.0g          Mileage (mpg)
14 rep78         int    %8.0g          Repair Record 1978
15 headroom      float  %6.1f          Headroom (in.)
16 trunk         int    %8.0g          Trunk space (cu. ft.)
17 weight        int    %8.0gc         Weight (lbs.)
18 length        int    %8.0g          Length (in.)
19 turn          int    %8.0g          Turn Circle (ft.)
20 foreign       byte   %8.0g          origin    Car type
21 -----
22 Sorted by:  foreign
23     Note:  dataset has changed since last save

```

If we wanted to make this change permanent, we could save the file as `auto2.dta` as shown below.

```

1 save auto2
2
3 file auto2.dta saved

```

5.3.2 Keeping and dropping observations

The above showed how to use `keep` and `drop` variables to eliminate variables from your data file. The `keep if` and `drop if` commands can be used to keep and drop observations. Thinking of your data like a spreadsheet, the `keep if`

and `drop if` commands can be used to eliminate rows of your data. Let's illustrate this with the auto data. Let's use the auto file and `clear` out the data currently in memory.

```
1 sysuse auto, clear
```

The variable `rep78` has values 1 to 5, and also has some missing values, as shown below.

```
1 tabulate rep78 , missing
2
3      Repair |
4 Record 1978 |      Freq.      Percent      Cum.
5 -----+-----
6           1 |          2       2.70       2.70
7           2 |          8      10.81      13.51
8           3 |         30      40.54      54.05
9           4 |         18      24.32      78.38
10          5 |         11      14.86      93.24
11          . |          5       6.76     100.00
12 -----+-----
13          Total |         74     100.00
```

We may want to eliminate the observations which have missing values using `drop if` as shown below. The portion after the `drop if` specifies which observations that should be eliminated.

```
1 drop if missing(rep78)
2
3 (5 observations deleted)
```

Using the `tabulate` command again shows that these observations have been eliminated.

```
1 tabulate rep78 , missing
2
3      rep78 |      Freq.      Percent      Cum.
4 -----+-----
5           1 |          2       2.90       2.90
6           2 |          8      11.59      14.49
7           3 |         30      43.48      57.97
8           4 |         18      26.09      84.06
9           5 |         11      15.94     100.00
10 -----+-----
11          Total |         69     100.00
```

We could make this change permanent by using the `save` command to save the file. Let's illustrate using `keep if` to eliminate observations. First let's clear out the current file and `use` the auto data file.

```
1 sysuse auto, clear
```

The `keep if` command can be used to eliminate observations, except that the part after the `keep if` specifies which observations should be kept. Suppose we want to keep just the cars which had a repair rating of 3 or less. The easiest way to do this would be using the `keep if` command, as shown below.

```
1 keep if (rep78 <= 3)
2
3 (34 observations deleted)
```

The `tabulate` command shows that this was successful.

```
1 tabulate rep78, missing
2
3      rep78 |      Freq.      Percent      Cum.
```

4	-----+-----			
5	1		2	5.00 5.00
6	2		8	20.00 25.00
7	3		30	75.00 100.00
8	-----+-----			
9	Total		40	100.00

Before we go on to the next section, let's clear out the data that is currently in memory.

```
1 clear
```

5.3.3 Selecting variables and observations with use

The above sections showed how to use `keep`, `drop`, `keep if`, and `drop if` for eliminating variables and observations. Sometimes, you may want to use a data file which is bigger than you can fit into memory and you would wish to eliminate variables and/or observations as you use the file. This is illustrated below with the auto data file. Selecting variables. You can specify just the variables you wish to bring in on the `use` command. For example, let's `use` the auto data file with just *make price* and *mpg*.

```
1 use make price mpg using "http://www.stata-press.com/data/r10/auto.dta"
```

The `describe` command shows us that this worked.

```
1 describe
2
3 Contains data from http://www.stata-press.com/data/r10/auto.dta
4 obs:          74              1978 Automobile Data
5 vars:          3              13 Apr 2007 17:45
6 size:        1,924 (99.8% of memory free)  (_dta has notes)
7
8           storage  display  value
9 variable name  type   format   label    variable label
10 -----
11 make          str18  %-18s           Make and Model
12 price          int    %8.0gc          Price
13 mpg           int    %8.0g           Mileage (mpg)
14 -----
15 Sorted by:
```

Let's clear out the data before the next example.

```
1 clear
```

Suppose we want to just bring in the observations where *rep78* is 3 or less. We can do this as shown below.

```
1 use "http://www.stata-press.com/data/r10/auto.dta" if (rep78 <= 3)
```

We can use `tabulate` to double check that this worked.

```
1 tabulate rep78, missing
```

2				
3	rep78		Freq.	Percent Cum.
4	-----+-----			
5	1		2	5.00 5.00
6	2		8	20.00 25.00
7	3		30	75.00 100.00
8	-----+-----			
9	Total		40	100.00

Let's clear out the data before the next example.

```
1 clear
```

Let's show another example. Lets read in just the cars that had a rating of 4 or higher.

```
1 use "http://www.stata-press.com/data/r10/auto.dta" if (rep78 >= 4) & (rep78 <.)
```

Let's check this using the `tabulate` command.

```
1 tabulate rep78, missing
2
3      rep78 |      Freq.      Percent      Cum.
4 -----+-----
5          4 |          18          62.07          62.07
6          5 |          11          37.93         100.00
7 -----+-----
8      Total |          29         100.00
```

Let's clear out the data before the next example.

```
1 clear
```

You can both eliminate variables and observations with the `use` command. Let's read in just make mpg price and rep78 for the cars with a repair record of 3 or lower.

```
1 use make mpg price rep78 if (rep78 <= 3) using "http://www.stata-press.com/data/r10/auto.dta"
```

Let's check this using `describe` and `tabulate`.

```
1 describe
2
3 Contains data from http://www.stata-press.com/data/r10/auto.dta
4   obs:          40          1978 Automobile Data
5  vars:           4          13 Apr 2007 17:45
6  size:       1,120 (99.9% of memory free)  (_dta has notes)
7
8      storage  display  value
9  variable name  type  format  label  variable label
10 -----
11 make          str18  %-18s          Make and Model
12 price          int   %8.0gc          Price
13 mpg            int   %8.0g          Mileage (mpg)
14 rep78          int   %8.0g          Repair Record 1978
15 -----
16 Sorted by:
17
18 tabulate rep78
19      rep78 |      Freq.      Percent      Cum.
20 -----+-----
21          1 |          2          5.00          5.00
22          2 |          8         20.00         25.00
23          3 |         30         75.00         100.00
24 -----+-----
25      Total |         40         100.00
```

Let's clear out the data before the next example.

```
1 clear
```

Note that the ordering of `if` and `using` is arbitrary.

```
1 use make mpg price rep78 using "http://www.stata-press.com/data/r10/auto.dta" if (rep78
   <= 3)
```

Let's check this using `describe` and `tabulate`.

```
1 describe
2
3 Contains data from http://www.stata-press.com/data/r10/auto.dta
4   obs:           40                1978 Automobile Data
5   vars:           4                13 Apr 2007 17:45
6   size:           1,120 (99.9% of memory free)  (_dta has notes)
7
8 -----
9           storage  display  value
variable name  type  format   label  variable label
10 -----
11 make         str18  %-18s                Make and Model
12 price        int    %8.0gc              Price
13 mpg          int    %8.0g               Mileage (mpg)
14 rep78        int    %8.0g               Repair Record 1978
15 -----
16 Sorted by:
17
18 tabulate rep78
19      rep78 |      Freq.      Percent      Cum.
20 -----+-----
21          1 |          2          5.00          5.00
22          2 |          8         20.00         25.00
23          3 |         30         75.00        100.00
24 -----+-----
25      Total |         40        100.00
```

Have a look at this command. Do you think it will work?

```
1 use make mpg if (rep78 <= 3) using "http://www.stata-press.com/data/r10/auto.dta"
2 rep78 not found
3 r(111);
```

You see, `rep78` was not one of the variables read in, so it could not be used in the `if` portion. To use a variable in the `if` portion, it has to be one of the variables that is read in.

5.3.4 Summary

- Using `keep`/`drop` to eliminate variables

```
1 keep make price mpg
2
3 drop displ gear_ratio
```

- Using `keep if`/`drop if` to eliminate observations

```
1 drop if missing(rep78)
2
3 keep if (rep78 <= 3)
```

- Eliminating variables and/or observations with `use`


```
1 use make mpg price rep78 using auto
2 use auto if (rep78 <= 3)
3
4 use make mpg price rep78 using auto if (rep78 <= 3)
```

6 Intermediate Data Management in Stata

6.1 Collapsing data across observations

Sometimes you have data files that need to be collapsed to be useful to you. For example, you might have student data but you really want classroom data, or you might have weekly data but you want monthly data, etc. We will illustrate this using an example showing how you can collapse data across kids to make family level data.

Here is a file containing information about the kids in three families. There is one record per kid. *Birth* is the order of birth (i.e., 1 is first), *age* *wt* and *sex* are the child's age, weight and sex. We will use this file for showing how to collapse data across observations.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear
2 list
3           famid    kidname    birth    age    wt    sex
4 1.         1      Beth      1        9    60    f
5 2.         1      Bob       2        6    40    m
6 3.         1      Barb      3        3    20    f
7 4.         2      Andy      1        8    80    m
8 5.         2      Al        2        6    50    m
9 6.         2      Ann       3        2    20    f
10 7.         3      Pete      1        6    60    m
11 8.         3      Pam       2        4    40    f
12 9.         3      Phil      3        2    20    m
```

Consider the `collapse` command below. It collapses across all of the observations to make a single record with the average age of the kids.

```
1 collapse age
2 list
3           age
4 1.  5.111111
```

The above `collapse` command was not very useful, but you can combine it with the `by(famid)` option, and then it creates one record for each family that contains the average age of the kids in the family.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear
2 collapse age, by(famid)
3 list
4           famid    age
5 1.         1        6
6 2.         2  5.333333
7 3.         3        4
```

The following `collapse` command does the exact same thing as above, except that the average of age is named `avgage` and we have explicitly told the `collapse` command that we want it to compute the mean.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear
2 collapse (mean) avgage=age, by(famid)
3 list
4           famid    avgage
5 1.         1        6
6 2.         2  5.333333
7 3.         3        4
```

We can request averages for more than one variable. Here we get the average for *age* and for *wt* all in the same command.

```

1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear
2 collapse (mean) avgage=age avgwt=wt, by(famid)
3 list
4      famid      avgage      avgwt
5  1.         1          6         40
6  2.         2      5.333333         50
7  3.         3          4         40

```

This command gets the average of `age` and `wt` like the command above, and also computes `numkids` which is the count of the number of kids in each family (obtained by counting the number of observations with valid values of `birth`).

```

1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear
2 collapse (mean) avgage=age avgwt=wt (count) numkids=birth, by(famid)
3 list
4      famid      avgage      avgwt      numkids
5  1.         1          6         40           3
6  2.         2      5.333333         50           3
7  3.         3          4         40           3

```

Suppose you wanted a count of the number of boys and girls in the family. We can do that with one extra step. We will create a dummy variable that is 1 if the kid is a boy (0 if not), and a dummy variable that is 1 if the kid is a girl (and 0 if not). The sum of the boy dummy variable is the number of boys and the sum of the `girl` dummy variable is the number of girls.

First, let's use the kids file (and clear out the existing data).

```

1 use "http://www.ats.ucla.edu/stat/stata/modules/kids", clear

```

We use `tabulate` with the `generate` option to make the dummy variables.

```

1 tabulate sex, generate(sexdum)
2      sex |      Freq.      Percent      Cum.
3 -----+-----
4      f |          4       44.44       44.44
5      m |          5       55.56      100.00
6 -----+-----
7      Total |          9      100.00

```

We can look at the dummy variables. `Sexdum1` is the dummy variable for girls. `Sexdum2` is the dummy variable for boys. The sum of `sexdum1` is the number of girls in the family. The sum of `sexdum2` is the number of boys in the family.

```

1 list famid sex sexdum1 sexdum2
2      famid      sex      sexdum1      sexdum2
3  1.         1      f          1          0
4  2.         1      m          0          1
5  3.         1      f          1          0
6  4.         2      m          0          1
7  5.         2      m          0          1
8  6.         2      f          1          0
9  7.         3      m          0          1
10  8.         3      f          1          0
11  9.         3      m          0          1

```

The command below creates `girls` which is the number of girls in the family, and `boys` which is the number of boys in the family.

```
1 collapse (count) numkids=birth (sum) girls=sexdum1 boys=sexdum2, by(famid)
```

We can list out the data to confirm that it worked correctly.

```
1 list famid boys girls numkids
2      famid      boys      girls      numkids
3  1.         1         1         2         3
4  2.         2         2         1         3
5  3.         3         2         1         3
```

6.1.1 Summary

- To create one record per family (famid) with the average of age within each family.

```
1 collapse age, by(famid)
```

- To create one record per family (famid) with the average of age (called avgage) and average weight (called avgwt) within each family.

```
1 collapse (mean) avgage=age avgwt=wt, by(famid)
```

- Same as above example, but also counts the number of kids within each family calling that numkids.

```
1 collapse (mean) avgage=age avgwt=wt (count) numkids=birth, by(famid)
```

- Counts the number of boys and girls in each family by using tabulate to create dummy variables based on sex and then summing the dummy variables within each family.

```
1 tabulate sex, generate(sexdum)
2 collapse (sum) girls=sexdum1 boys=sexdum2, by(famid)
```

6.2 Working across variables using foreach

6.2.1 Introduction

This module illustrates

1. how to create and recode variables manually and
2. how to use **foreach** to ease the process of creating and recoding variables.

Consider the sample program below, which reads in income data for twelve months.

```
1 input famid inc1-inc12
2 1 3281 3413 3114 2500 2700 3500 3114 3319 3514 1282 2434 2818
3 2 4042 3084 3108 3150 3800 3100 1531 2914 3819 4124 4274 4471
4 3 6015 6123 6113 6100 6100 6200 6186 6132 3123 4231 6039 6215
5 end
6
7 list
```

The output is shown below

```
1 famid inc1 inc2 inc3 inc4 inc5 inc6 inc7 inc8 inc9 inc10 inc11 inc12
2 1      3281 3413 3114 2500 2700 3500 3114 3319 3514 1282 2434 2818
3 2      4042 3084 3108 3150 3800 3100 1531 2914 3819 4124 4274 4471
4 3      6015 6123 6113 6100 6100 6200 6186 6132 3123 4231 6039 6215
```

6.2.2 Computing variables (manually)

Say that we wanted to compute the amount of tax (10%) paid for each month, the simplest way to do this is to compute 12 variables (*taxinc1-taxinc12*) by multiplying each of the (*inc1-inc12*) by .10 as illustrated below. As you see, this requires entering a command computing the tax for each month of data (for months 1 to 12) via the **generate** command.

```
1 generate taxinc1 = inc1 * .10
2 generate taxinc2 = inc2 * .10
3 generate taxinc3 = inc3 * .10
4 generate taxinc4 = inc4 * .10
5 generate taxinc5 = inc5 * .10
6 generate taxinc6 = inc6 * .10
7 generate taxinc7 = inc7 * .10
8 generate taxinc8 = inc8 * .10
9 generate taxinc9 = inc9 * .10
10 generate taxinc10 = inc10 * .10
11 generate taxinc11 = inc11 * .10
12 generate taxinc12 = inc12 * .10
```

The output is omitted.

6.2.3 Computing variables (using the **foreach** command)

Another way to compute 12 variables representing the amount of tax paid (10%) for each month is to use the **foreach** command. In the example below we use the **foreach** command to cycle through the variables *inc1* to *inc12* and compute the taxable income as *taxinc1 - taxinc12*.

```
1 foreach var of varlist inc1-inc12 {
2     generate tax`var' = `var' * .10
3 }
```

The initial **foreach** statement tells Stata that we want to cycle through the variables *inc1* to *inc12* using the statements that are surrounded by the curly braces. The first time we cycle through the statements, the value of *var* will be *inc1* and the second time the value of *var* will be *inc2* and so on until the final iteration where the value of *var* will be *inc12*. Each statement within the loop (in this case, just the one **generate** statement) is evaluated and executed. When we are inside the **foreach** loop, we can access the value of *var* by surrounding it with the funny quotation marks like this ``var'`. The ``` is the quote right below the `~` on your keyboard and the `'` is the quote below the `"` on your keyboard. The first time through the loop, ``var'` is replaced with *inc1*, so the statement

```
1 generate tax`var' = `var' * .10
```

becomes

```
1 generate taxinc1 = inc1 * .10
```

This is repeated for *inc2* and then *inc3* and so on until *inc12*. So, this **foreach** loop is the equivalent of executing the 12 **generate** statements manually, but much easier and less error prone.

6.2.4 Collapsing across variables (manually)

Often one needs to sum across variables (also known as collapsing across variables). For example, let's say the quarterly income for each observation is desired. In order to get this information, four quarterly variables *incqtr1-incqtr4* need to be computed. Again, this can be achieved manually or by using the **foreach** command. Below is an example of

how to compute 4 quarterly income variables *incqtr1-incqtr4* by simply adding together the months that comprise a quarter.

```
1 generate incqtr1 = inc1 + inc2 + inc3
2 generate incqtr2 = inc4 + inc5 + inc6
3 generate incqtr3 = inc7 + inc8 + inc9
4 generate incqtr4 = inc10+ inc11+ inc12
5
6 list incqtr1 - incqtr4
```

The output is shown below.

```
1      +-----+
2      | incqtr1  incqtr2  incqtr3  incqtr4 |
3      |-----|
4  1. |    9808    8700    9947    6534 |
5  2. |   10234   10050    8264   12869 |
6  3. |   18251   18400   15441   16485 |
7      +-----+
```

6.2.5 Collapsing across variables (using the **foreach** command)

This same result as above can be achieved using the **foreach** command. The example below illustrates how to compute the quarterly income variables *incqtr1-incqtr4* using the **foreach** command.

```
1 foreach qtr of numlist 1/4 {
2   local m3 = `qtr'*3
3   local m2 = (`qtr'*3)-1
4   local m1 = (`qtr'*3)-2
5   generate incqtr`qtr' = inc`m1' + inc`m2' + inc`m3'
6 }
7 list incqtr1 - incqtr4
```

The output is shown below.

```
1      +-----+
2      | incqtr1  incqtr2  incqtr3  incqtr4 |
3      |-----|
4  1. |    9808    8700    9947    6534 |
5  2. |   10234   10050    8264   12869 |
6  3. |   18251   18400   15441   16485 |
7      +-----+
```

In this example, instead of cycling across variables, the **foreach** command is cycling across numbers, 1, 2, 3 then 4 which we refer to as *qtr* which represent the 4 quarters of variables that we wish to create. The trick is the relationship between the quarter and the month numbers that compose the quarter and to create a kind of formula that relates the quarters to the months. For example, quarter 1 of data corresponds to months 3, 2 and 1, so we can say that when the quarter (*qtr*) is 1 we want the months represented by *qtr*3*, (*qtr*3*)-1 and (*qtr*3*)-2, yielding 3, 2, and 1. This is what the statements below from the **foreach** loop are doing. They are relating the quarter to the months.

```
1   local m3 = `qtr'*3
2   local m2 = (`qtr'*3)-1
3   local m1 = (`qtr'*3)-2
```

So, when *qtr* is 1, the value for *m3* is $1 * 3$, the value for *m2* is $(1 * 3) - 1$ and the value for *m1* is $(1 * 3) - 2$. Then, imagine all of those values being substituted into the following statement from the **foreach** loop.

```
1 generate incqtr`qtr' = inc`m1' + inc`m2' + inc`m3'
```

This then becomes

```
1 generate incqtr1 = inc3 + inc2 + inc1
```

and for the next quarter (when *qtr* becomes 2) the statement would become

```
1 generate incqtr2 = inc6 + inc5 + inc4
```

In this example, with only 4 quarters of data, it would probably be easier to simply write out the 4 **generate** statements manually, however if you had 40 quarters of data, then the **foreach** loop can save you considerable time, effort and mistakes.

6.2.6 Identifying patterns across variables (using the **foreach** command)

The **foreach** command can also be used to identify patterns across variables of a dataset. Let's say, for example, that one needs to know which months had income that was less than the income of the previous month. To obtain this information, dummy indicators can be created to indicate in which months this occurred. Note that only 11 dummy indicators are needed for a 12 month period because the interest is in the change from one month to the next. When a month has income that is less than the income of the previous month, the dummy indicators *lowinc2-lowinc12* get assigned a "1". When this is not the case, they are assigned a "0". This program is illustrated below (note for simplicity we assume no missing data on income).

```
1 foreach curmon of numlist 2/12 {
2   local lastmon = `curmon' - 1
3   generate lowinc`curmon' = 1 if ( inc`curmon' < inc`lastmon' )
4   replace lowinc`curmon' = 0 if ( inc`curmon' >= inc`lastmon' )
5 }
```

We can list out the original values of *inc* and *lowinc* and verify that this worked properly

```
1 list famid inc1-inc12, clean noobs
2 famid inc1 inc2 inc3 inc4 inc5 inc6 inc7 inc8 inc9 inc10 inc11 inc12
3 1      3281 3413 3114 2500 2700 3500 3114 3319 3514 1282 2434 2818
4 2      4042 3084 3108 3150 3800 3100 1531 2914 3819 4124 4274 4471
5 3      6015 6123 6113 6100 6100 6200 6186 6132 3123 4231 6039 6215
6 list famid lowinc2-lowinc12, clean noobs
7
8 famid lowinc2 lowinc3 lowinc4 lowinc5 lowinc6 lowinc7 lowinc8 lowinc9 lowinc10 //omitted
9 1          0          1          1          0          0          1          0          0          1
10 2          1          0          0          0          1          1          0          0          0
11 3          0          1          1          0          0          1          1          1          0
```

This time we used the **foreach** loop to compare the current month, represented by *curmon*, and the prior month, computed as *`curmon'-1* creating *lastmon*. So, for the first pass through the **foreach** loop the value for *curmon* is 2 and the value for *lastmon* is 1, so the **generate** and **replace** statements become

```
1 generate lowinc2 = 1 if ( inc2 < inc1 )
2 replace lowinc2 = 0 if ( inc2 >= inc1 )
```

The process is repeated until *curmon* is 12, and then the **generate** and **replace** statements become

```
1 generate lowinc12 = 1 if ( inc12 < inc11 )
2 replace lowinc12 = 0 if ( inc12 >= inc11 )
```

If you were using `foreach` to span a large range of values (say 1/1000) then it is more efficient to use `forvalues` since it is designed to quickly increment through a sequential list, for example

```
1 forvalues curmon = 2/12 {
2     local lastmon = `curmon' - 1
3     generate lowinc`curmon' = 1 if ( inc`curmon' < inc`lastmon' )
4     replace lowinc`curmon' = 0 if ( inc`curmon' >= inc`lastmon' )
5 }
```

6.3 Combining data

This module will illustrate how you can combine files in Stata. Examples will include appending files, one to one match merging, and one to many match merging.

6.3.1 Appending data files

When you have two data files, you may want to combine them by stacking them one on top of the other. For example, we have a file containing *dads* and a file containing *moms* as shown below.

```
1 input famid str4 name inc
2 "Art" 22000
3 "Bill" 30000
4 "Paul" 25000
5 end
6 save dads, replace
7 list
8           famid      name      inc
9   1.          2      Art      22000
10  2.          1      Bill      30000
11  3.          3      Paul      25000
12 clear
13 input famid str4 name inc
14 "Bess" 15000
15 "Pat" 50000
16 "Amy" 18000
17 end
18 save moms, replace
19 list
20           famid      name      inc
21  1.          1      Bess      15000
22  2.          3      Pat       50000
23  3.          2      Amy       18000
```

If we wanted to combine these files by stacking them one atop the other, we can use the `append` command as shown below.

```
1 use dads, clear
2 append using moms
```

We can use the `list` command to see if this worked correctly.

```
1 list
2           famid      name      inc
3   1.          2      Art      22000
4   2.          1      Bill      30000
5   3.          3      Paul      25000
```


6	4.	1	Bess	15000
7	5.	3	Pat	50000
8	6.	2	Amy	18000

The `append` worked properly ...the *dads* and *moms* are stacked together in one file. But, there is a little problem. We can't tell the *dads* from the *moms*. Let's try doing this again, but first we will create a variable called *momdad* in the *dads* and *moms* data file which will contain dad for the *dads* data file and mom for the *moms* data file. When we combine the two files together, the *momdad* variable will tell us who the *moms* and *dads* are.

Here we make *momdad* variable for the *dads* data file. We save the file calling it *dads1*.

```
1 use dads, clear
2 generate str3 momdad = "dad"
3 save dads1
4 file dads1.dta saved
```

Here we make *momdad* variable for the *moms* data file. We `save` the file calling it *moms1*.

```
1 use moms, clear
2 generate str3 momdad = "mom"
3 save moms1
4 file moms1.dta saved
```

Now, let's append *dads1* and *moms1* together.

```
1 use dads1, clear
2 append using moms1
```

Now, when we list the data the *momdad* variable shows who the moms and dads are.

```
1 list
2      famid      name      inc      momdad
3  1.         2      Art      22000      dad
4  2.         1      Bill      30000      dad
5  3.         3      Paul      25000      dad
6  4.         1      Bess      15000      mom
7  5.         3      Pat       50000      mom
8  6.         2      Amy       18000      mom
```

6.3.2 Match merging

Another way of combining data files is match merging. Say that we wanted to combine the *dads* with the *faminc* data file, having the *dads* information and the family information side by side. We can do this with a match merge.

Let's have a look at the *dads* and *faminc* file.

```
1 use dads, clear
2 list
3      famid      name      inc
4  1.         2      Art      22000
5  2.         1      Bill      30000
6  3.         3      Paul      25000
7 clear
8 input famid faminc96 faminc97 faminc98
9 3 75000 76000 77000
10 1 40000 40500 41000
11 2 45000 45400 45800
12 end
13 save faminc, replace
```

```

14 list
15           famid   faminc96   faminc97   faminc98
16 1.           3       75000       76000       77000
17 2.           1       40000       40500       41000
18 3.           2       45000       45400       45800

```

We want to combine the data files so they look like this.

```

1 famid name inc   faminc96 faminc97 faminc98
2 1   Bill 30000 40000    40500    41000
3 2   Art  22000 45000    45400    45800
4 3   Paul 25000 75000    76000    77000

```

Notice that the *famid* variable is used to associate the observation from the *dads* file with the appropriate observation from the *faminc* file. The strategy for merging the files goes like this.

1. sort *dads* on *famid* and *save* that file (calling it *dads2*).
2. sort *faminc* on *famid* and *save* that file (calling it *faminc2*).
3. use the *dads2* file.
4. merge the *dads2* file with the *faminc2* file using *famid* to match them.

Here are those four steps.

1. Sort the *dads* file by *famid* and *save* it as *dads2*

```

1 use dads, clear
2 sort famid
3 save dads2
4 file dads2.dta saved

```

2. Sort the *faminc* file by *famid* and *save* it as *faminc2*.

```

1 use faminc, clear
2 sort famid
3 save faminc2
4 file faminc2.dta saved

```

3. Use the *dads2* file

```

1 use dads2, clear

```

4. Merge with the *faminc2* file using *famid* as the key variable.

```

1 merge famid using faminc2

```

It seems like this worked just fine, but what is that *_merge* variable?

```

1 list, nodisplay noobs
2           famid   name   inc   faminc96   faminc97   faminc98   _merge
3           1   Bill   30000   40000    40500    41000         3
4           2   Art   22000   45000    45400    45800         3
5           3   Paul  25000   75000    76000    77000         3

```

The *_merge* variable indicates, for each observation, how the merge went. This is useful for identifying mismatched records. *_merge* can have one of three values

1. The record contains information from file1 only (e.g., a *dad2* record with no corresponding *faminc2* record).
2. The record contains information from file2 only (e.g., a *faminc2* record with no corresponding *dad2* record).
3. The record contains information from both files (e.g., the *dad2* and *faminc2* records matched up).

When you have many records, tabulating `_merge` is very useful to summarize how many mismatched you have. In our case, all of the records match so the value for `_merge` was always 3.

1	<code>tabulate _merge</code>			
2	<code>_merge</code>	<code>Freq.</code>	<code>Percent</code>	<code>Cum.</code>
3	-----+-----			
4	3	3	100.00	100.00
5	-----+-----			
6	Total	3	100.00	

6.3.3 One-to-many match merging

Another kind of merge is called a *one to many* merge. Our *one to one* merge matched up *dads* and *faminc* and there was a one to one matching of the files. If we merge *dads* with *kids*, there can be multiple kids per dad and hence this is a *one to many* merge.

As you see below, the strategy for the *one to many* merge is really the same as the *one to one* merge.

1. sort *dads* on *famid* and `save` that file as *dads3*
2. sort *kids* on *famid* and `save` that file as *kids3*
3. use the *dads3* file
4. merge the *dads3* file with the *kids3* file using *famid* to match them.

The 4 steps are shown below.

1. Sort the *dads* data file on *famid* and `save` that file as *dads3*.

```

1 use dads, clear
2 sort famid
3 save dads3
4 file dads3.dta saved
5 list
6      famid      name      inc
7  1.         1      Bill    30000
8  2.         2       Art    22000
9  3.         3      Paul    25000

```

2. Sort the *kids* data file on *famid* and `save` that file as *kids3*.

```

1 clear
2 input famid str4 kidname birth age wt str1 sex
3 1 "Beth" 1 9 60 "f"
4 2 "Andy" 1 8 40 "m"
5 3 "Pete" 1 6 20 "f"
6 1 "Bob" 2 6 80 "m"
7 1 "Barb" 3 3 50 "m"
8 2 "Al" 2 6 20 "f"
9 2 "Ann" 3 2 60 "m"
10 3 "Pam" 2 4 40 "f"
11 3 "Phil" 3 2 20 "m"
12 end
13
14 sort famid
15
16 save kids3
17 file kids3.dta saved
18
19 list
20      famid      kidname      birth      age      wt      sex

```

21	1.	1	Beth	1	9	60	f
22	2.	1	Bob	2	6	40	m
23	3.	1	Barb	3	3	20	f
24	4.	2	Andy	1	8	80	m
25	5.	2	Al	2	6	50	m
26	6.	2	Ann	3	2	20	f
27	7.	3	Pete	1	6	60	m
28	8.	3	Pam	2	4	40	f
29	9.	3	Phil	3	2	20	m

3. Use the *dads3* file.

```
1 use dads3, clear
```

4. Merge the *dads3* file with the *kids3* file using *famid* to match them.

```
1 merge famid using kids3
```

Let's list out the results.

```
1 list famid name kidname birth age _merge
2      famid  name  kidname  birth  age  _merge
3  1.      1  Bill    Barb    3    3    3
4  2.      2  Art     Al      2    6    3
5  3.      3  Paul    Pam     2    4    3
6  4.      1  Bill    Bob     2    6    3
7  5.      1  Bill    Beth    1    9    3
8  6.      2  Art     Andy    1    8    3
9  7.      2  Art     Ann     3    2    3
10 8.      3  Paul    Phil    3    2    3
11 9.      3  Paul    Pete    1    6    3
```

The results are a bit easier to read if we sort the data on *famid* and *birth*.

```
1 sort famid birth
2 list famid name kidname birth age _merge
3      famid  name  kidname  birth  age  _merge
4  1.      1  Bill    Beth    1    9    3
5  2.      1  Bill    Bob     2    6    3
6  3.      1  Bill    Barb    3    3    3
7  4.      2  Art     Andy    1    8    3
8  5.      2  Art     Al      2    6    3
9  6.      2  Art     Ann     3    2    3
10 7.      3  Paul    Pete    1    6    3
11 8.      3  Paul    Pam     2    4    3
12 9.      3  Paul    Phil    3    2    3
```

As you see, this is basically the same as a *one to one* merge. You may wonder if the order of the files on the merge statement is relevant. Here, we switch the order of the files and the results are the same. The only difference is the order of the records after the merge.

```
1 use kids3, clear
2 merge famid using dads3
3 list famid name kidname birth age
4      famid  name  kidname  birth  age
5  1.      1  Bill    Beth    1    9
6  2.      1  Bill    Bob     2    6
7  3.      1  Bill    Barb    3    3
```

8	4.	2	Art	Andy	1	8
9	5.	2	Art	Al	2	6
10	6.	2	Art	Ann	3	2
11	7.	3	Paul	Pete	1	6
12	8.	3	Paul	Pam	2	4
13	9.	3	Paul	Phil	3	2

6.3.4 Summary

- Appending data example

```
1 use dads, clear
2 append using moms
```

- Match merge example steps (one-to-one and one-to-many)
 1. sort *dads* on *famid* and save that file
 2. sort *kids* on *famid* and save that file
 3. use the *dads* file
 4. merge the *dads* file with the *kids* file using *famid* to match them.

- Match merge example program

```
1 use dads, clear
2 sort famid
3 save dads2
4
5 use faminc, clear
6 sort famid
7 save faminc2
8
9 use dads2, clear
10 merge famid using faminc2
```

6.4 Reshaping data wide to long

This module illustrates the power (and simplicity) of Stata in its ability to reshape data files. These examples take **wide** data files and reshape them into **long** form. These show common examples of reshaping data, but do not exhaustively demonstrate the different kinds of data reshaping that you could encounter.

6.4.1 Example 1: Reshaping data wide to long

Consider the family income data file below.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/faminc.dta", clear
2 list
3
4      famid  faminc96  faminc97  faminc98
5  1.      3      75000      76000      77000
6  2.      1      40000      40500      41000
7  3.      2      45000      45400      45800
```

This is called a **wide** format since the years of data are wide. We may want the data to be **long**, where each year of data is in a separate observation. The **reshape** command can accomplish this, as shown below.

```
1 reshape long faminc, i(famid) j(year)
2 (note: j = 96 97 98)
3
```

```

4 Data                                wide  ->  long
5 -----
6 Number of obs.                      3  ->    9
7 Number of variables                  4  ->    3
8 j variable (3 values)                ->  year
9 xij variables:
10      faminc96 faminc97 faminc98  ->  faminc
11 -----

```

The `list` command shows that the data are now in **long** form, where each *year* is represented as its own observation.

```

1 list
2      famid      year      faminc
3  1.         1        96      40000
4  2.         1        97      40500
5  3.         1        98      41000
6  4.         2        96      45000
7  5.         2        97      45400
8  6.         2        98      45800
9  7.         3        96      75000
10 8.         3        97      76000
11 9.         3        98      77000

```

Let's look at the **wide** format and contrast it with the **long** format.

The `reshape wide` command puts the data back into **wide** format. We then list out the **wide** file.

```

1 reshape wide
2 (note: j = 96 97 98)
3
4 Data                                long  ->  wide
5 -----
6 Number of obs.                      9  ->    3
7 Number of variables                  3  ->    4
8 j variable (3 values)                year  ->  (dropped)
9 xij variables:
10      faminc  ->  faminc96 faminc97 faminc98
11 -----
12 list
13      famid  faminc96  faminc97  faminc98
14  1.         1      40000      40500      41000
15  2.         2      45000      45400      45800
16  3.         3      75000      76000      77000

```

The `reshape long` command puts the data back into **long** format. We then list out the **long** file.

```

1 reshape long
2 (note: j = 96 97 98)
3
4 Data                                wide  ->  long
5 -----
6 Number of obs.                      3  ->    9
7 Number of variables                  4  ->    3
8 j variable (3 values)                ->  year
9 xij variables:
10      faminc96 faminc97 faminc98  ->  faminc
11 -----
12 list

```

	famid	year	faminc
13			
14	1.	1	96
15	2.	1	97
16	3.	1	98
17	4.	2	96
18	5.	2	97
19	6.	2	98
20	7.	3	96
21	8.	3	97
22	9.	3	98

Now let's look at the pieces of the original `reshape` command.

```
1 reshape long faminc, i(famid) j(year)
```

- `long` tells `reshape` that we want to go from wide to long
- `faminc` tells Stata that the `stem` of the variable to be converted from **wide** to **long** is `faminc`
- `i(famid)` option tells `reshape` that `famid` is the unique identifier for records in their **wide** format
- `j(year)` tells `reshape` that the suffix of `faminc` (i.e., 96 97 98) should be placed in a variable called `year`

6.4.2 Example 2: Reshaping data wide to long

Consider the file containing the kids and their heights at 1 year of age (`ht1`) and at 2 years of age (`ht2`).

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/kidshtwt.dta", clear
2 list famid birth ht1 ht2
3
4      famid      birth      ht1      ht2
5  1.         1         1       2.8     3.4
6  2.         1         2       2.9     3.8
7  3.         1         3       2.2     2.9
8  4.         2         1         2     3.2
9  5.         2         2       1.8     2.8
10 6.         2         3       1.9     2.4
11 7.         3         1       2.2     3.3
12 8.         3         2       2.3     3.4
13 9.         3         3       2.1     2.9
```

Lets reshape this data into a long format. The critical questions are:

Q: What is the stem of the variable going from **wide** to long.

A: The stem is `ht`

Q: What variable uniquely identifies an observation when it is in the **wide** form.

A: `famid` and `birth` together uniquely identify the **wide** observations.

Q: What do we want to call the variable which contains the suffix of `ht`, i.e., 1 and 2.

A: Lets call the suffix `age`.

With the answers to these questions, the reshape command will look like this.

```
1 reshape long ht, i(famid birth) j(age)
```

Let's look at the **wide** data, and then the data reshaped to be **long**.

```
1 list famid birth ht1 ht2
2
3      famid      birth      ht1      ht2
4  1.         1         1       2.8     3.4
5  2.         1         2       2.9     3.8
6  3.         1         3       2.2     2.9
7  4.         2         1         2     3.2
```

```

7      5.      2      2      1.8      2.8
8      6.      2      3      1.9      2.4
9      7.      3      1      2.2      3.3
10     8.      3      2      2.3      3.4
11     9.      3      3      2.1      2.9
12 reshape long ht, i(famid birth) j(age)
13 (note: j = 1 2)
14
15 Data                                wide  ->  long
16 -----
17 Number of obs.                      9  ->    18
18 Number of variables                  7  ->    7
19 j variable (2 values)                ->    age
20 xij variables:
21                                ht1 ht2  ->   ht
22 -----
23 list famid birth age ht
24      famid      birth      age      ht
25  1.      1      1      1      2.8
26  2.      1      1      2      3.4
27  3.      1      2      1      2.9
28  4.      1      2      2      3.8
29  5.      1      3      1      2.2
30  6.      1      3      2      2.9
31  7.      2      1      1      2
32  8.      2      1      2      3.2
33  9.      2      2      1      1.8
34 10.      2      2      2      2.8
35 11.      2      3      1      1.9
36 12.      2      3      2      2.4
37 13.      3      1      1      2.2
38 14.      3      1      2      3.3
39 15.      3      2      1      2.3
40 16.      3      2      2      3.4
41 17.      3      3      1      2.1
42 18.      3      3      2      2.9

```

6.4.3 Example 3: Reshaping data wide to long

The file with the kids heights at *age 1* and *age 2* also contains their weights at *age 1* and *age 2* (called *wt1* and *wt2*).

```

1 use "http://www.ats.ucla.edu/stat/stata/modules/kidshtwt.dta", clear
2 list famid birth ht1 ht2 wt1 wt2
3      famid      birth      ht1      ht2      wt1      wt2
4  1.      1      1      2.8      3.4      19      28
5  2.      1      2      2.9      3.8      21      28
6  3.      1      3      2.2      2.9      20      23
7  4.      2      1      2      3.2      25      30
8  5.      2      2      1.8      2.8      20      33
9  6.      2      3      1.9      2.4      22      33
10  7.      3      1      2.2      3.3      22      28
11  8.      3      2      2.3      3.4      20      30
12  9.      3      3      2.1      2.9      22      31

```

Let's reshape this data into a **long** format. This is basically the same as the previous command except that *ht* is replaced with *wt*.


```
1 reshape long ht wt, i(famid birth) j(age)
```

Let's look at the **wide** data, and then the data reshaped to be **long**.

```
1 list famid birth ht1 ht2 wt1 wt2
2
3      famid      birth      ht1      ht2      wt1      wt2
4  1.         1         1       2.8       3.4       19       28
5  2.         1         2       2.9       3.8       21       28
6  3.         1         3       2.2       2.9       20       23
7  4.         2         1         2       3.2       25       30
8  5.         2         2       1.8       2.8       20       33
9  6.         2         3       1.9       2.4       22       33
10 7.         3         1       2.2       3.3       22       28
11 8.         3         2       2.3       3.4       20       30
12 9.         3         3       2.1       2.9       22       31
13 reshape long ht wt, i(famid birth) j(age)
14 (note: j = 1 2)
15
16 Data                                wide  ->  long
17 -----
18 Number of obs.                      9  ->    18
19 Number of variables                  7  ->     6
20 j variable (2 values)                ->   age
21 xij variables:
22      ht1 ht2  ->  ht
23      wt1 wt2  ->  wt
24 -----
25 list famid birth age ht wt
26
27      famid      birth      age      ht      wt
28  1.         1         1         1       2.8       19
29  2.         1         1         2       3.4       28
30  3.         1         2         1       2.9       21
31  4.         1         2         2       3.8       28
32  5.         1         3         1       2.2       20
33  6.         1         3         2       2.9       23
34  7.         2         1         1         2       25
35  8.         2         1         2       3.2       30
36  9.         2         2         1       1.8       20
37 10.         2         2         2       2.8       33
38 11.         2         3         1       1.9       22
39 12.         2         3         2       2.4       33
40 13.         3         1         1       2.2       22
41 14.         3         1         2       3.3       28
42 15.         3         2         1       2.3       20
43 16.         3         2         2       3.4       30
44 17.         3         3         1       2.1       22
45 18.         3         3         2       2.9       31
```

6.4.4 Example 4: Reshaping data wide to long with character suffixes

It also is possible to reshape a wide data file to be long when there are character suffixes. Look at the `dadmomw` file below.

```
1 use "http://www.ats.ucla.edu/stat/stata/modules/dadmomw.dta", clear
2 list
3      famid      named      incd      namem      incm
```

4	1.	1	Bill	30000	Bess	15000
5	2.	2	Art	22000	Amy	18000
6	3.	3	Paul	25000	Pat	50000

We would like to make *name* and *inc* into **long** formats but their suffixes are characters (d & m) instead of numbers. Stata can handle that as long as you use **string** in the command to indicate that the suffix is a character.

```
1 reshape long name inc, i(famid) j(dadmom) string
```

Let's look at the data before and after reshaping.

```
1 list
2      famid      named      incd      namem      incm
3      1.         1      Bill      30000      Bess      15000
4      2.         2       Art      22000       Amy      18000
5      3.         3      Paul      25000       Pat      50000
6 reshape long name inc, i(famid) j(dadmom) string
7 (note: j = d m)
8
9 Data
10
11 ----- wide -> long -----
12 Number of obs.              3 -> 6
13 Number of variables         5 -> 4
14 j variable (2 values)      -> dadmom
15 xij variables:
16      named namem -> name
17      incd incm -> inc
18
19 list
20      famid      dadmom      name      inc
21      1.         1         d      Bill      30000
22      2.         1         m      Bess      15000
23      3.         2         d      Art       22000
24      4.         2         m      Amy       18000
25      5.         3         d      Paul      25000
26      6.         3         m      Pat       50000
```

6.4.5 Summary reshaping data wide to long

```
1 Wide format
2      famid      faminc96      faminc97      faminc98
3      1.         1         40000         40500         41000
4      2.         2         45000         45400         45800
5      3.         3         75000         76000         77000
6
7 reshape long faminc, i(famid) j(year)
8
9 Long Format
10      famid      year      faminc
11      1.         1         96      40000
12      2.         1         97      40500
13      3.         1         98      41000
14      4.         2         96      45000
15      5.         2         97      45400
16      6.         2         98      45800
17      7.         3         96      75000
```

18	8.	3	97	76000
19	9.	3	98	77000

The general syntax of `reshape long` can be expressed as...

```
1 reshape long stem-of-wide-vars, i(wide-id-var) j(var-for-suffix)
```

where

- stem-of-wide-vars: is the stem of the wide variables, e.g., faminc
- wide-id-var: is the variable that uniquely identifies wide observations, e.g., famid
- var-for-suffix: is the variable that will contain the suffix of the wide variables, e.g., year

6.5 Reshaping data long to wide

This module illustrates the power (and simplicity) of Stata in its ability to reshape data files. These examples take **long** data files and reshape them into **wide** form. These examples cover some common examples, but this is only part of the features and options of the Stata `reshape` command.

6.5.1 Example 1: Reshaping data long to wide

The `reshape` command can be used to make data from a **long** format to a **wide** format. Consider the `kids` file (to make things simple at first, we will drop the variables `kidname`, `sex` and `wt`).

```
1 use kids, clear
2 drop kidname sex wt
3 list
4           famid      birth      age
5  1.           1          1          9
6  2.           1          2          6
7  3.           1          3          3
8  4.           2          1          8
9  5.           2          2          6
10 6.           2          3          2
11 7.           3          1          6
12 8.           3          2          4
13 9.           3          3          2
```

Let's make `age` in this file wide, making one record per family which would contain `age1 age2 age3`, the ages of the kids in the family (`age2` would be missing if there is only one kid, and `age3` would be missing if there are only two kids). Let's look at the data before and after reshaping.

```
1 list
2           famid      birth      age
3  1.           1          1          9
4  2.           1          2          6
5  3.           1          3          3
6  4.           2          1          8
7  5.           2          2          6
8  6.           2          3          2
9  7.           3          1          6
10 8.           3          2          4
11 9.           3          3          2
12
13 reshape wide age, i(famid) j(birth)
14
15 (note:  j = 1 2 3)
```

16					
17	Data		long	->	wide
18					
19	Number of obs.		9	->	3
20	Number of variables		3	->	4
21	j variable (3 values)		birth	->	(dropped)
22	xij variables:				
23			age	->	age1 age2 age3
24					
25					
26	list				
27					
28		famid	age1	age2	age3
29	1.	1	9	6	3
30	2.	2	8	6	2
31	3.	3	6	4	2

Let's look at the pieces of the reshape command.

```
1 reshape wide age, j(birth) i(famid)
```

- **wide**: tells reshape that we want to go from long to wide
- **age**: tells Stata that the variable to be converted from long to wide is *age*
- **i(famid)**: tells reshape that *famid* uniquely identifies observations in the wide form
- **j(birth)**: tells reshape that the suffix of *age* (1 2 3) should be taken from the variable *birth*

6.5.2 Example 2: Reshaping data long to wide with more than one variable

The **reshape** command can work on more than one variable at a time. In the example above, we just reshaped the *age* variable. In the example below, we reshape the variables *age*, *wt* and *sex* like this

```
1 reshape wide age wt sex, i(famid) j(birth)
```

Let's look at the data before and after reshaping.

1	use kids, clear					
2	list					
3		famid	kidname	birth	age	wt
4	1.	1	Beth	1	9	60
5	2.	1	Bob	2	6	40
6	3.	1	Barb	3	3	20
7	4.	2	Andy	1	8	80
8	5.	2	Al	2	6	50
9	6.	2	Ann	3	2	20
10	7.	3	Pete	1	6	60
11	8.	3	Pam	2	4	40
12	9.	3	Phil	3	2	20
13						
14	reshape wide kidname age wt sex, i(famid) j(birth)					
15						
16	(note: j = 1 2 3)					
17						
18	Data		long	->	wide	
19						
20	Number of obs.		9	->	3	
21	Number of variables		6	->	13	
22	j variable (3 values)		birth	->	(dropped)	

```

23 xij variables:
24             kidname  ->  kidname1 kidname2 kidname3
25             age      ->  age1 age2 age3
26             wt       ->  wt1 wt2 wt3
27             sex      ->  sex1 sex2 sex3
28 -----
29
30 list
31
32 Observation 1
33
34     famid      1    kidname1    Beth    age1      9
35     wt1       60    sex1       f    kidname2    Bob
36     age2      6    wt2       40    sex2       m
37     kidname3  Barb    age3      3    wt3       20
38     sex3      f
39
40
41 Observation 2
42
43     famid      2    kidname1    Andy    age1      8
44     wt1       80    sex1       m    kidname2    Al
45     age2      6    wt2       50    sex2       m
46     kidname3  Ann    age3      2    wt3       20
47     sex3      f
48
49
50 Observation 3
51
52     famid      3    kidname1    Pete    age1      6
53     wt1       60    sex1       m    kidname2    Pam
54     age2      4    wt2       40    sex2       f
55     kidname3  Phil    age3      2    wt3       20
56     sex3      m

```

6.5.3 Example 3: Reshaping wide with character suffixes

The examples above showed how to reshape data using numeric suffixes, but `reshape` can handle character suffixes as well. Consider the `dadmom1` data file shown below.

```

1 use dadmom1, clear
2 list
3
4     famid      name      inc      dadmom
5  1.      2      Art      22000      dad
6  2.      1      Bill     30000      dad
7  3.      3      Paul     25000      dad
8  4.      1      Bess     15000      mom
9  5.      3      Pat      50000      mom
10 6.      2      Amy      18000      mom

```

Let's reshape this to be in a wide format, containing one record per family. The `reshape` command below uses `string` to tell reshape that the suffix is character.

```

1 reshape wide name inc, i(famid) j(dadmom) string

```

Let's look at the data before and after reshaping.

```

1 list
2      famid      name      inc      dadmom
3  1.         2      Art     22000      dad
4  2.         1      Bill    30000      dad
5  3.         3      Paul    25000      dad
6  4.         1      Bess    15000      mom
7  5.         3      Pat     50000      mom
8  6.         2      Amy     18000      mom
9
10 reshape wide name inc, i(famid) j(dadmom) string
11
12 (note: j = dad mom)
13
14 Data                                long  ->  wide
15 -----
16 Number of obs.                      6    ->    3
17 Number of variables                  4    ->    5
18 j variable (2 values)               dadmom -> (dropped)
19 xij variables:
20                                name  ->  namedad namemom
21                                inc   ->  incdad incmom
22 -----
23
24 list
25
26      famid      namedad      incdad      namemom      incmom
27  1.         1      Bill     30000      Bess     15000
28  2.         2      Art     22000      Amy     18000
29  3.         3      Paul    25000      Pat     50000

```

6.5.4 Summary

- Reshaping data long to wide

```

1 Long format
2      famid      birth      age
3  1.         1         1         9
4  2.         1         2         6
5  3.         1         3         3
6  4.         2         1         8
7  5.         2         2         6
8  6.         2         3         2
9  7.         3         1         6
10 8.         3         2         4
11 9.         3         3         2
12
13 reshape wide age, j(birth) i(famid)
14
15 Wide format
16      famid      age1      age2      age3
17  1.         1         9         6         3
18  2.         2         8         6         2
19  3.         3         6         4         2

```

- The general syntax of reshape wide can be expressed as:

```
1 reshape wide long-var(s), i( wide-id-var ) j( var-with-suffix )
```

where

- long-var(s): is the name of the long variable(s) to be made wide e.g. age
- wide-id-var: is the variable that uniquely identifies wide observations, e.g. famid
- var-with-suffix: is the variable from the long file that contains the suffix for the wide variables, e.g. birth